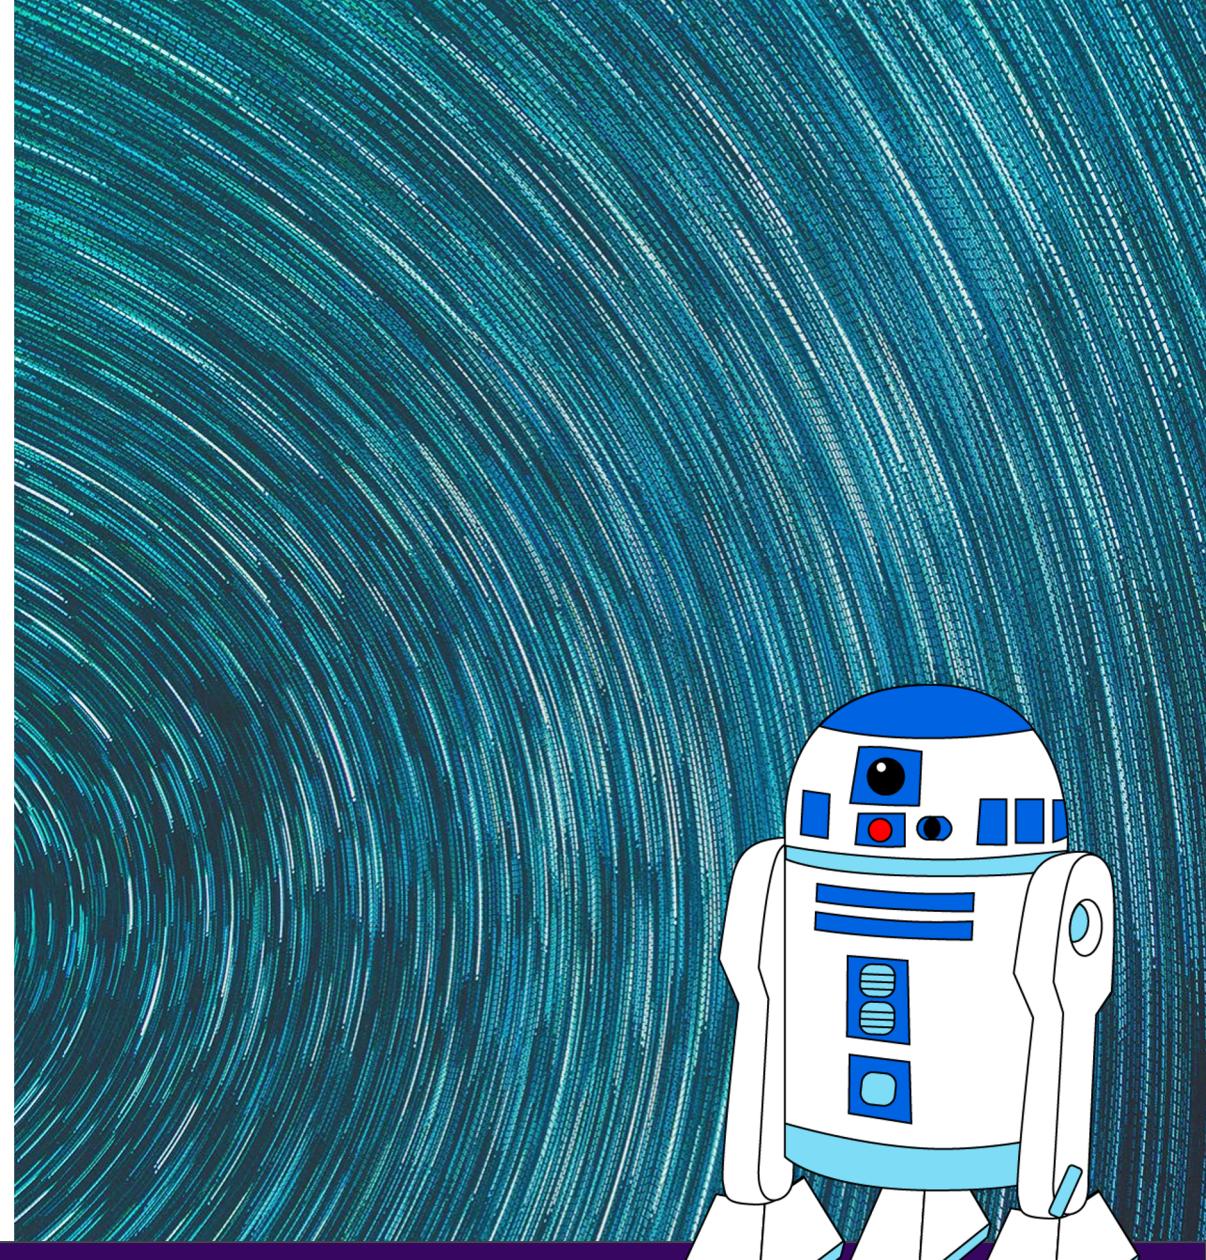


CIS 521:
ARTIFICIAL INTELLIGENCE

Expectimax and Utilities

Professor Chris Callison-Burch

Many of today's slides are courtesy of Dan Klein and
Pieter Abbeel of University of California, Berkeley



Review: Adversarial Search (Minimax)

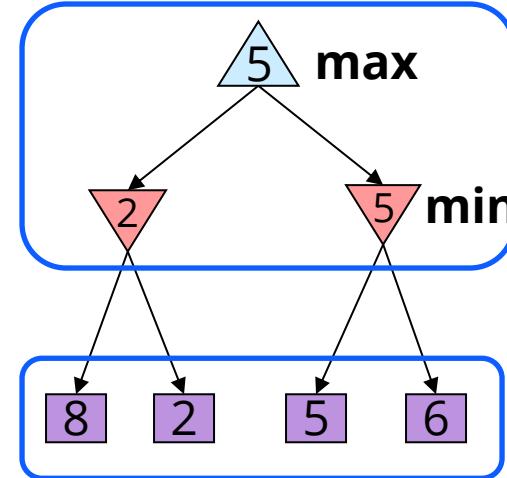
- Minimax search:

- A state-space search tree

- Players alternate turns

- Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary

**Minimax values:
computed recursively**



**Terminal values:
part of the game**

Review: Minimax Implementation

```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

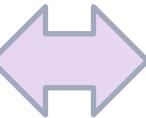


```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

Minimax Example



```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

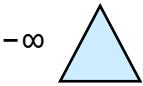


```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



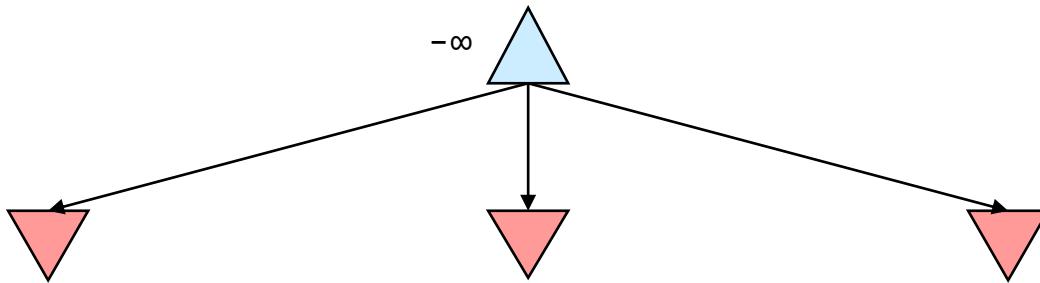
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



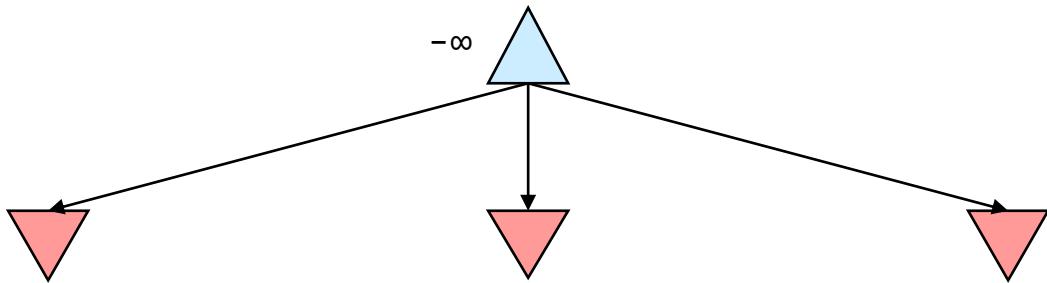
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



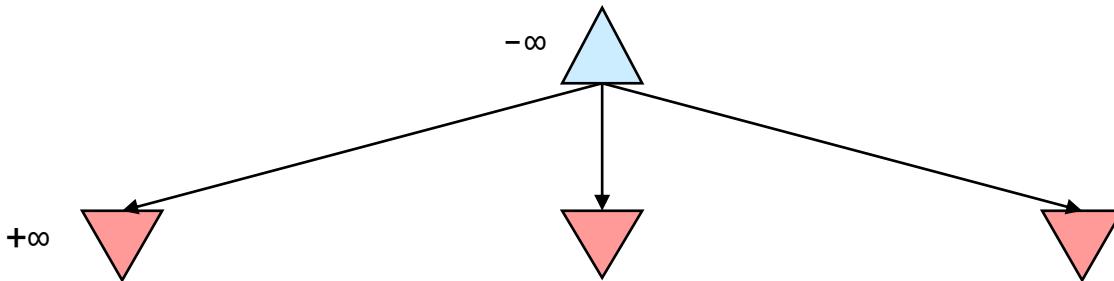
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



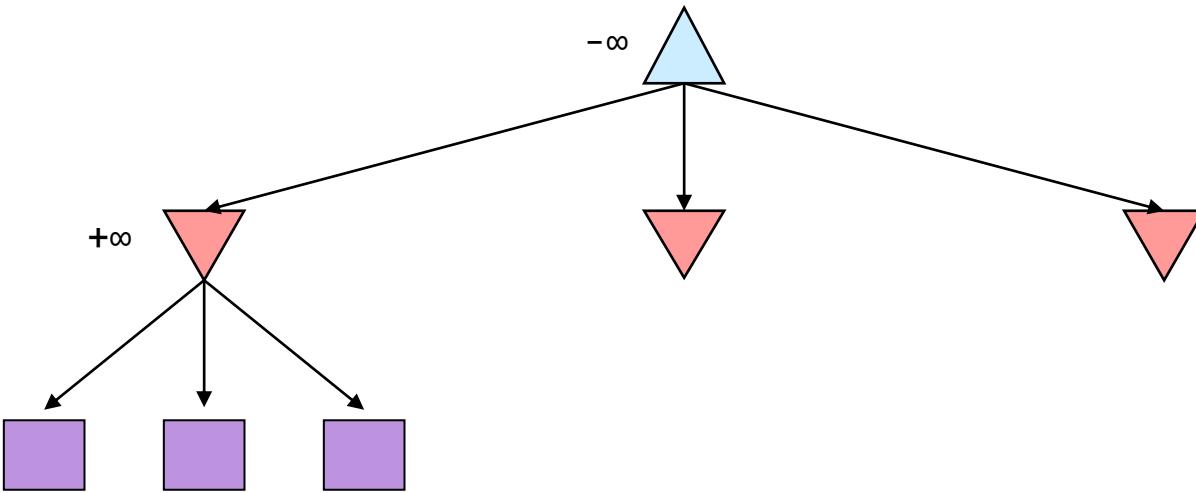
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = −∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



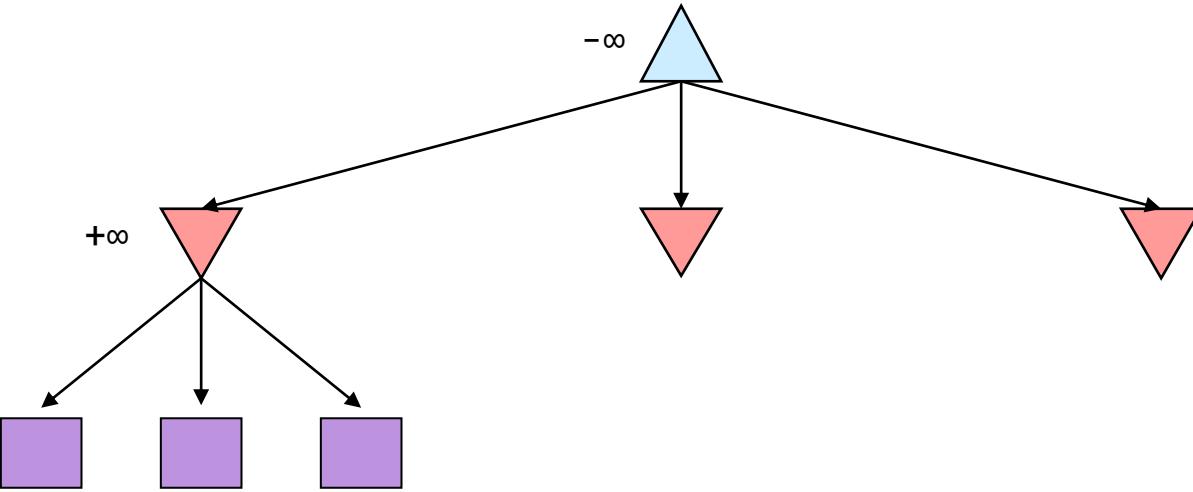
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



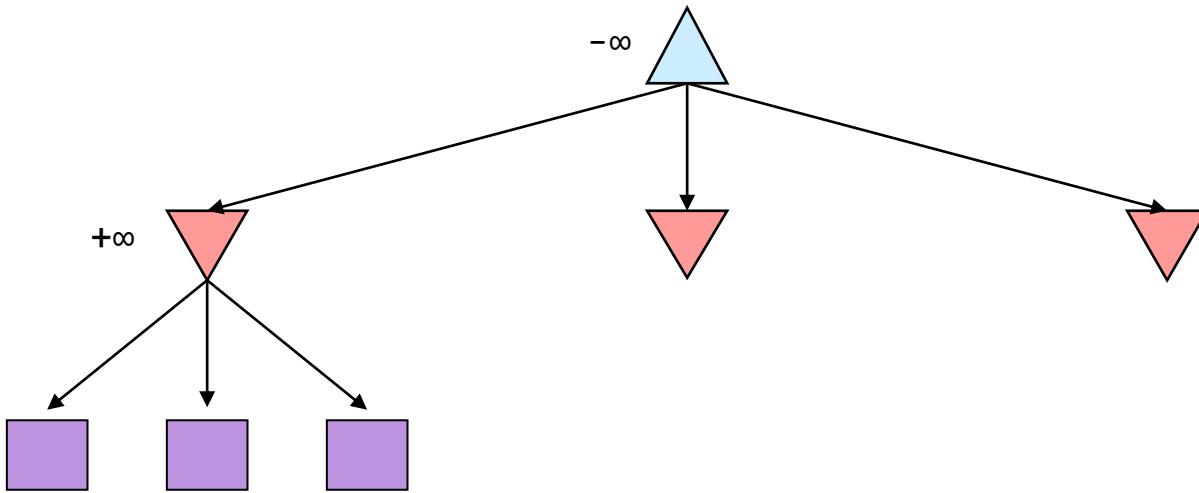
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



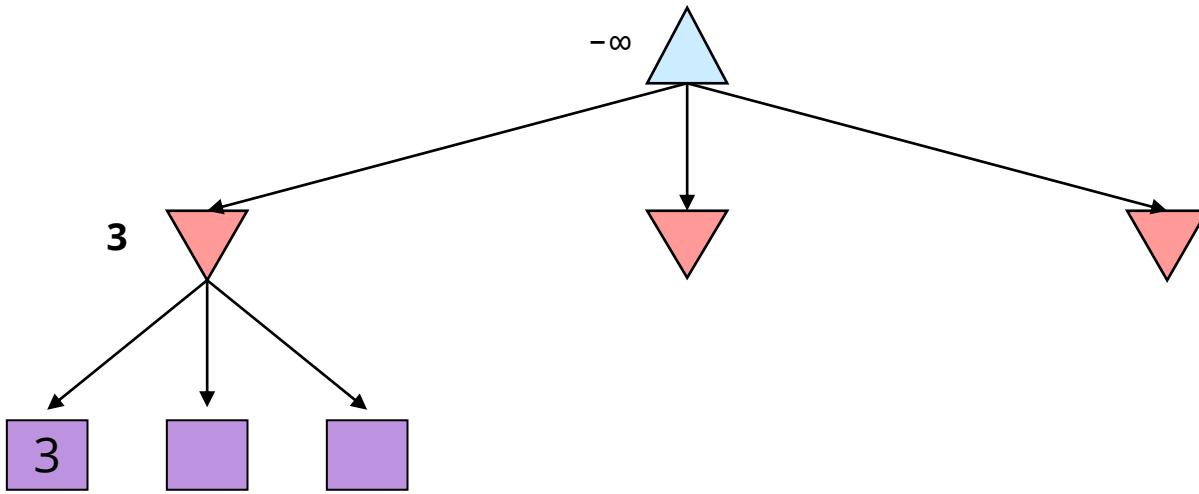
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



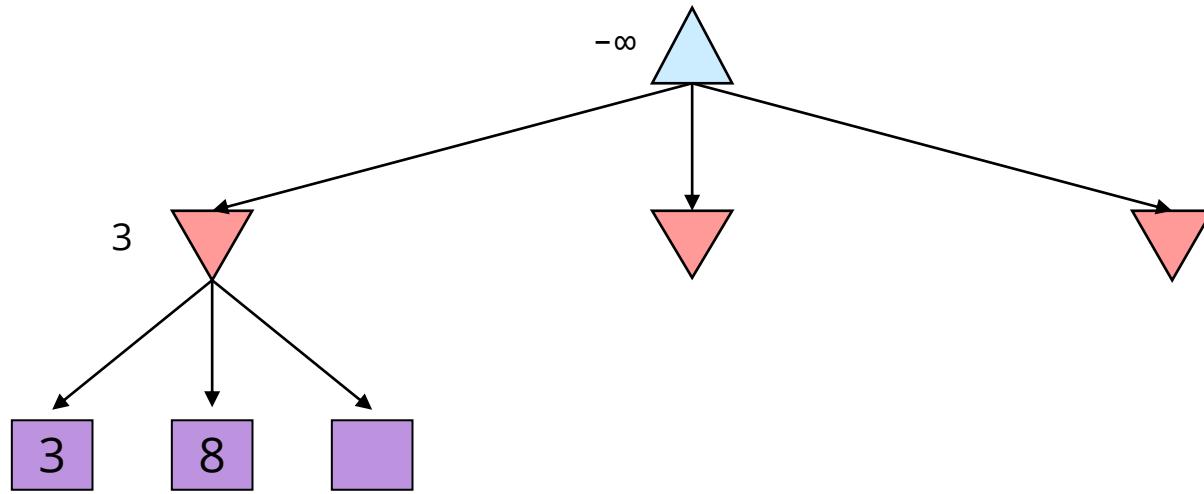
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



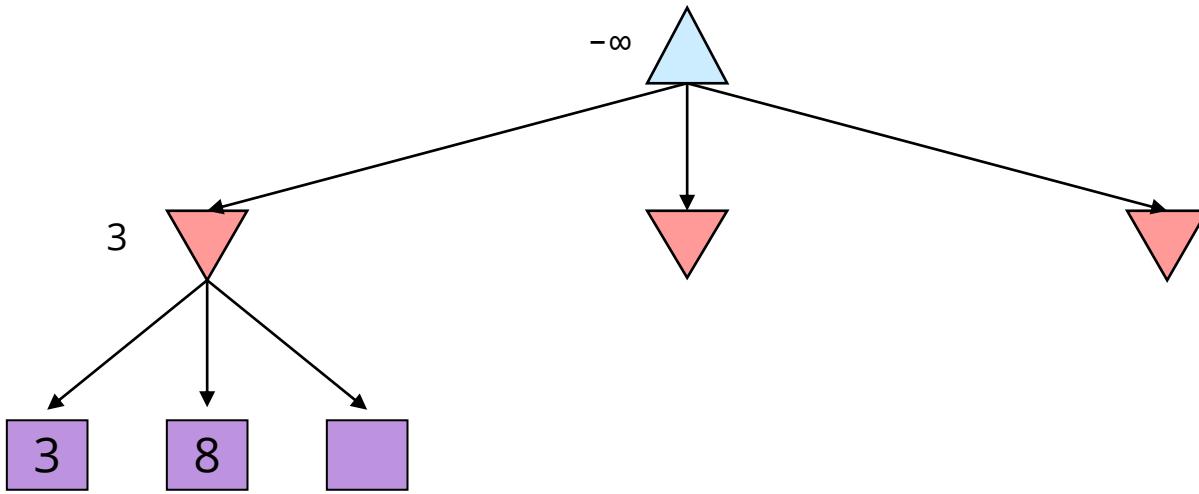
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



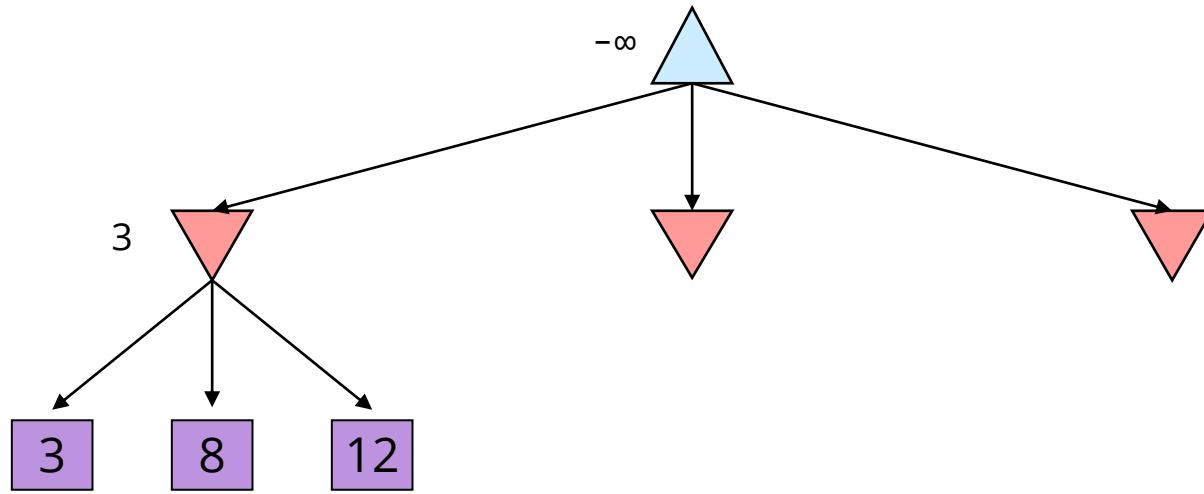
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



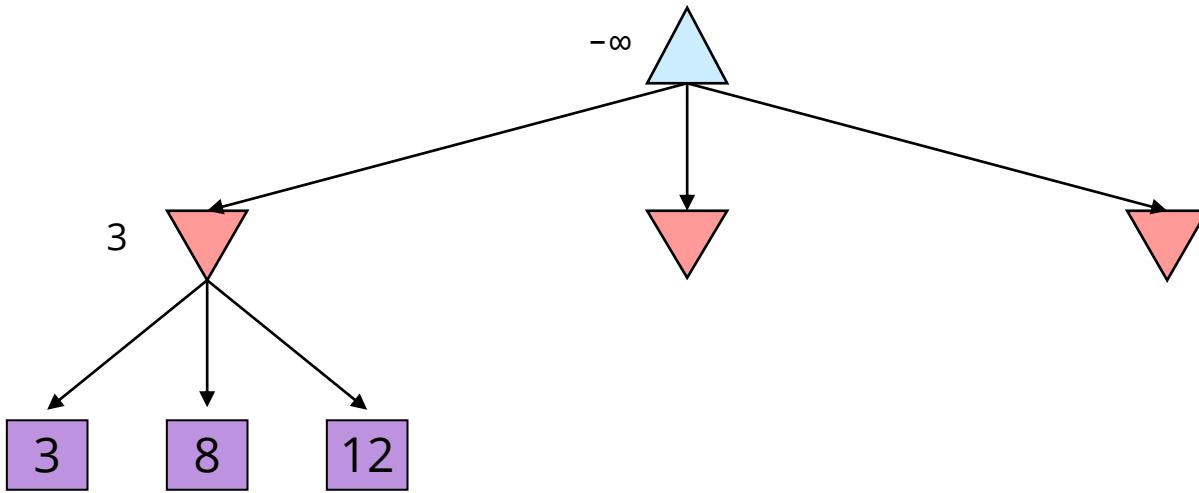
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



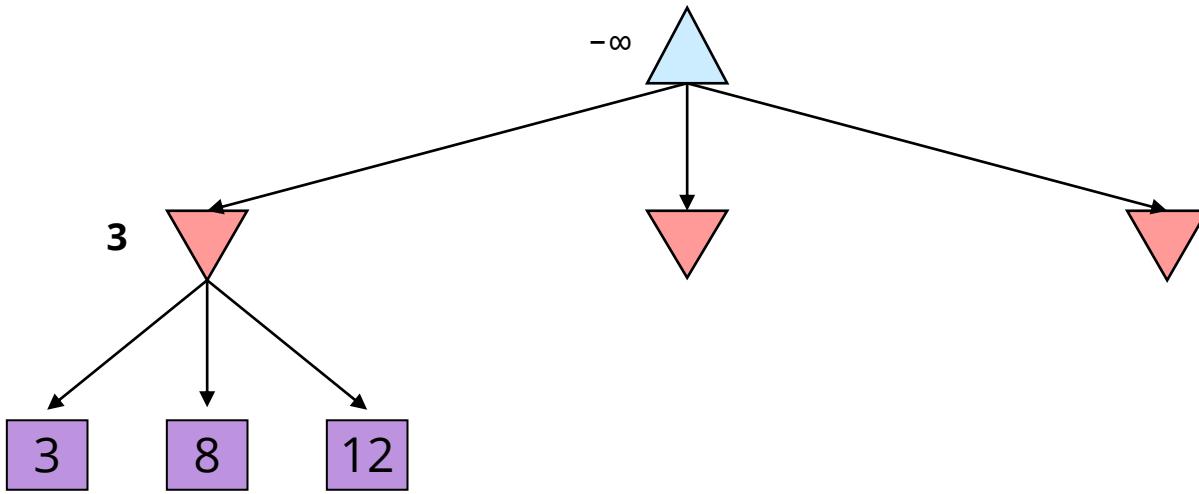
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



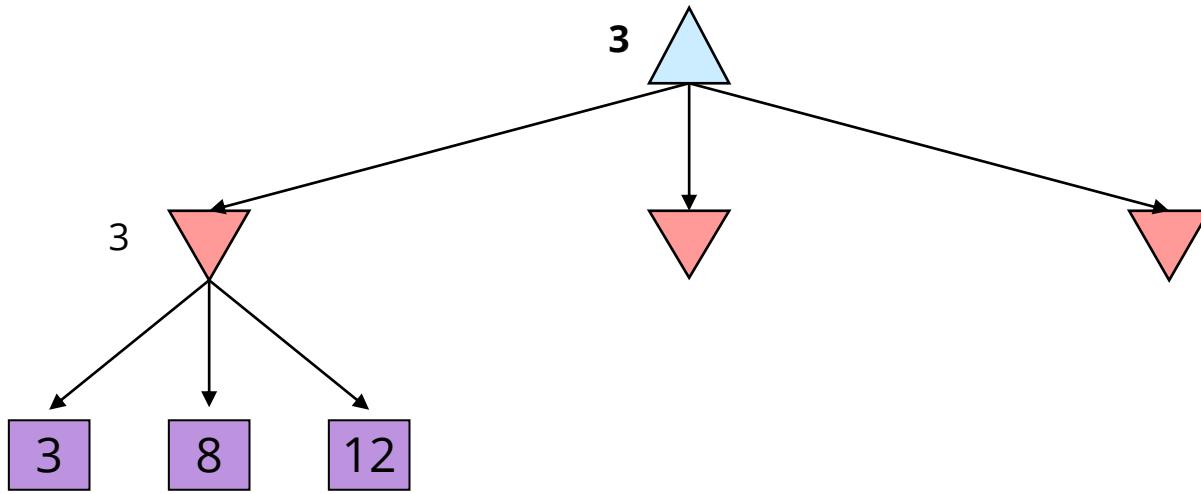
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



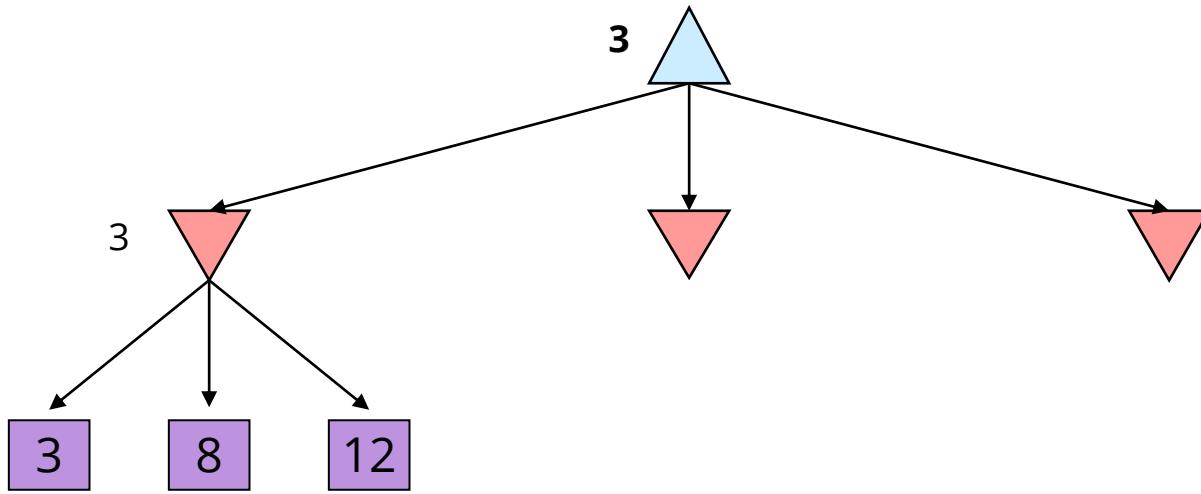
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



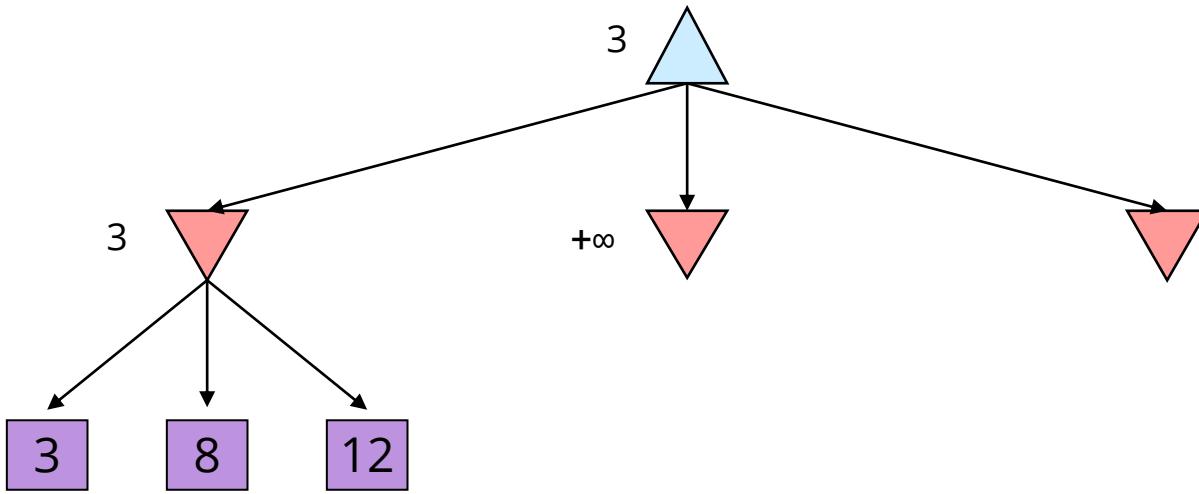
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



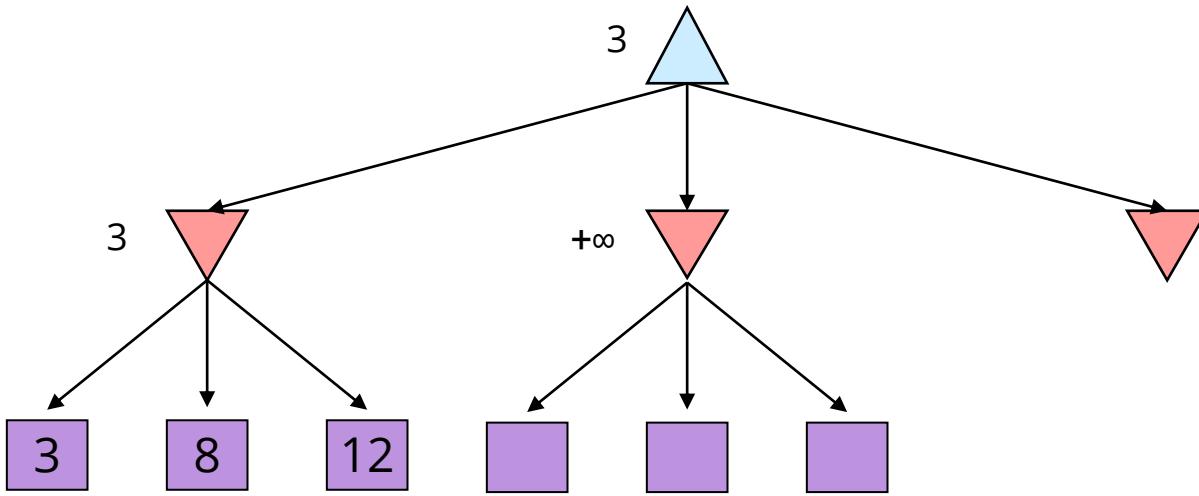
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



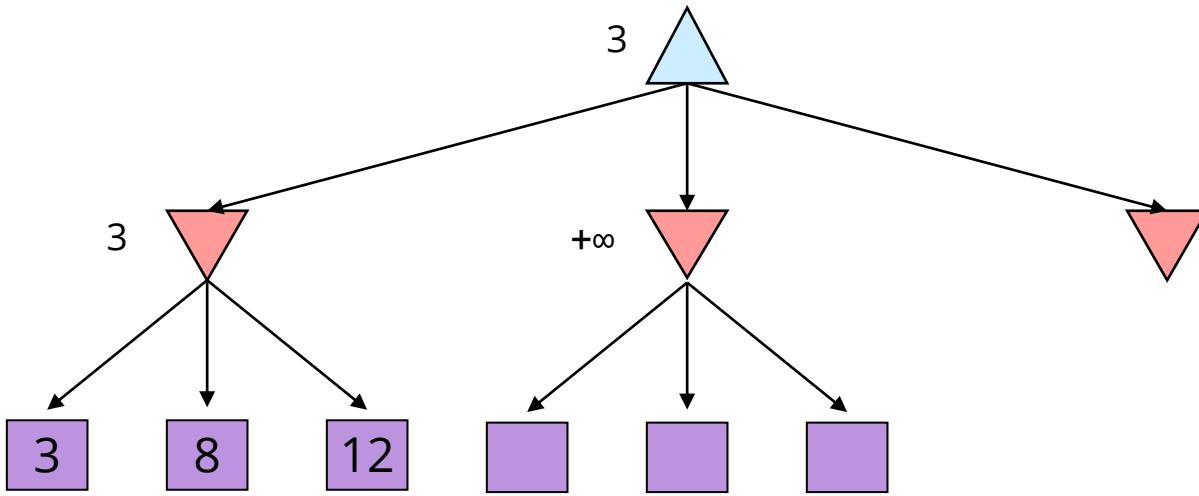
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



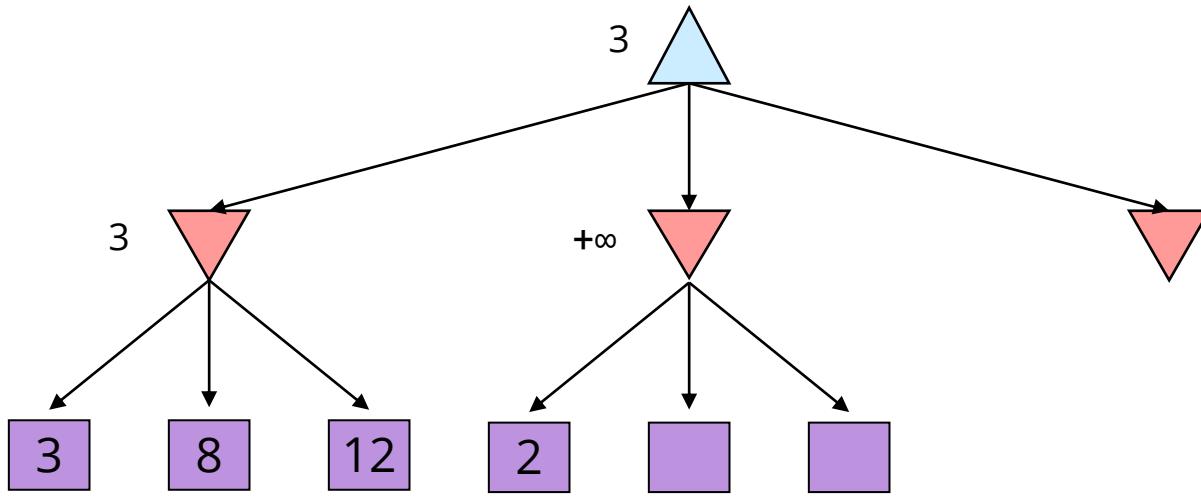
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



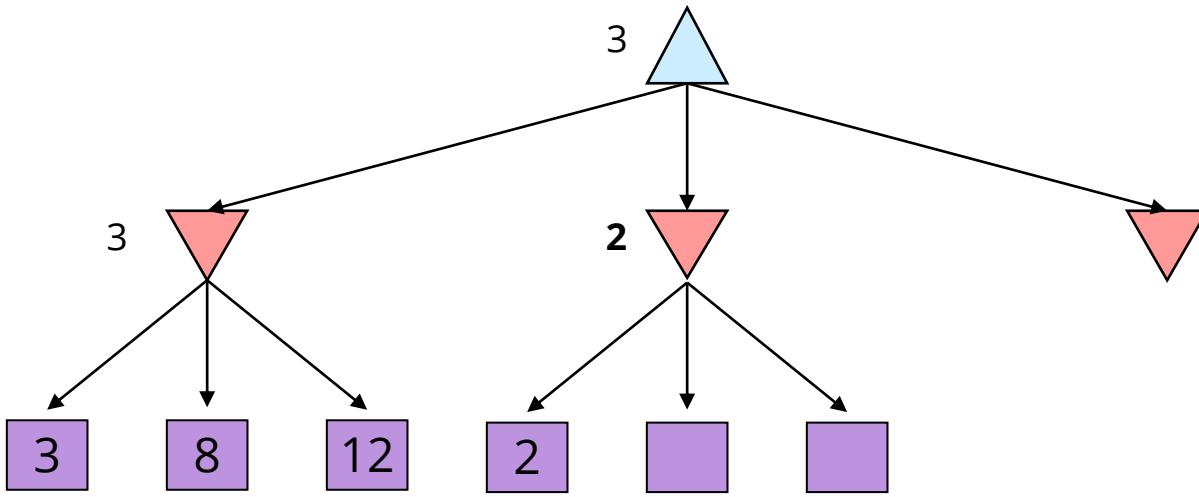
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



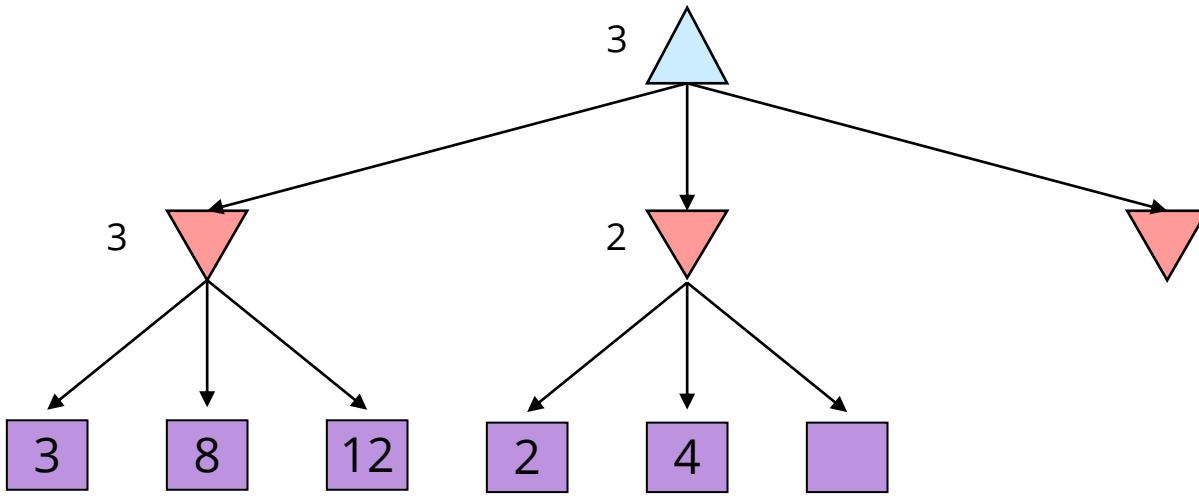
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



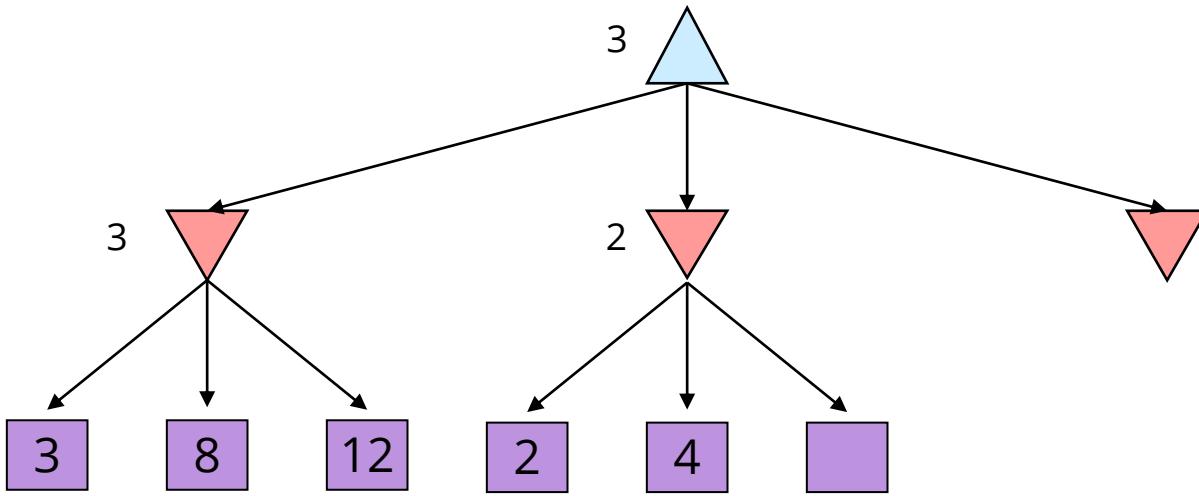
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



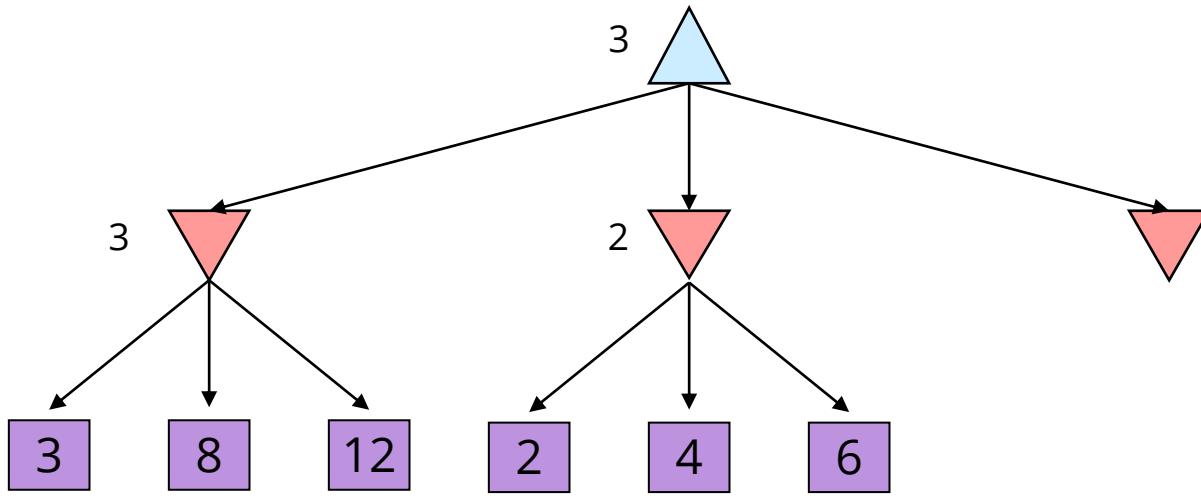
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



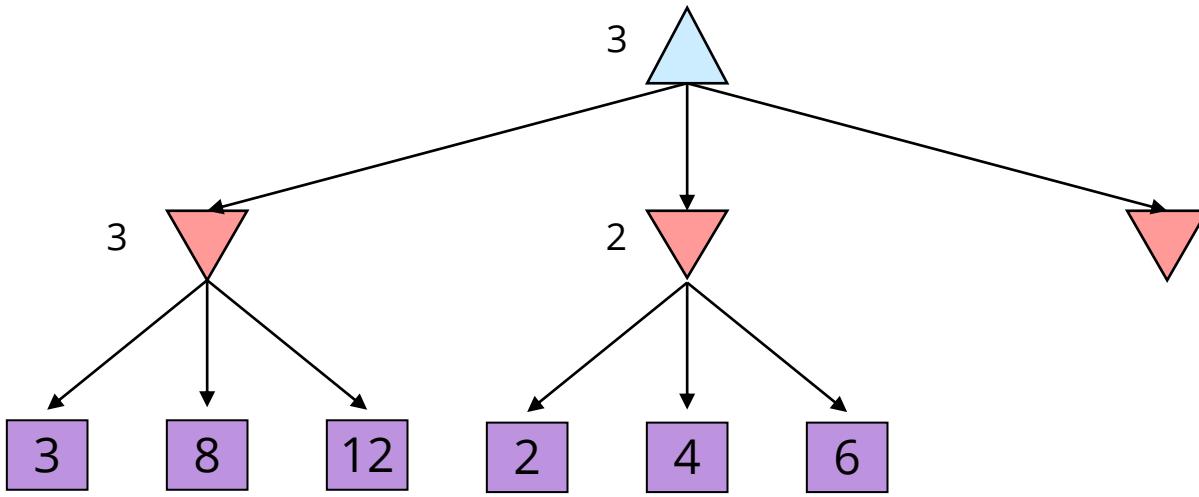
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



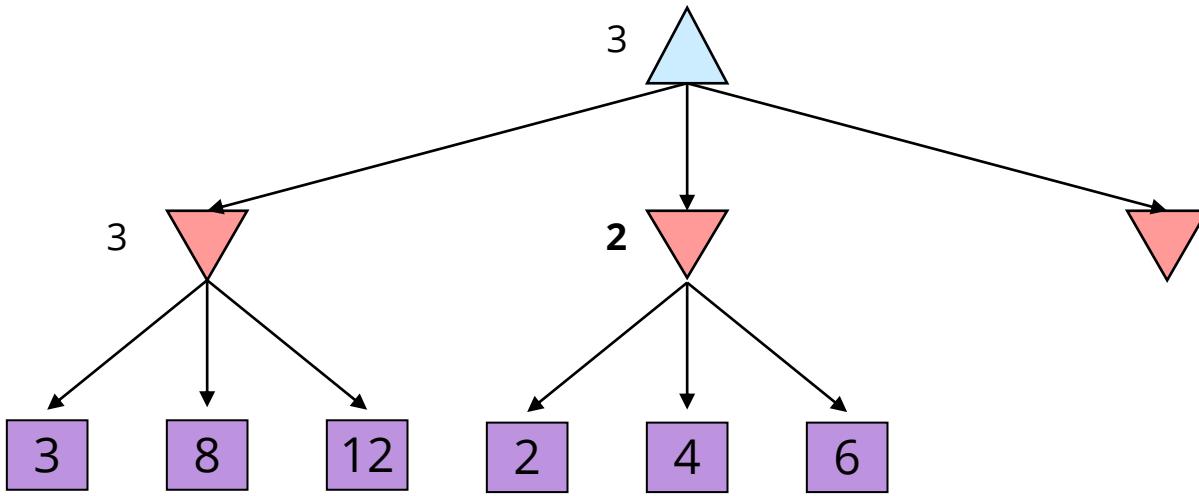
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



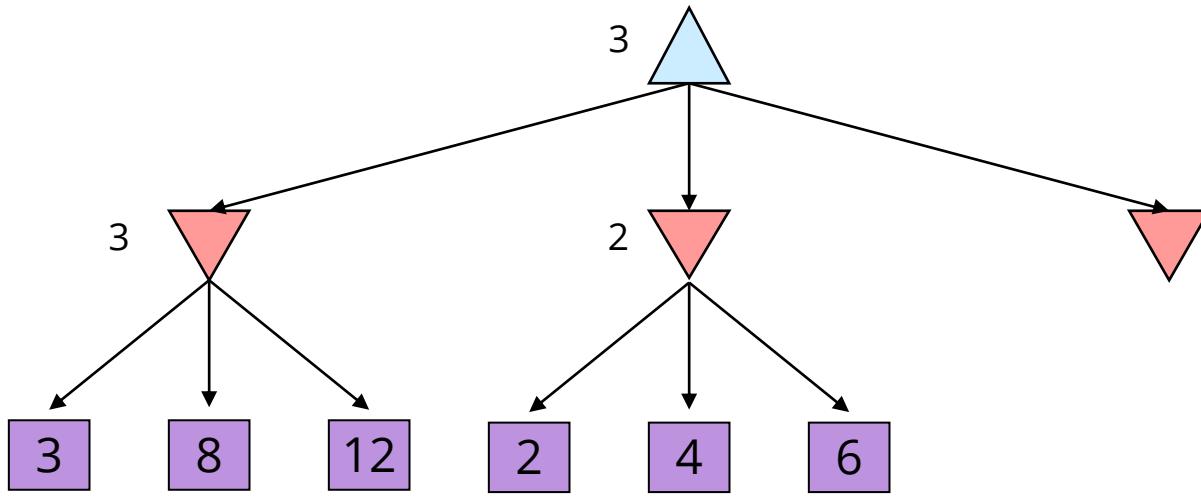
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



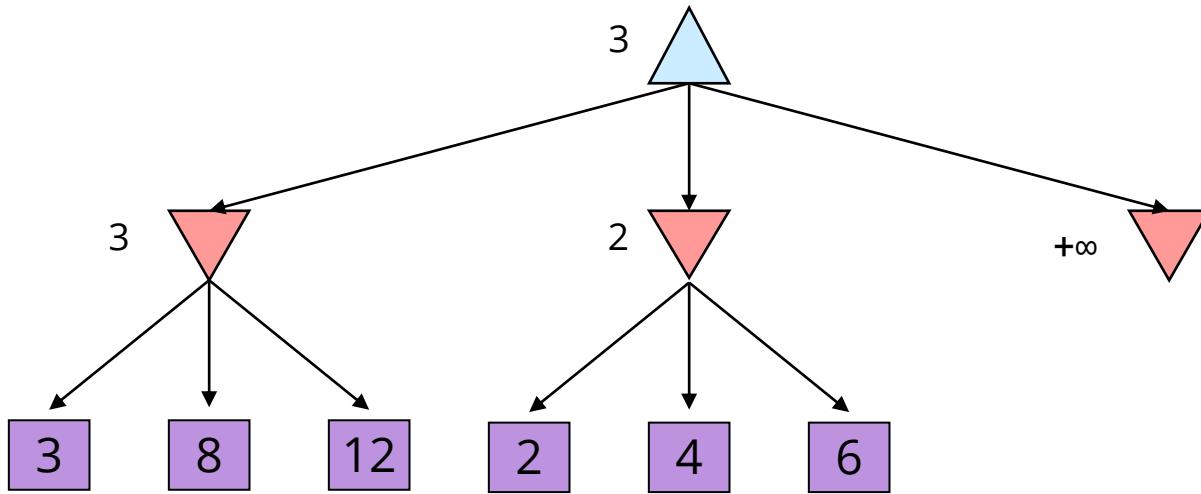
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



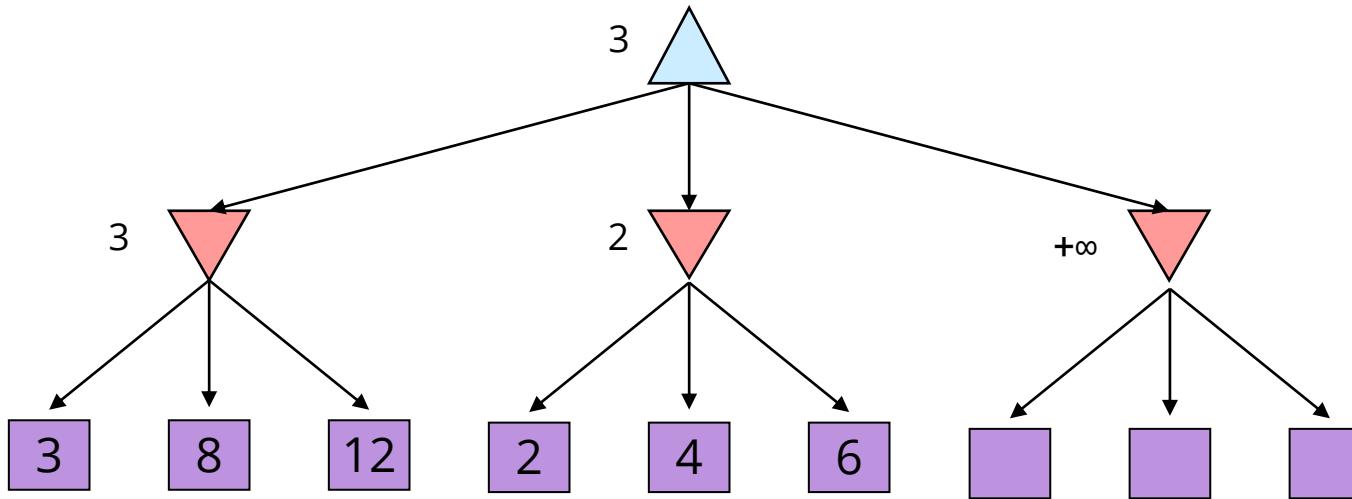
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



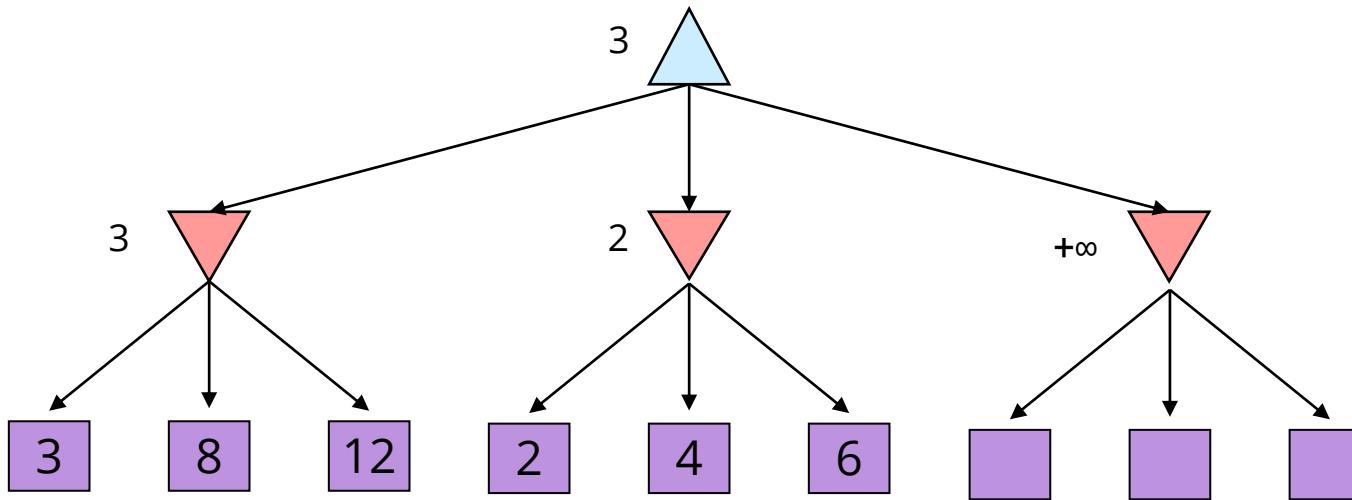
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



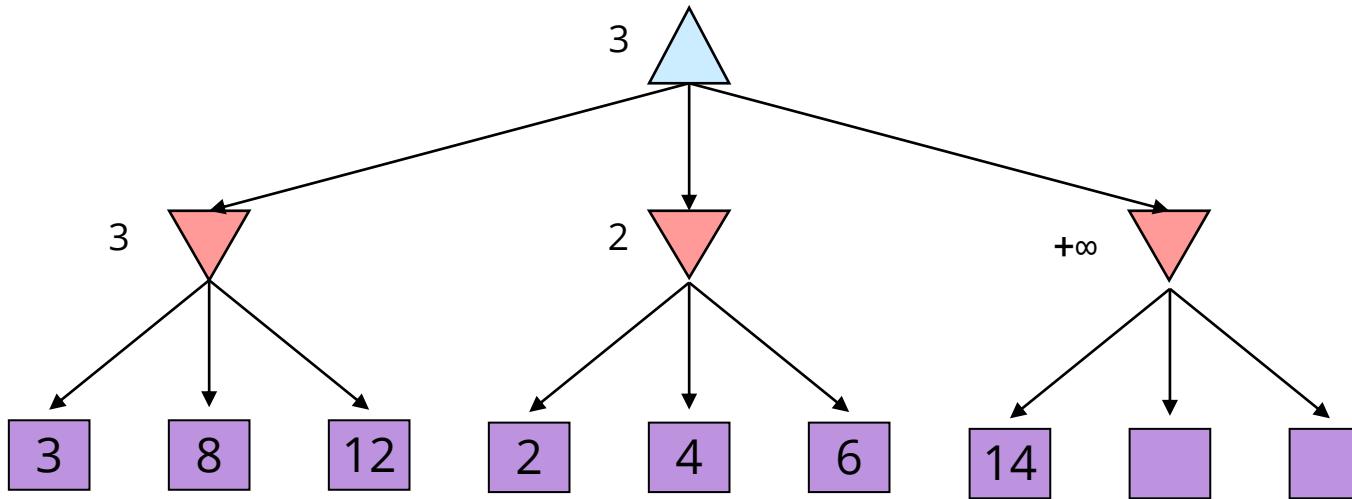
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



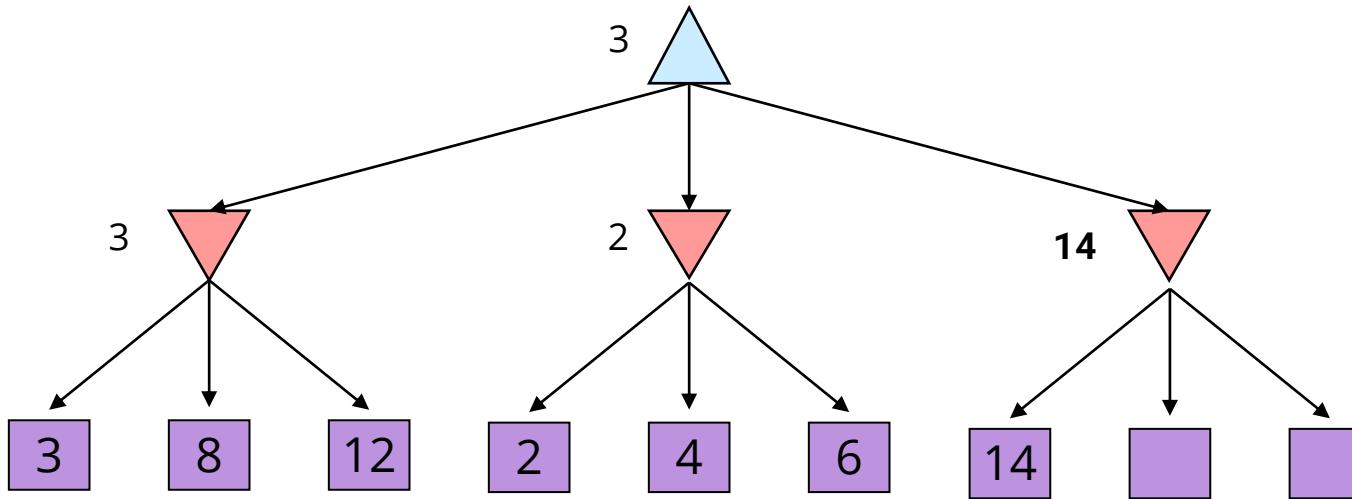
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



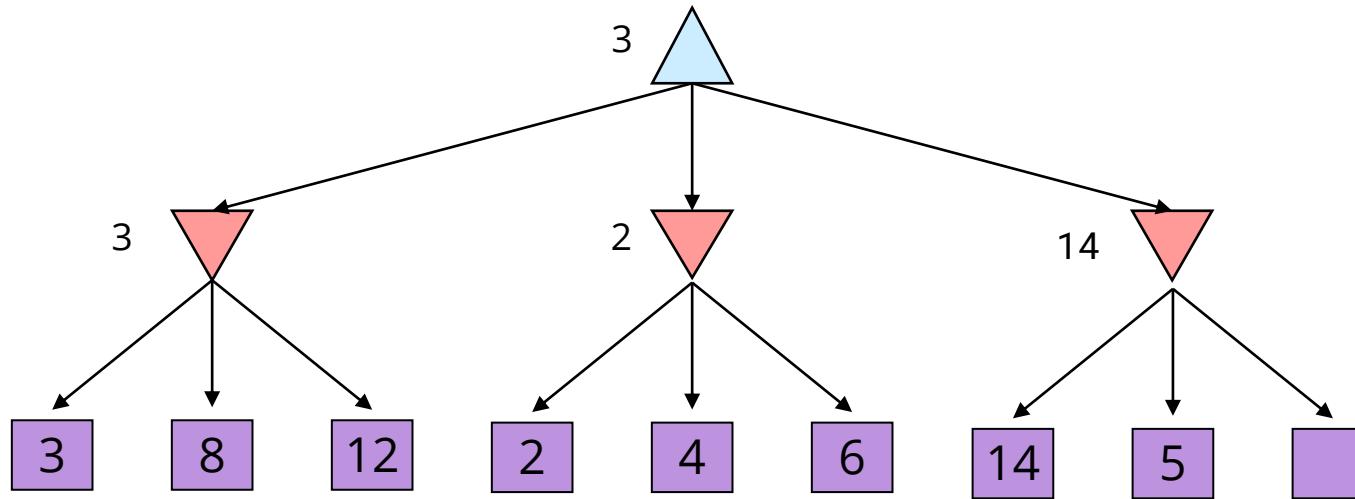
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



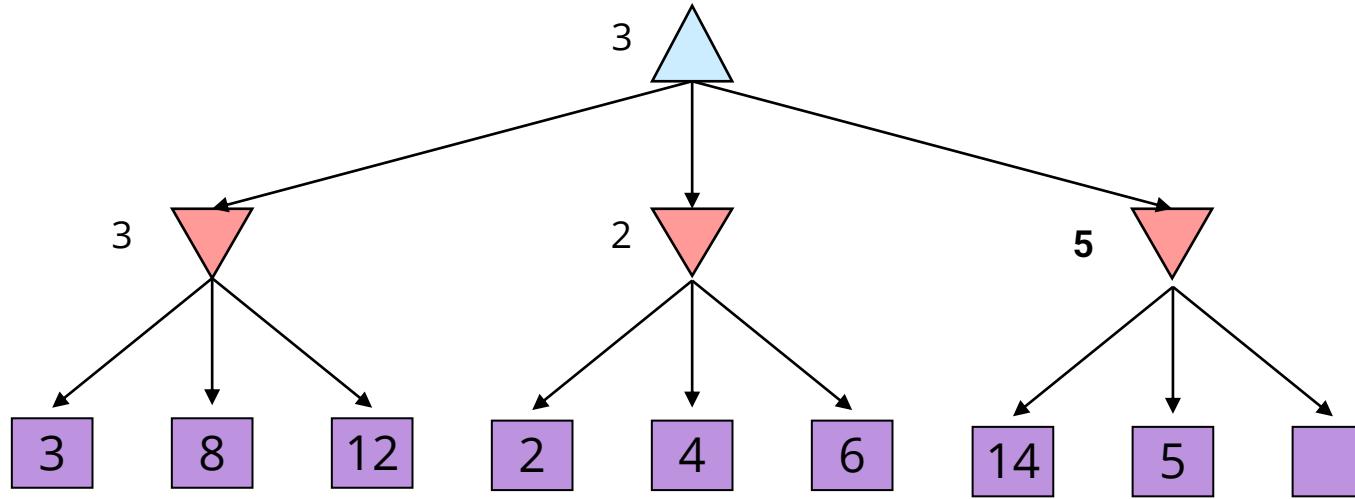
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



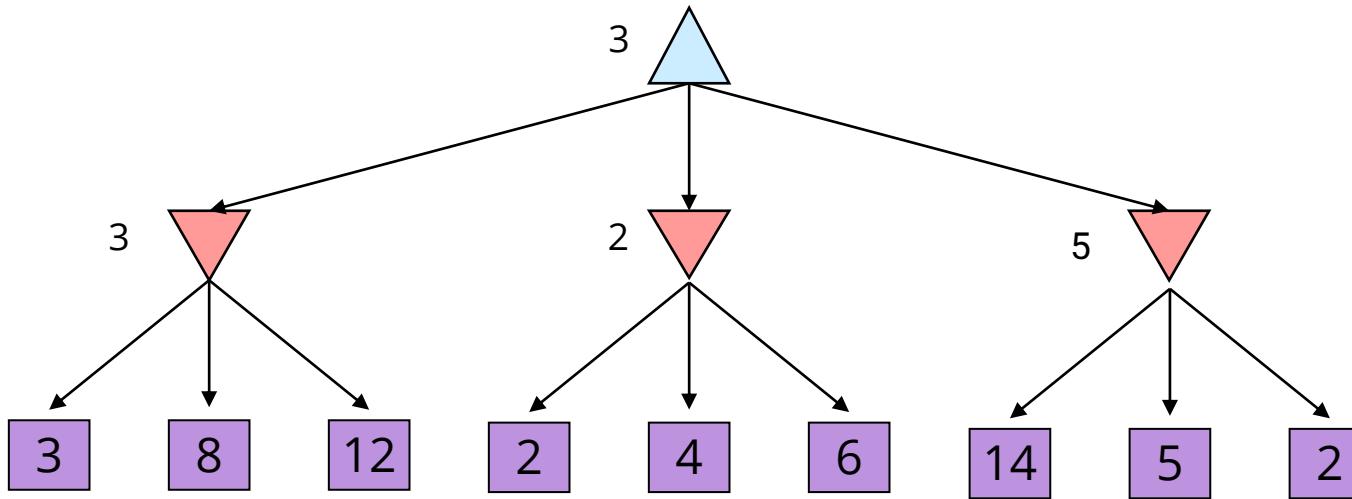
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



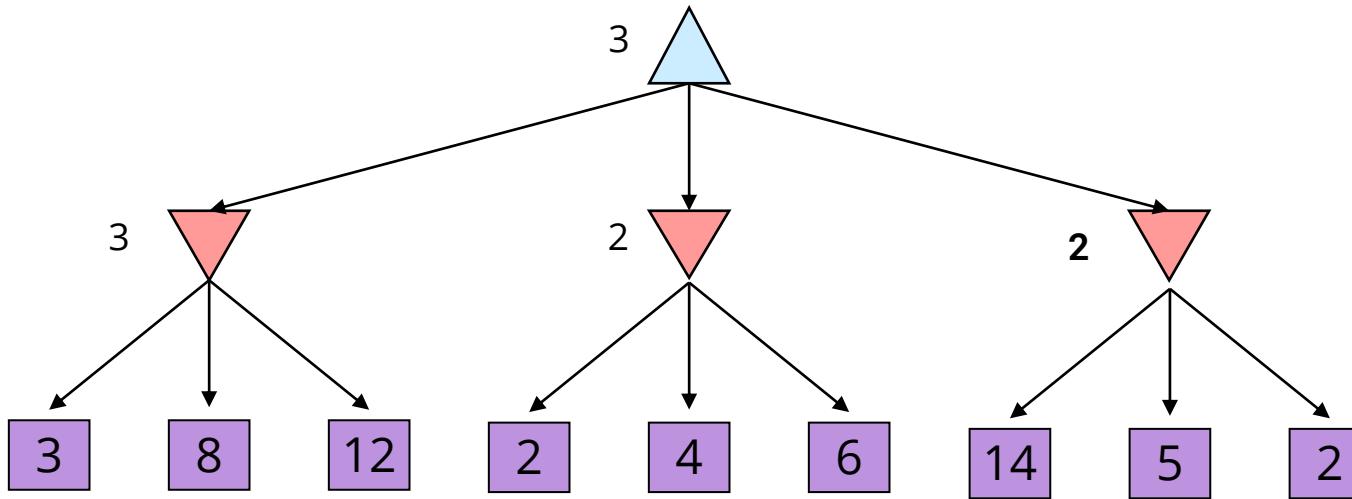
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



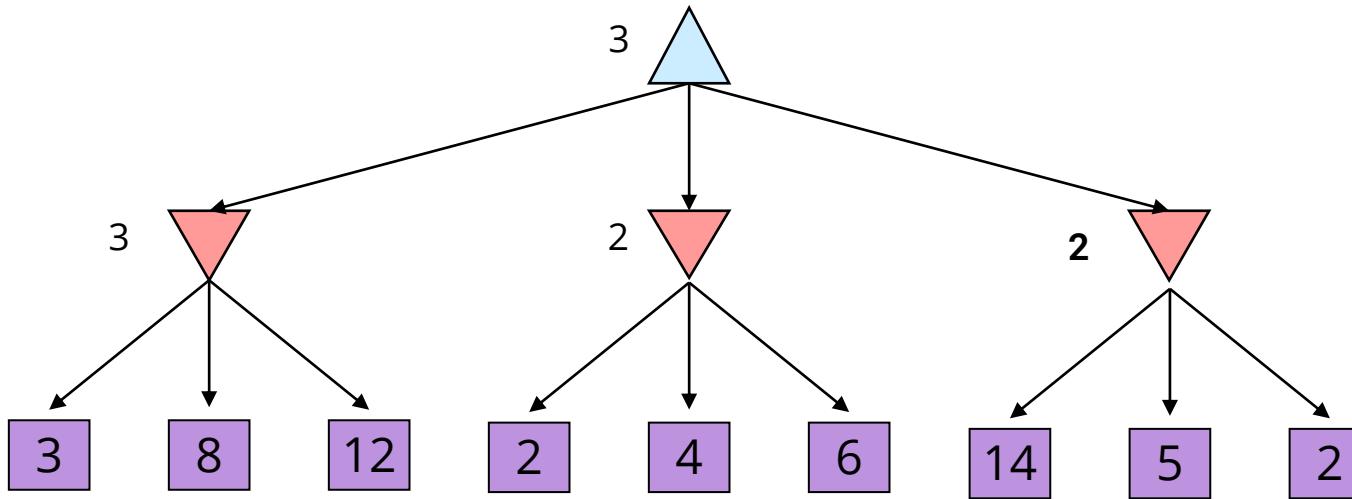
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



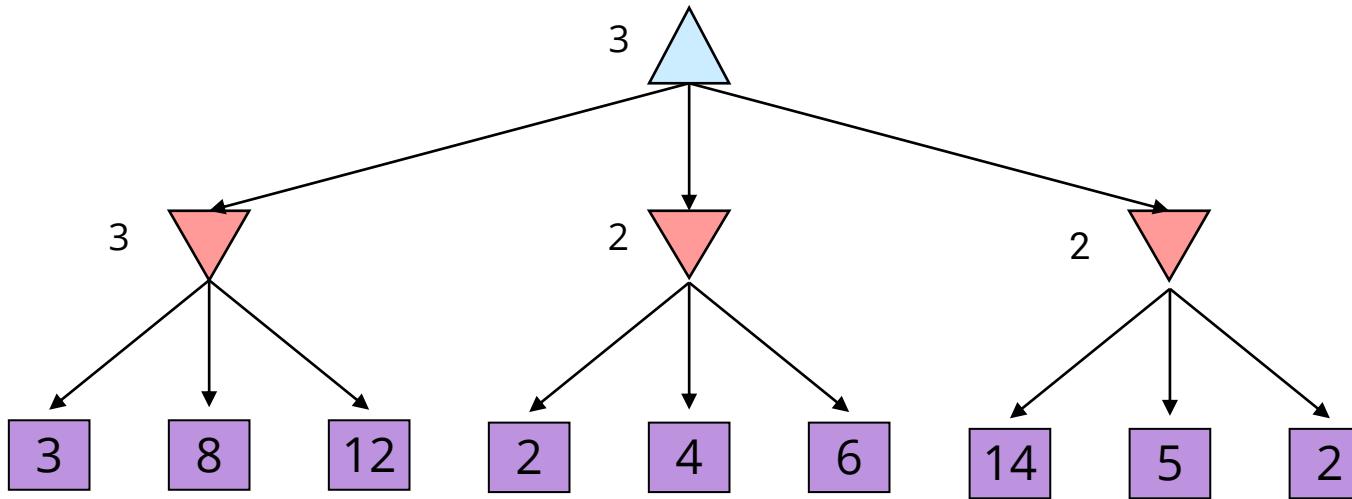
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



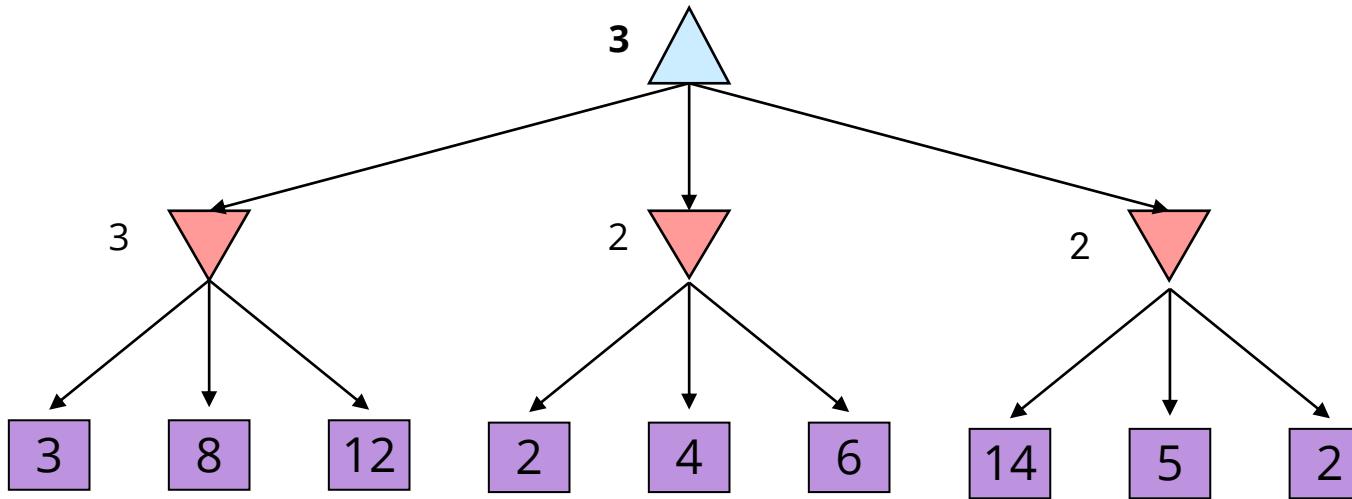
```
def max-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, min-value(successor))  
    return v
```

```
def min-value(state):  
    if the state is a terminal state:  
        return the state's utility  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, max-value(successor))  
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



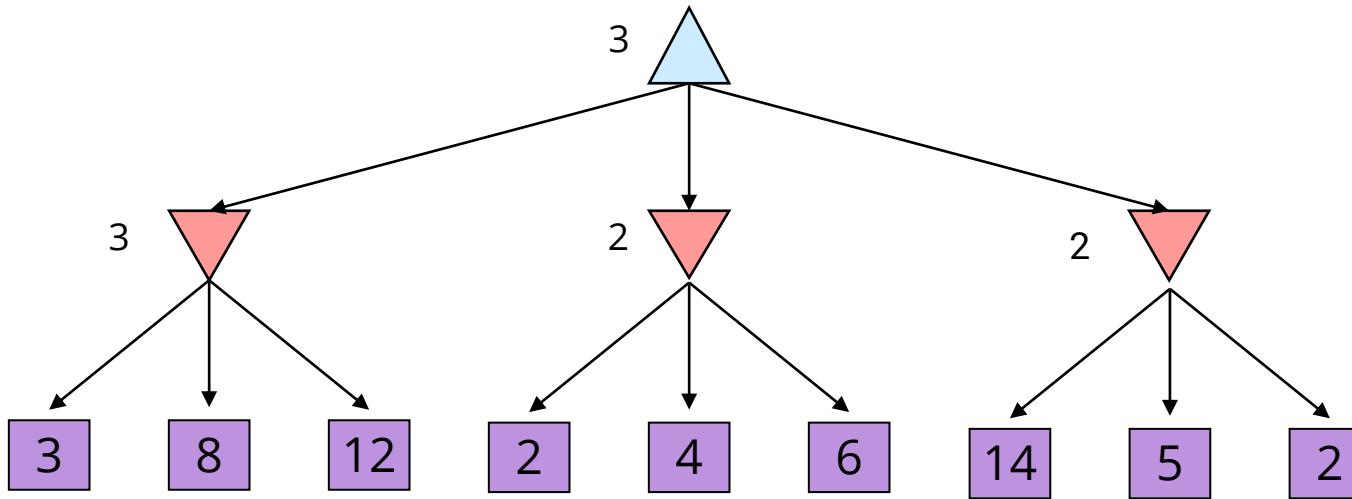
```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Example



```
def max-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    if the state is a terminal state:
        return the state's utility
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

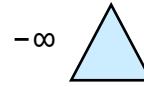
$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Alpha-Beta Pruning

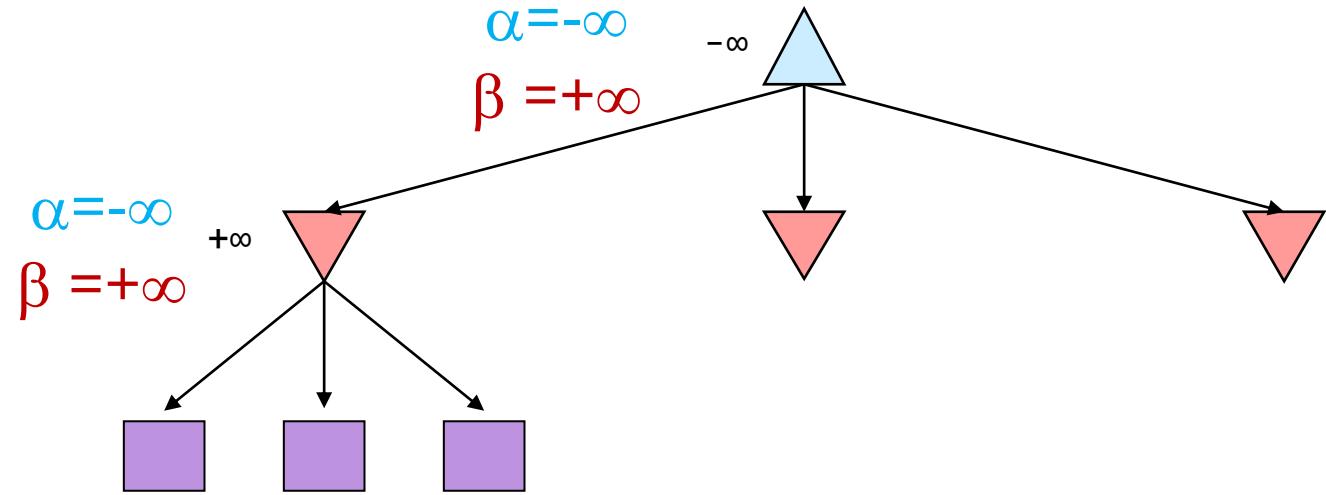
- During Minimax, keep track of two additional values:
 - α : MAX's current *lower* bound on MAX's outcome**
 - β : MIN's current *upper* bound on MIN's outcome**
- MAX will never allow a move that could lead to a worse score (for MAX) than α
- MIN will never allow a move that could lead to a better score (for MAX) than β
- Therefore, stop evaluating a branch whenever:
 - When evaluating a MAX node: a value $v \geq \beta$ is backed-up
 - MIN will never select that MAX node
 - When evaluating a MIN node: a value $v \leq \alpha$ is found
 - MAX will never select that MIN node

Alpha-Beta Pruning Example

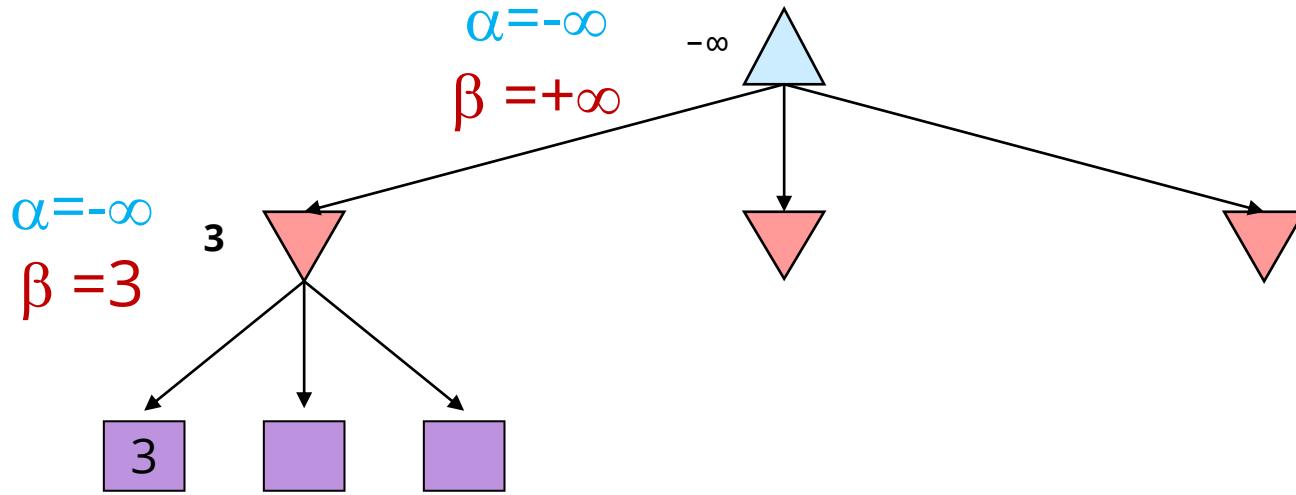
$\alpha = -\infty$
 $\beta = +\infty$



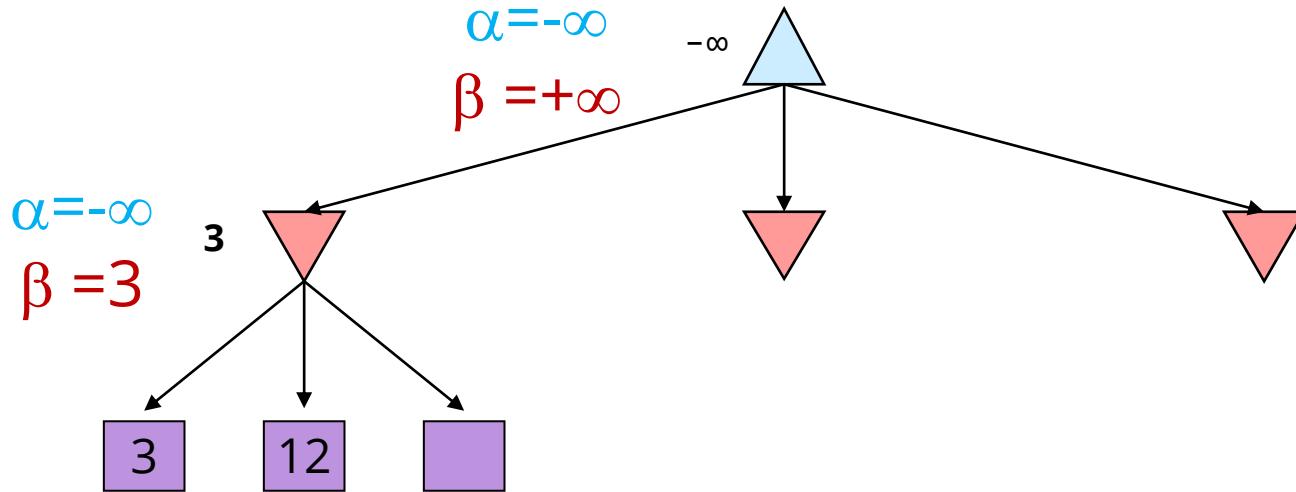
Alpha-Beta Pruning Example



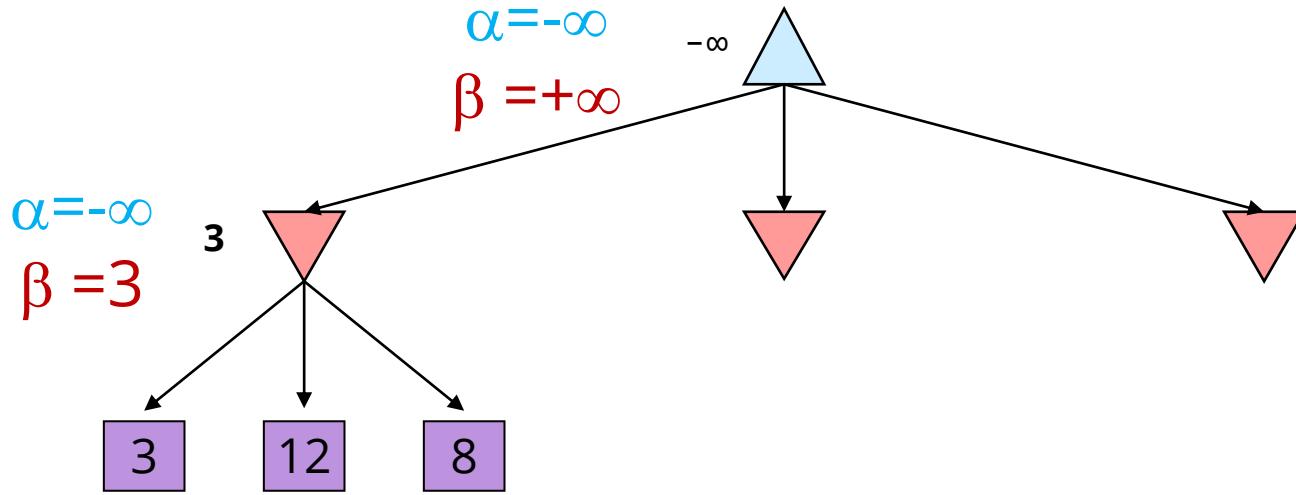
Alpha-Beta Pruning Example



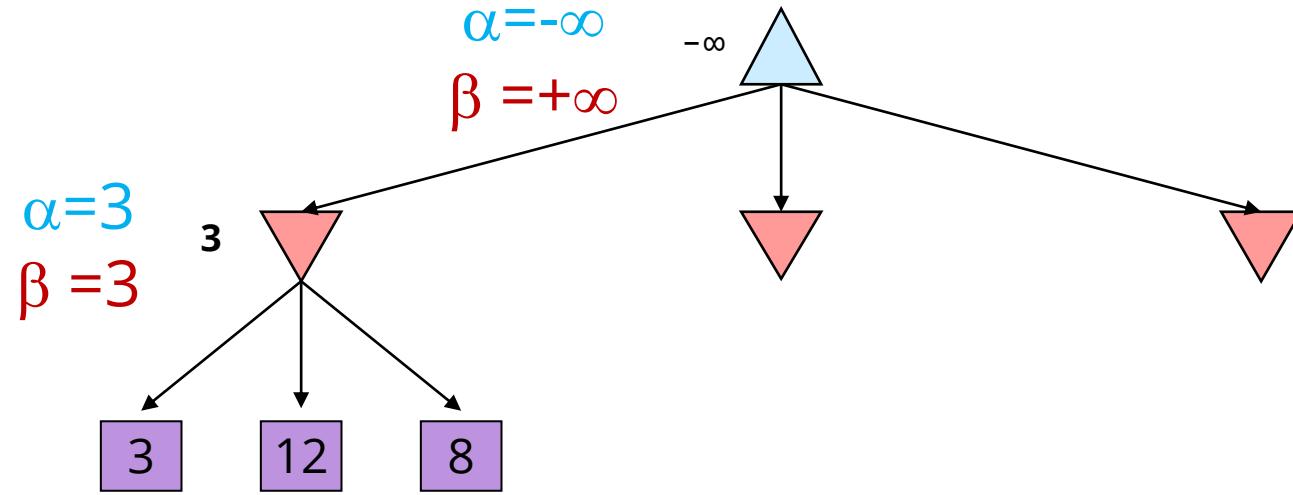
Alpha-Beta Pruning Example



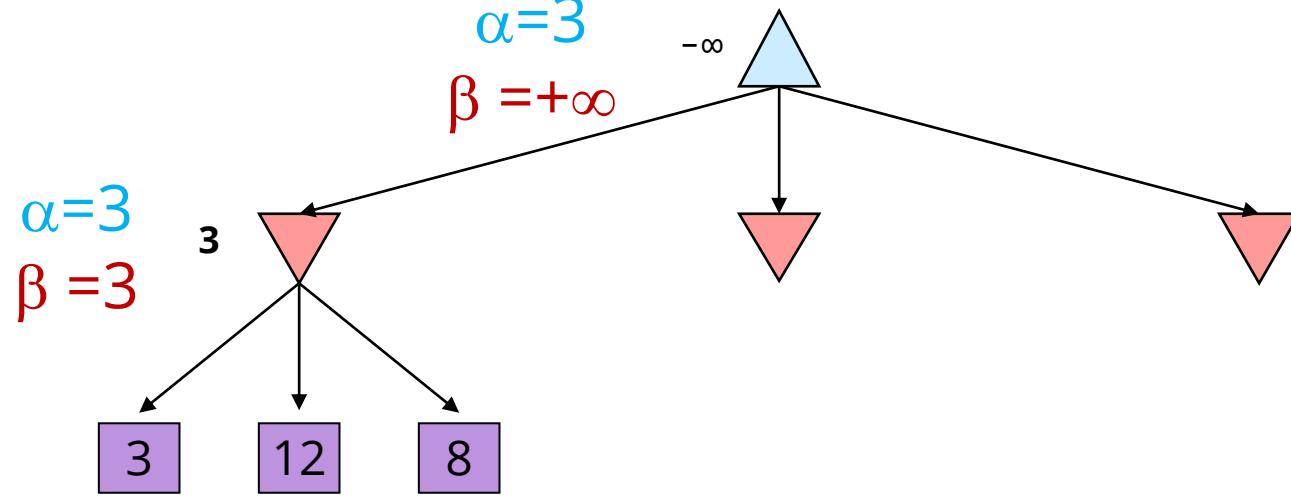
Alpha-Beta Pruning Example



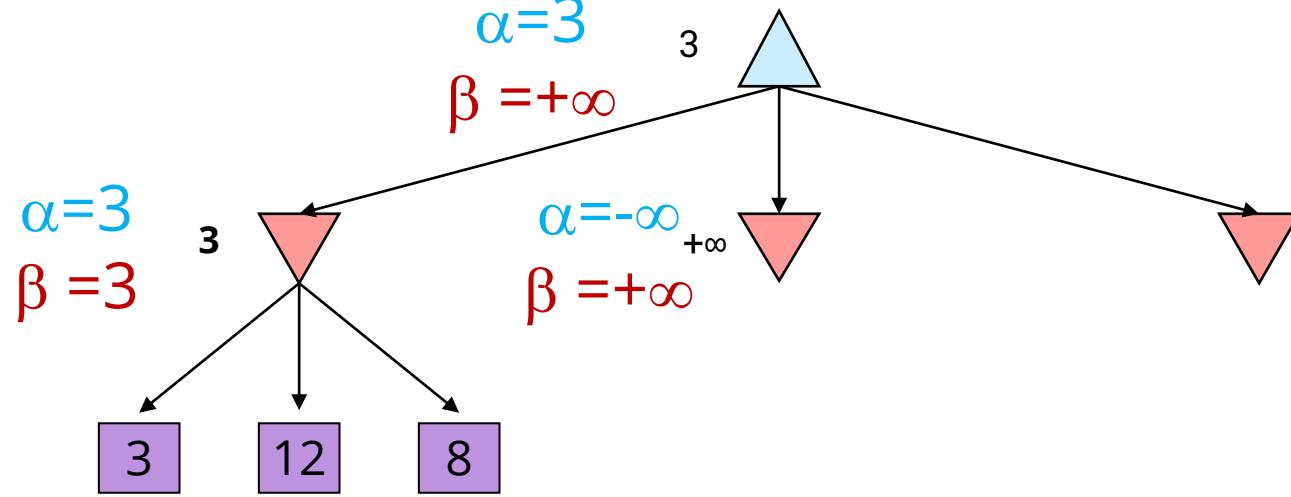
Alpha-Beta Pruning Example



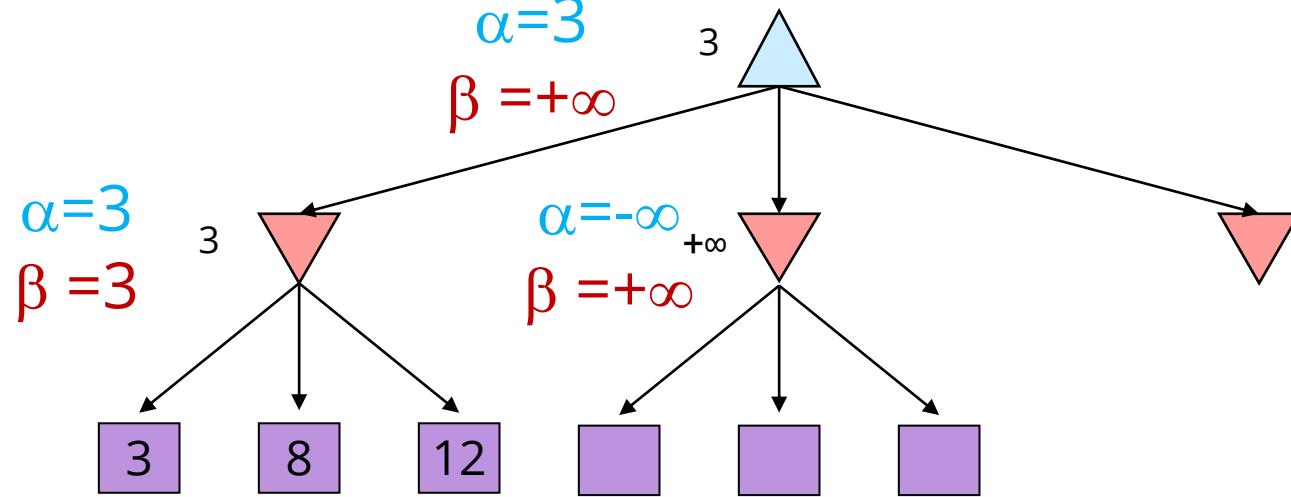
Alpha-Beta Pruning Example



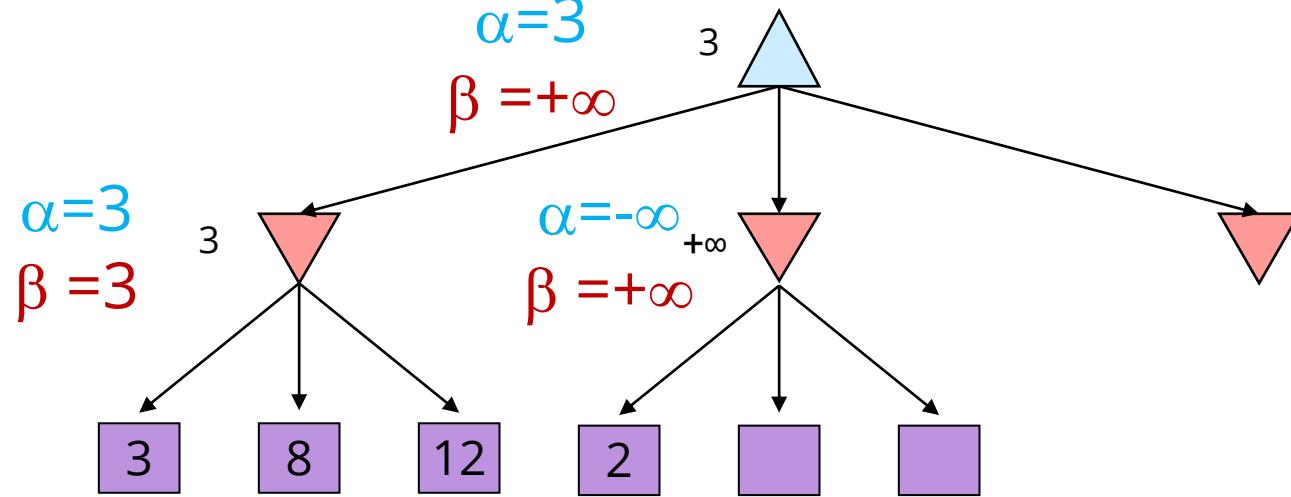
Alpha-Beta Pruning Example



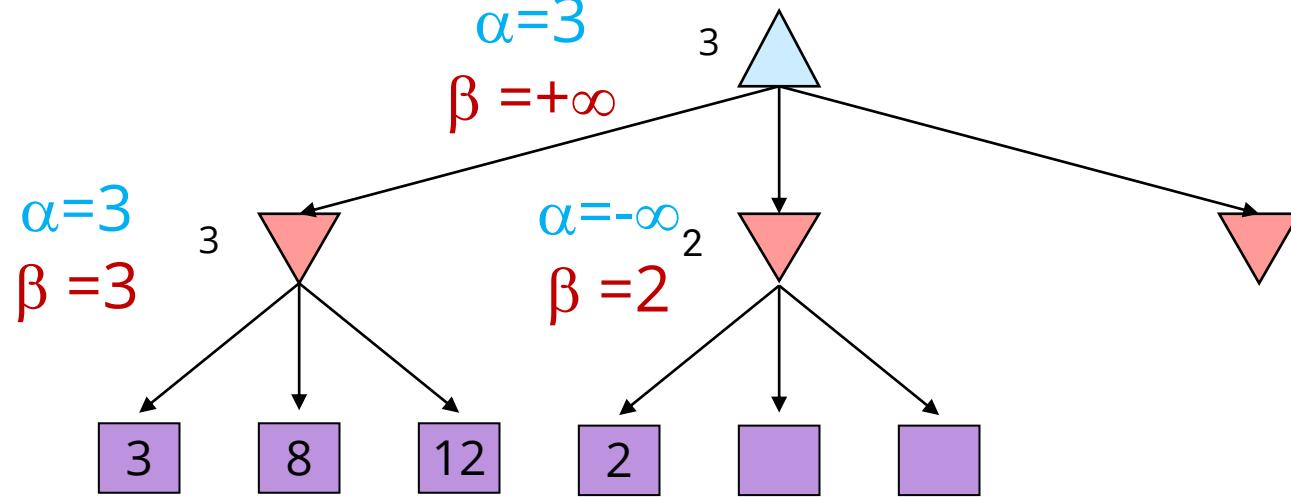
Alpha-Beta Pruning Example



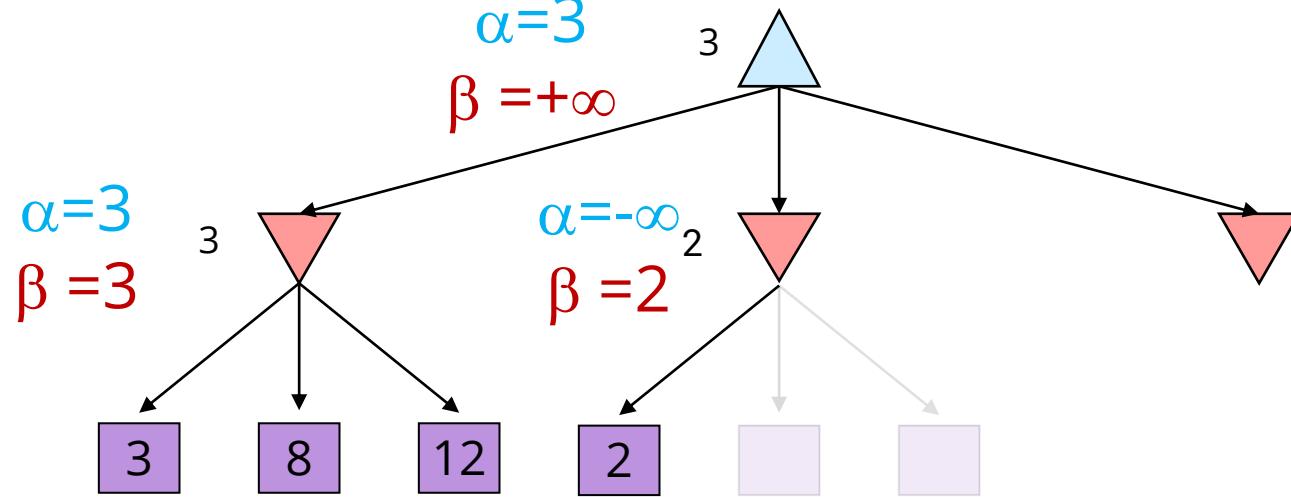
Alpha-Beta Pruning Example



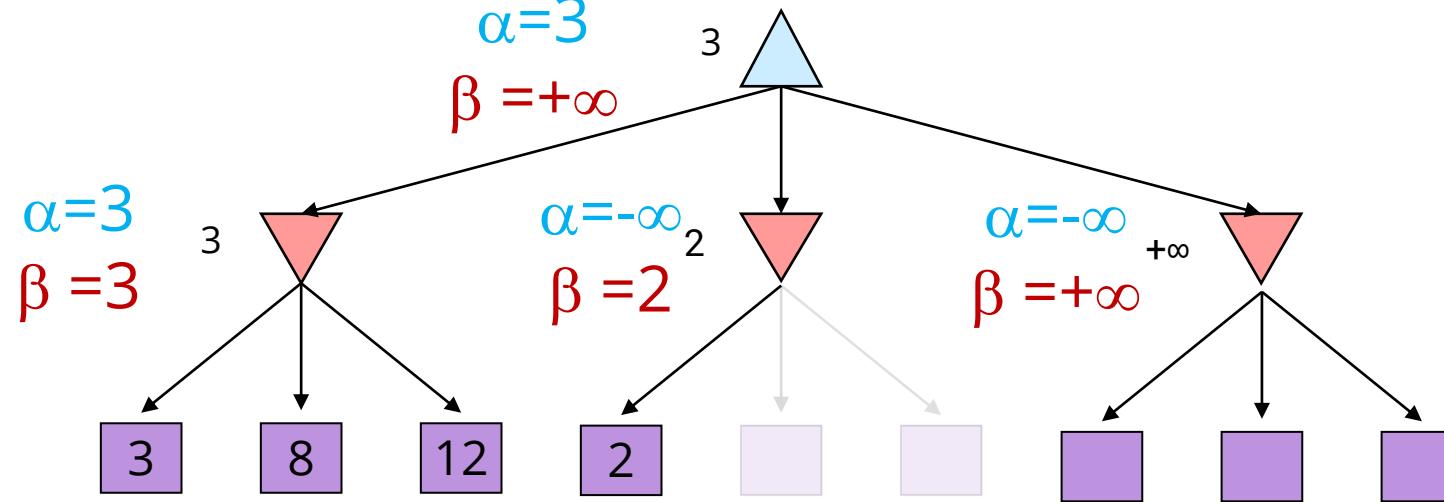
Alpha-Beta Pruning Example



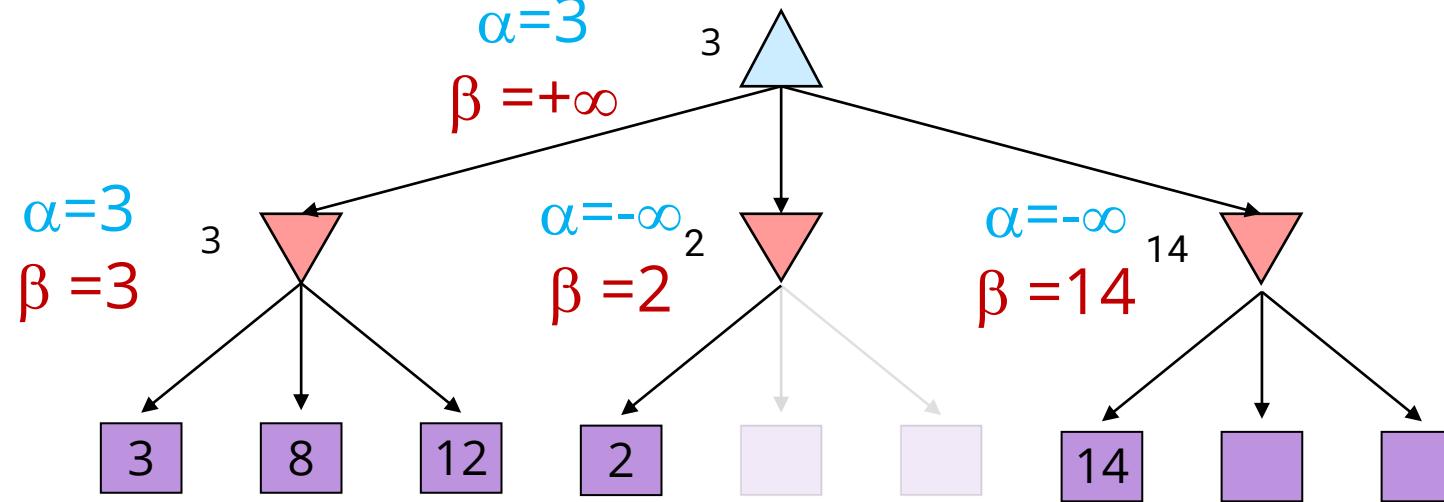
Alpha-Beta Pruning Example



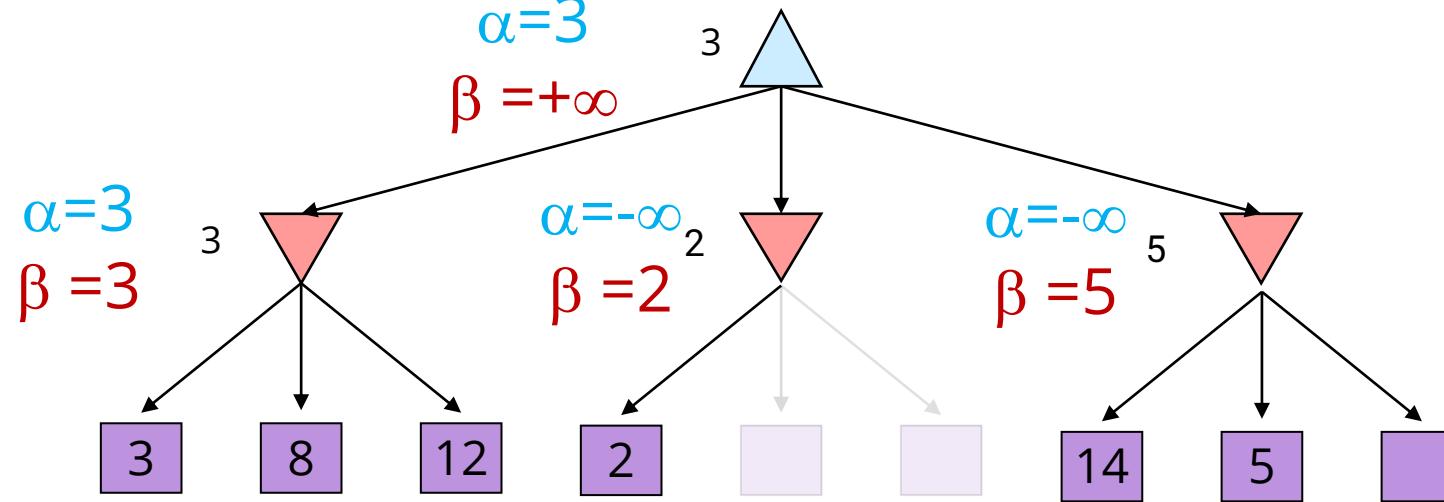
Alpha-Beta Pruning Example



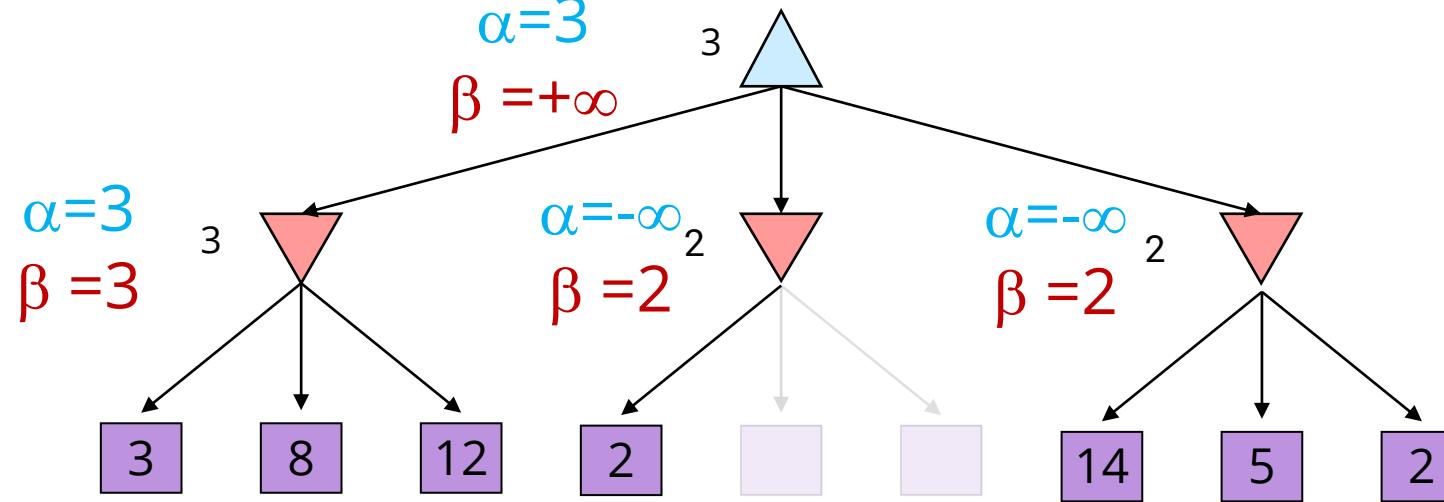
Alpha-Beta Pruning Example



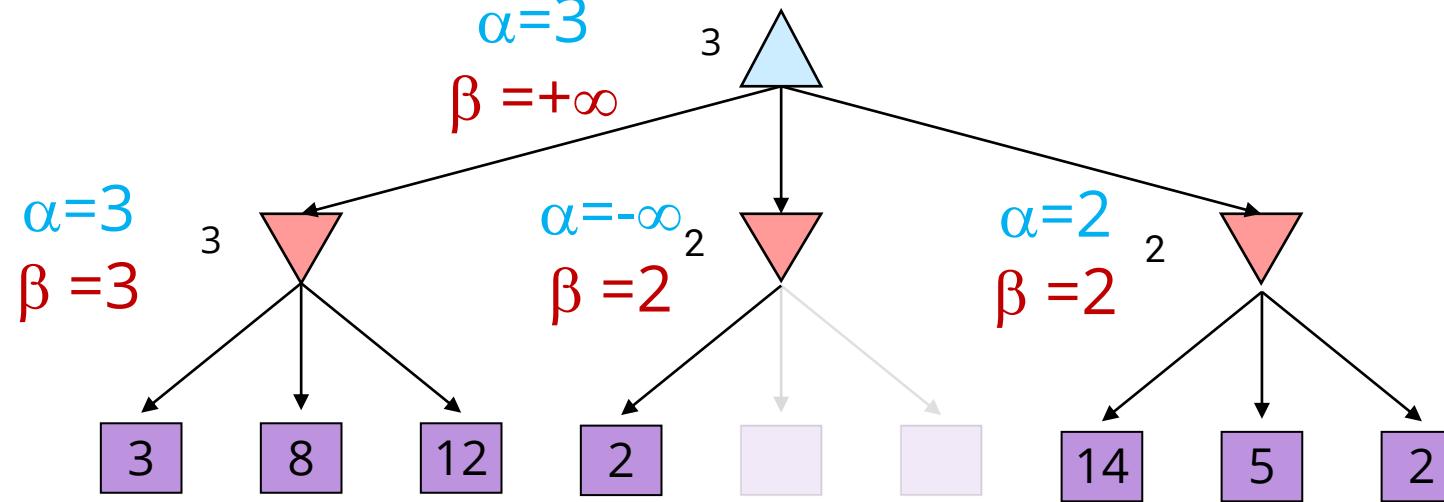
Alpha-Beta Pruning Example



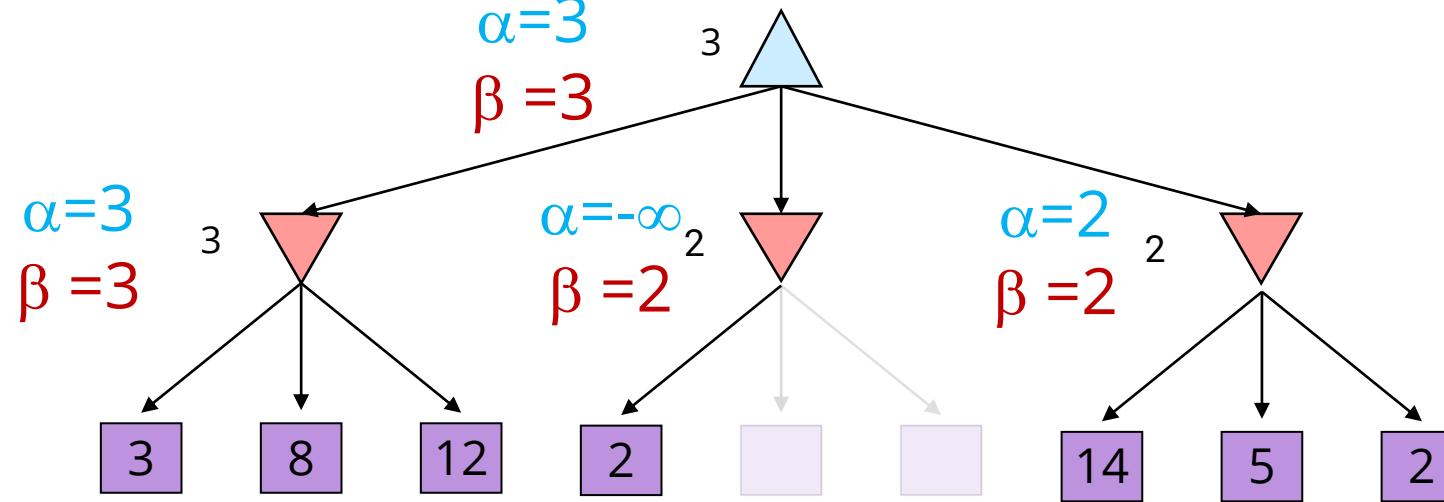
Alpha-Beta Pruning Example



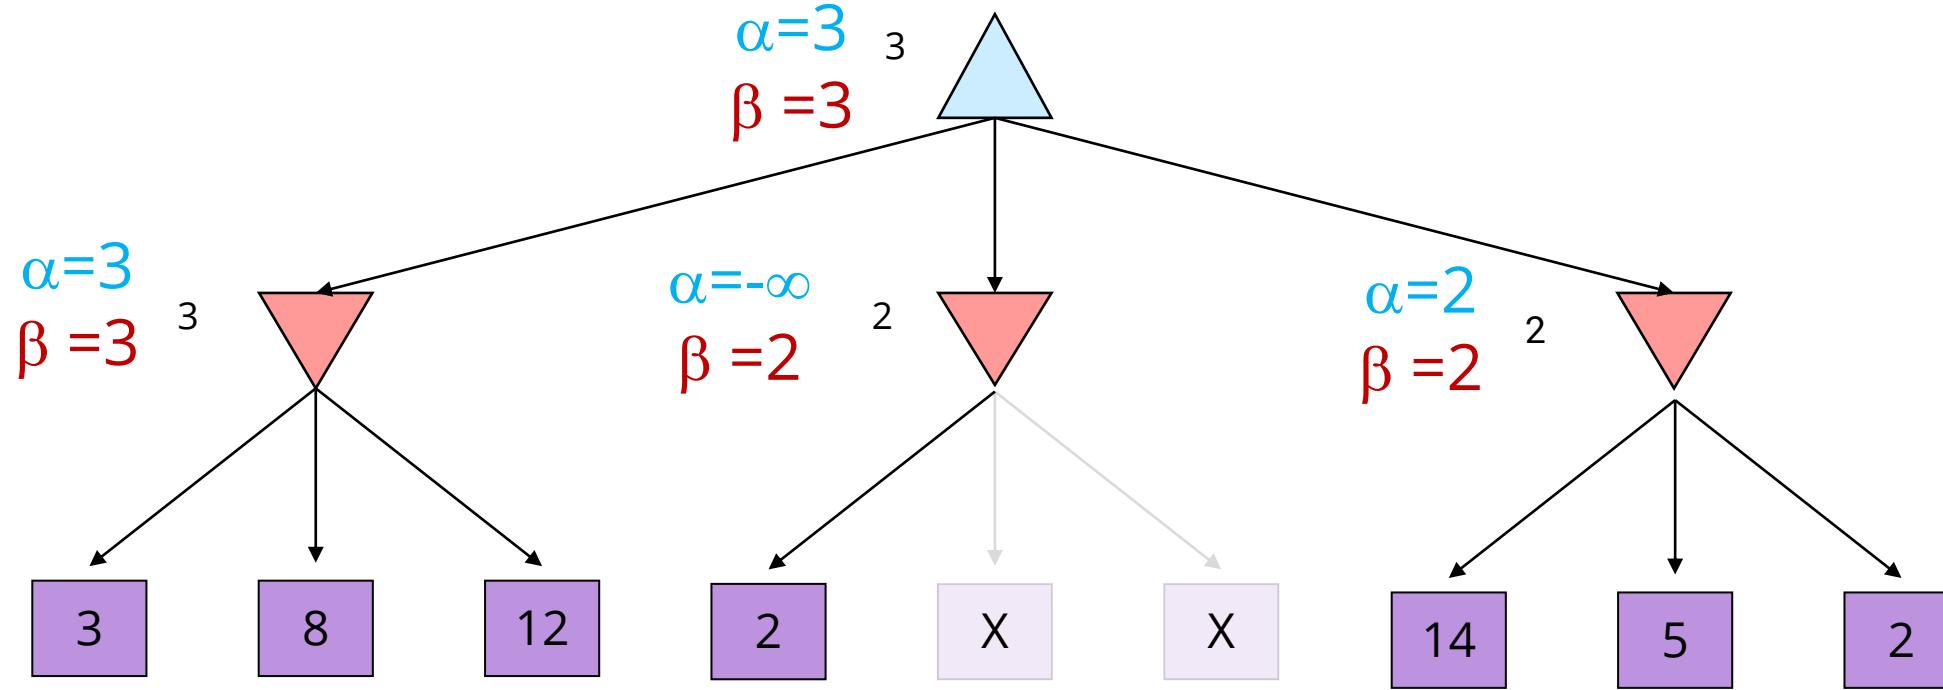
Alpha-Beta Pruning Example



Alpha-Beta Pruning Example



Alpha-Beta Pruning



α : MAX's current lower bound on MAX's outcome

β : MIN's current upper bound on MIN's outcome

α : MAX's best option on path to root

β : MIN's best option on path to root

Review: Evaluation functions

- Evaluates how good a ‘board position’ is
Based on *static features* of that board alone
- Zero-sum assumption lets us use one function to describe goodness for both players.
 - $f(n) > 0$ if MAX is winning in position n
 - $f(n) = 0$ if position n is tied
 - $f(n) < 0$ if MIN is winning in position n
- Build using expert knowledge,
Tic-tac-toe: $f(n) = (\# \text{ of 3 lengths open for MAX}) - (\# \text{ open for MIN})$

(AIMA 5..1)

Review: Chess Evaluation Functions

- Chess needs an evaluation function since it is impossible to search the game tree deeply enough to reach the terminal nodes
- $f(n) = (\text{sum of } A\text{'s piece values}) - (\text{sum of } B\text{'s piece values})$
- More complex: weighted sum of positional features:
$$\sum w_i \cdot \text{feature}_i(n)$$
- $f(n)$ can be a **weighted linear function**

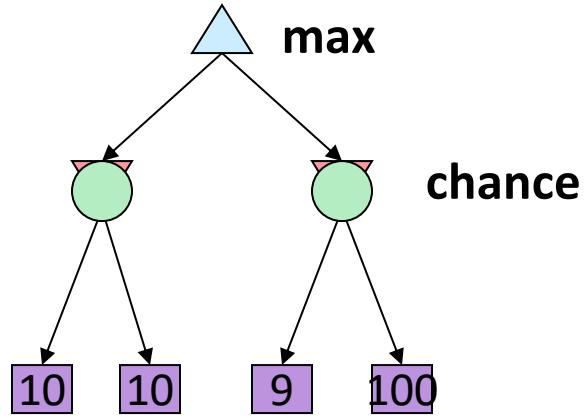
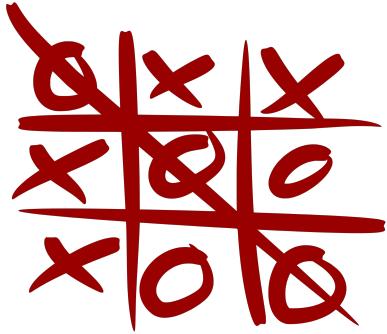
Pawn	1.0
Knight	3.0
Bishop	3.25
Rook	5.0
Queen	9.0

Pieces values for a simple evaluation function often taught to novice chess players

Uncertain Outcomes



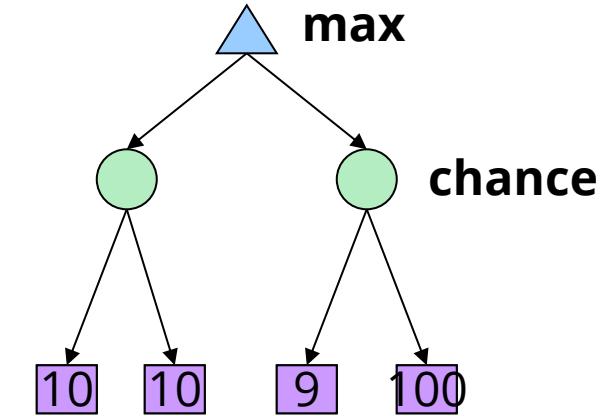
Title



Idea: Uncertain outcomes controlled by chance, not an adversary!

Expectimax Search

- Why wouldn't we know what the result of an action will be?
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the opponent isn't optimal
 - Actions can fail: when moving a robot, wheels might slip
- Values should now reflect average-case (expectimax) outcomes, not worst-case (minimax) outcomes
- **Expectimax search:** compute the average score under optimal play
 - Max nodes as in minimax search
 - Chance nodes are like min nodes but the outcome is uncertain
 - Calculate their **expected utilities**
 - i.e. take weighted average (expectation) of children
- Later, we'll learn how to formalize the underlying uncertain-result problems as **Markov Decision Processes**

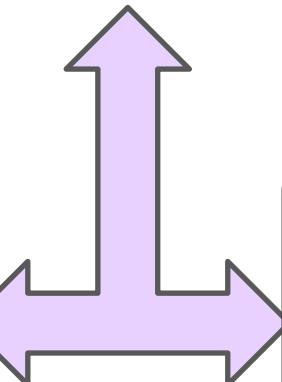


Expectimax Pseudocode

```
def value(state):
    if the state is a terminal state: return the state's utility
    if the next agent is MAX: return max-value(state)
    if the next agent is EXP: return exp-value(state)
```

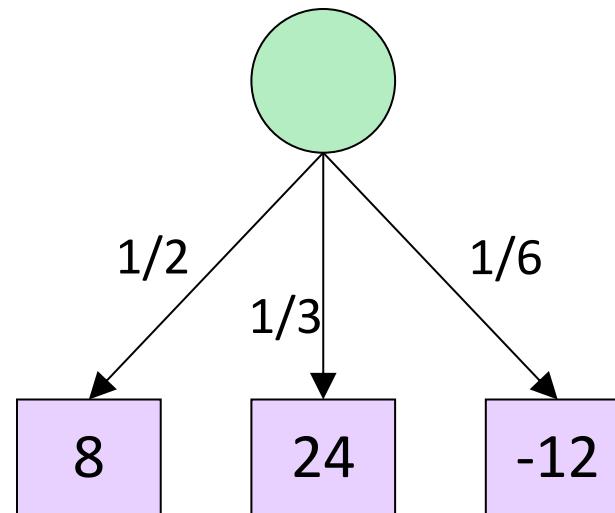
```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

```
def exp-value(state):
    initialize v = 0
    for each successor of state:
        p = probability(successor)
        v += p *
            value(successor)
    return v
```



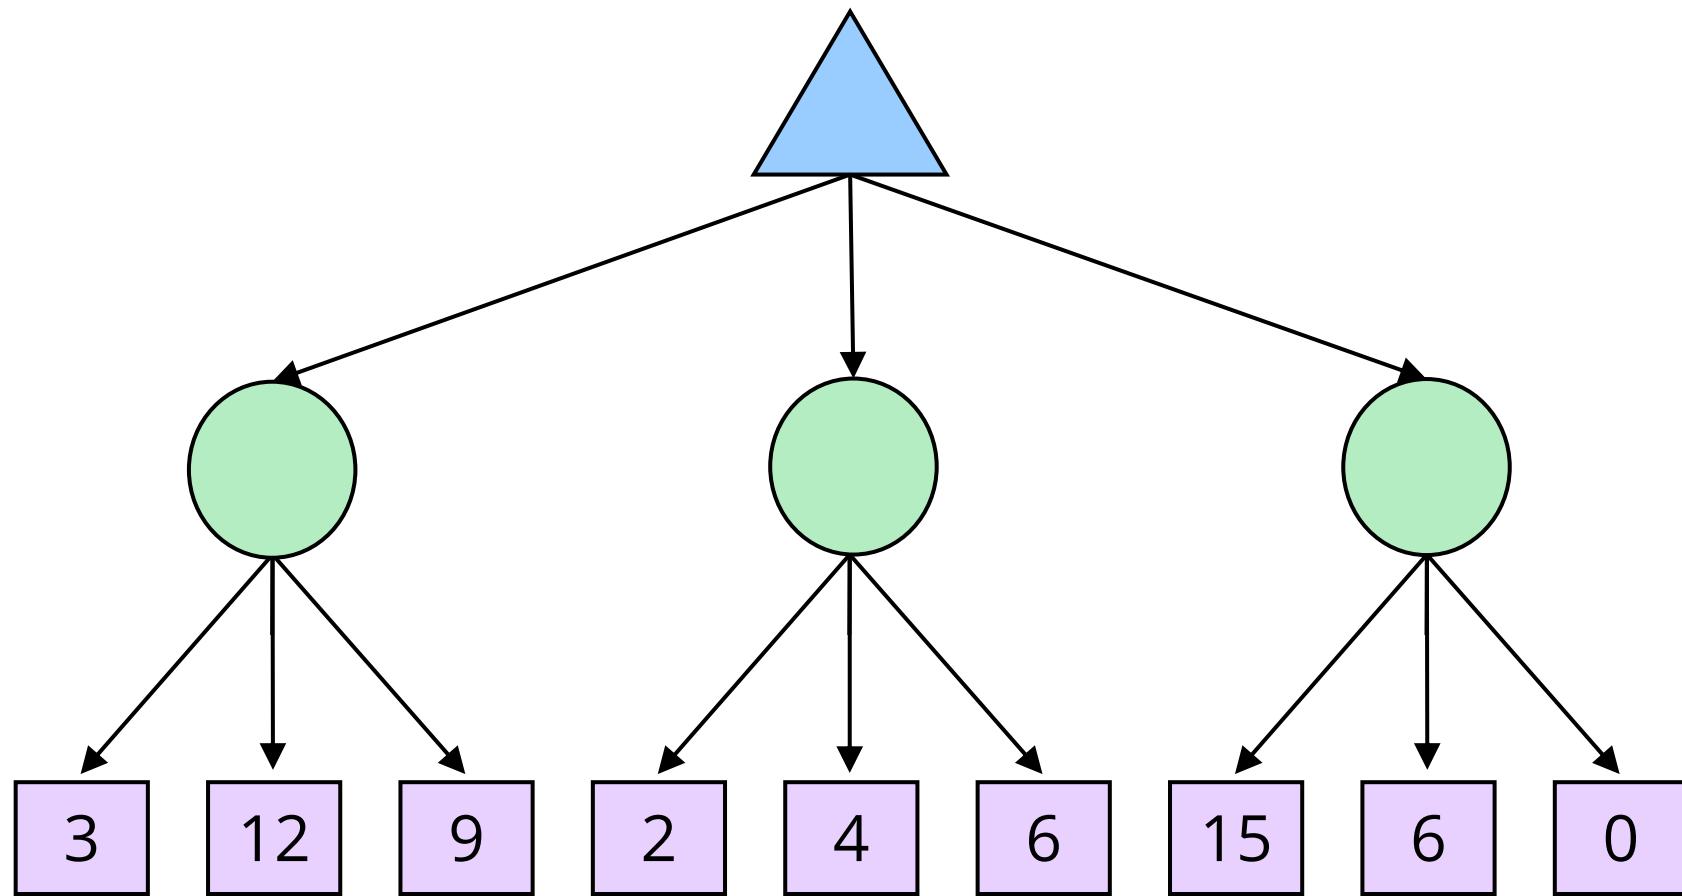
Expectimax Pseudocode

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p =  
        probability(successor)  
        v += p *  
        value(successor)  
    return v
```

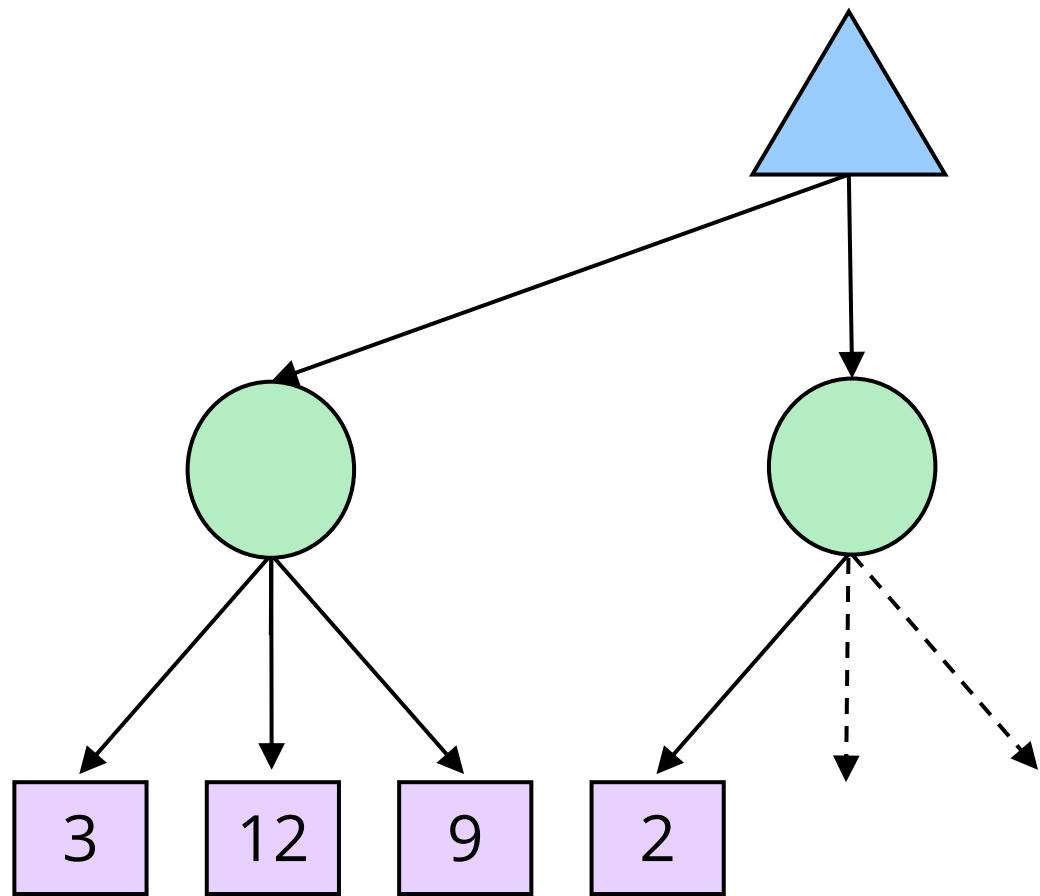


$$v = \frac{1}{2} \cdot (8) + \frac{1}{3} \cdot (24) + \frac{1}{6} \cdot (-12)$$

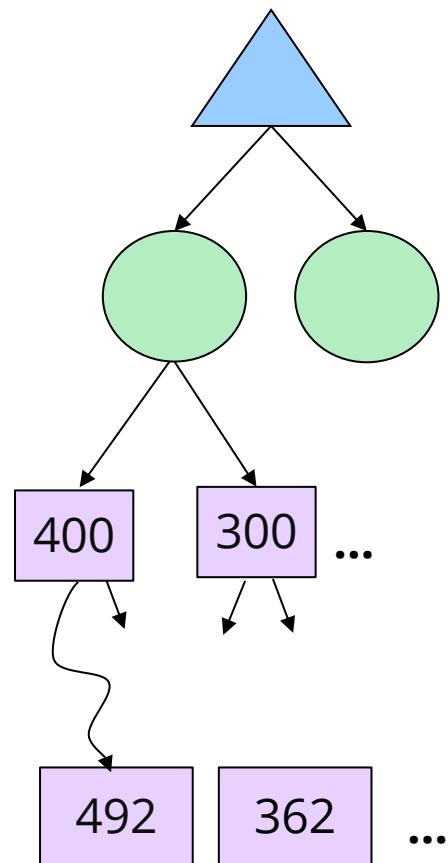
Title



Title

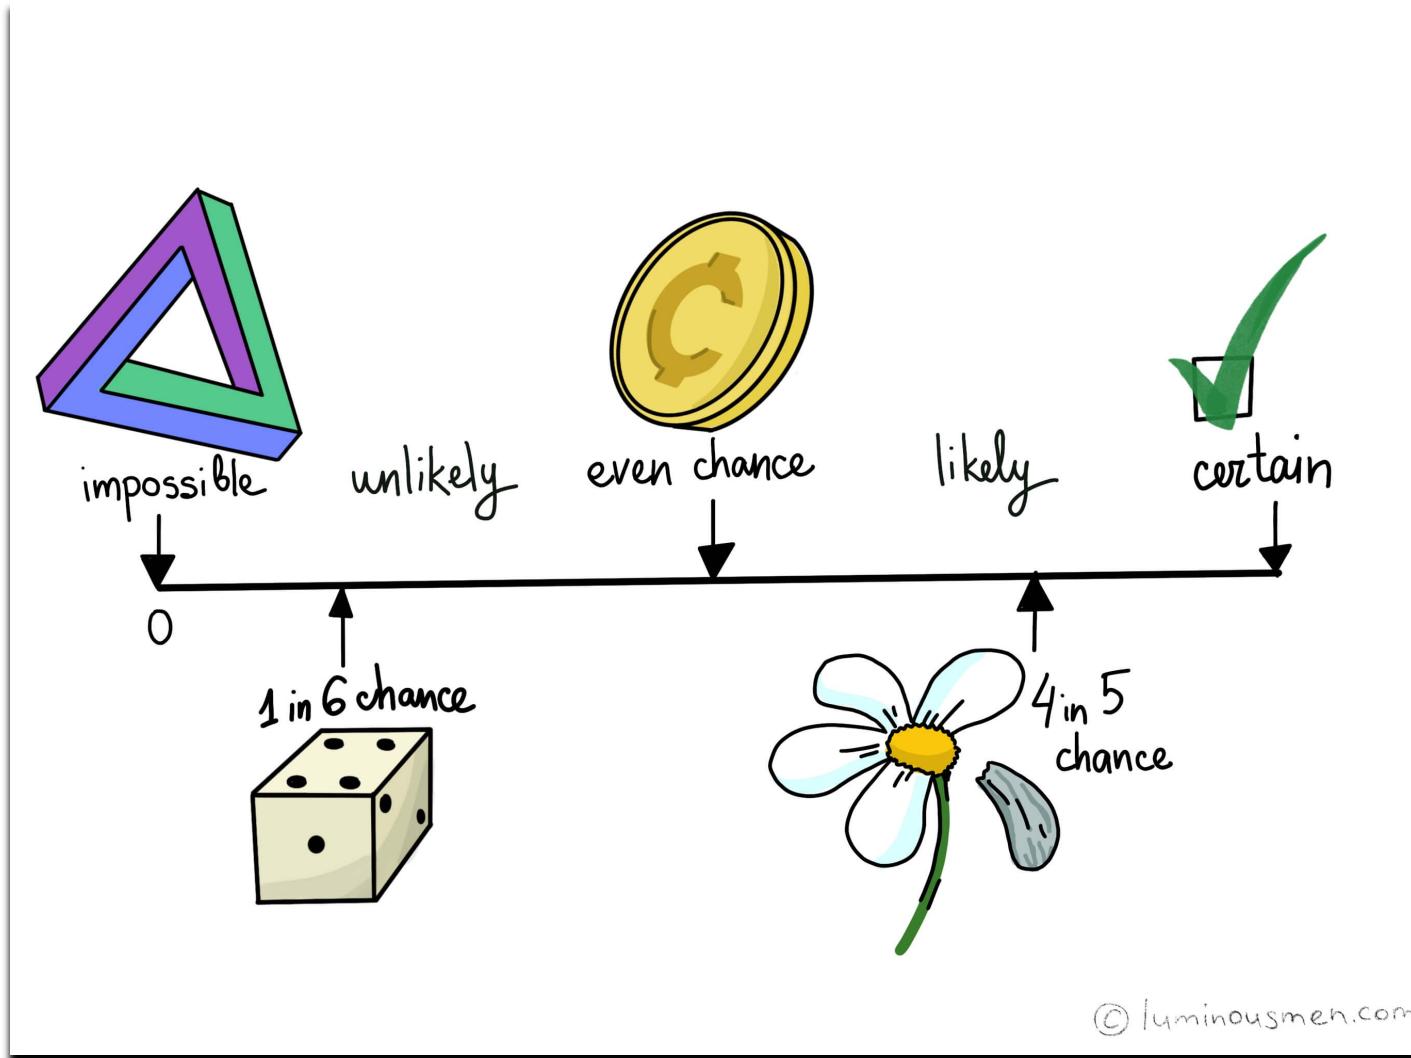


Title



Estimate of true expectimax value (which would require a lot of work to compute)

Probabilities



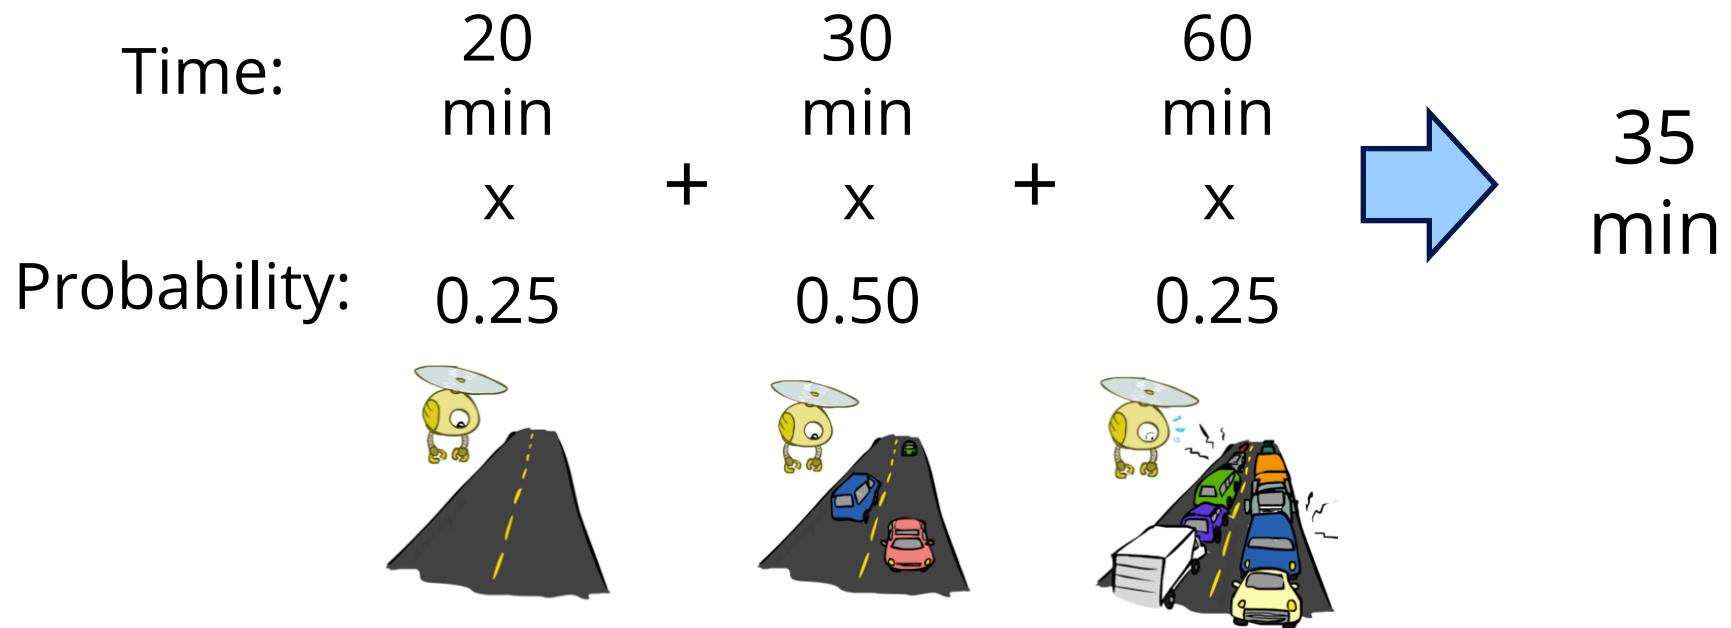
Probabilities

- A **random variable** represents an event whose outcome is unknown
- A **probability distribution** is an assignment of weights to outcomes
- Example: Traffic on freeway
 - Random variable: T = whether there's traffic
 - Outcomes: $T \in \{\text{none}, \text{light}, \text{heavy}\}$
 - Distribution: $P(T = \text{none}) = 0.25, P(T = \text{light}) = 0.50, P(T = \text{heavy}) = 0.25$
- Some laws of probability (more later):
 - Probabilities are always non-negative
 - Probabilities over all possible outcomes sum to one
- As we get more evidence, probabilities may change:
 $P(T = \text{heavy}) = 0.25, P(T = \text{heavy} | \text{Hour} = 8\text{am}) = 0.60$
We'll talk about methods for reasoning and updating probabilities later



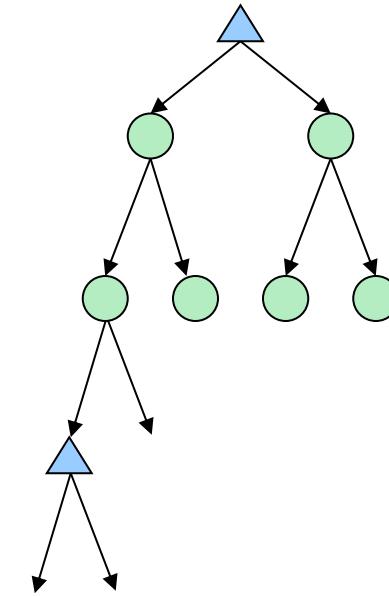
Probabilities

- The expected value of a function of a random variable is the average, weighted by the probability distribution over outcomes
- Example: How long to get to the airport?



What Probabilities to Use?

- In expectimax search, we have a probabilistic model of how the opponent (or environment) will behave in any state
 - Model could be a simple uniform distribution (roll a die)
 - Model could be sophisticated and require a great deal of computation
 - We have a chance node for any outcome out of our control: opponent or environment
 - The model might say that adversarial actions are likely!
- For now, assume each chance node magically comes along with probabilities that specify the distribution over its outcomes



Having a probabilistic belief about another agent's action does not mean that the agent is flipping any coins!

Title

- **Objectivist / frequentist answer:**

Averages over repeated *experiments*

E.g. empirically estimating $P(\text{rain})$ from historical observation

Assertion about how future experiments will go (in the limit)

New evidence changes the *reference class*

Makes one think of *inherently random* events, like rolling dice

- **Subjectivist / Bayesian answer:**

Degrees of belief about unobserved variables

E.g. an agent's belief that it's raining, given the temperature

E.g. agent's belief how an opponent will behave, given the state

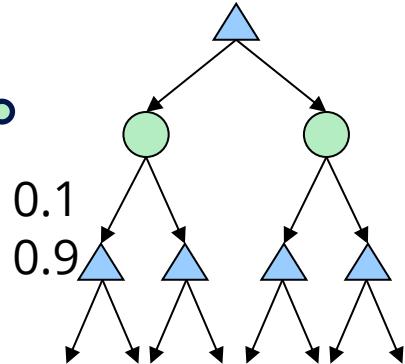
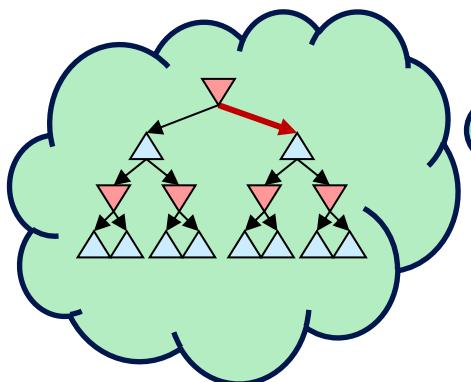
Often *learn* probabilities from past experiences (more later)

New evidence *updates beliefs* (more later)



Quiz: Informed Probabilities

- Let's say you know that your opponent is actually running a depth 2 minimax, using the result 80% of the time, and moving randomly otherwise
- Question: What tree search should you use?



- **Answer: Expectimax!**

- To figure out EACH chance node's probabilities, you have to run a simulation of your opponent
- This kind of thing gets very slow very quickly
- Even worse if you have to simulate your opponent simulating you...
- ... except for minimax, which has the nice property that it all collapses into one game tree

Title

- **Dice rolls increase b : 21 possible rolls with 2 dice**
 - Backgammon ≈ 20 legal moves
 - Depth 2 $\rightarrow 20 \times (21 \times 20)^3 = 1.2 \times 10^9$
- **As depth increases, probability of reaching a given search node shrinks**
 - So usefulness of search is diminished
 - So limiting depth is less damaging
 - But pruning is trickier...
- **Historic AI: TDGammon uses depth-2 search + very good evaluation function + reinforcement learning \rightarrow world-champion level play**
- **1st AI world champion in any game!**

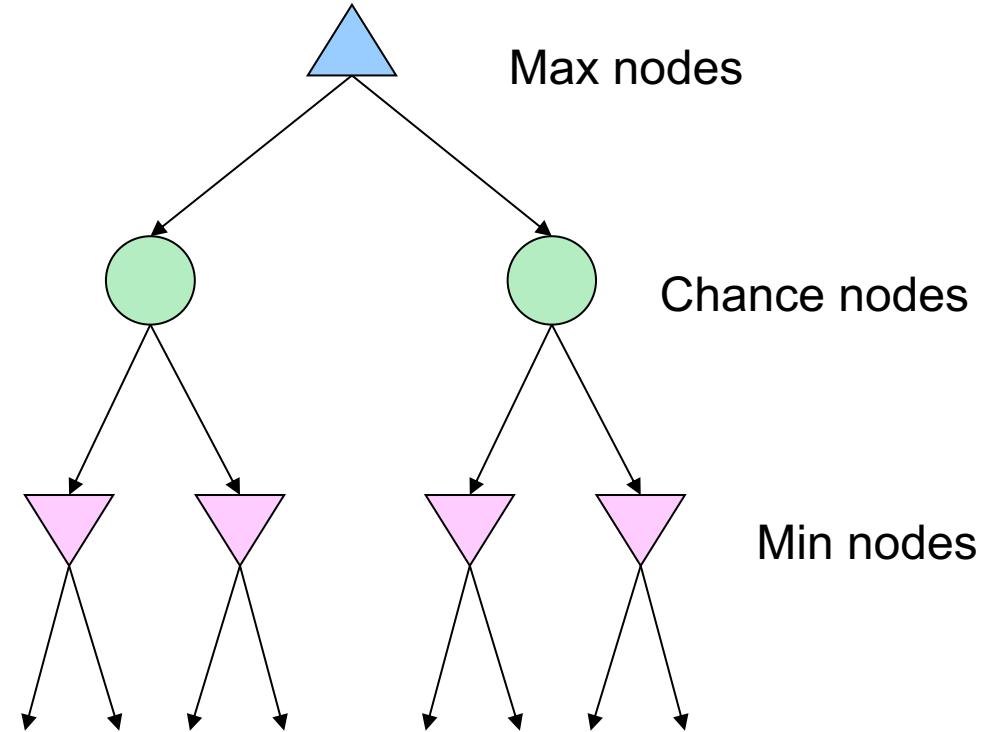


Title

- **E.g. Backgammon**
- **Expectiminimax**

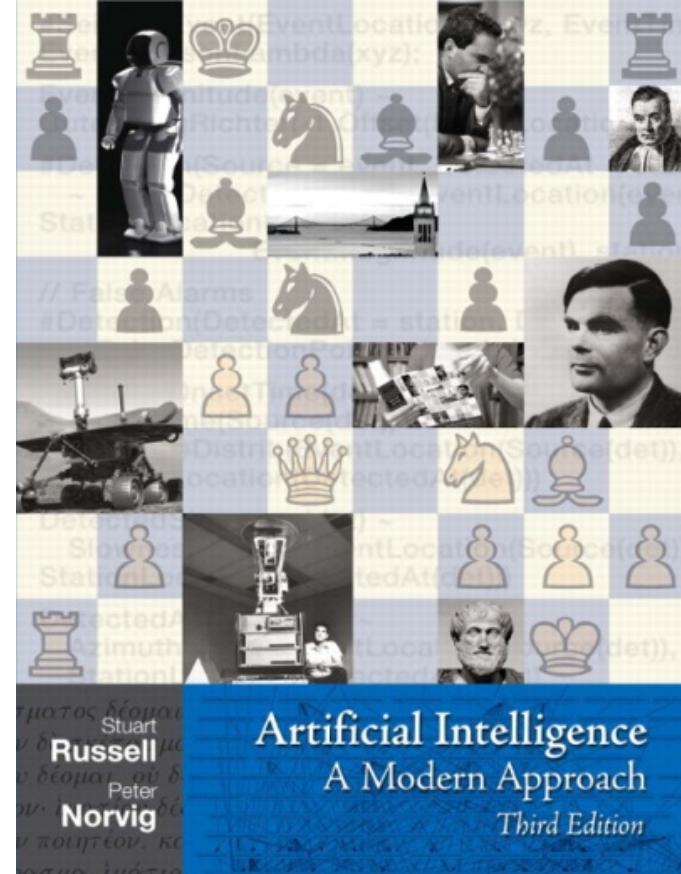
Environment is an extra “random agent” player that moves after each min/max agent

Each node computes the appropriate combination of its children



Utilities

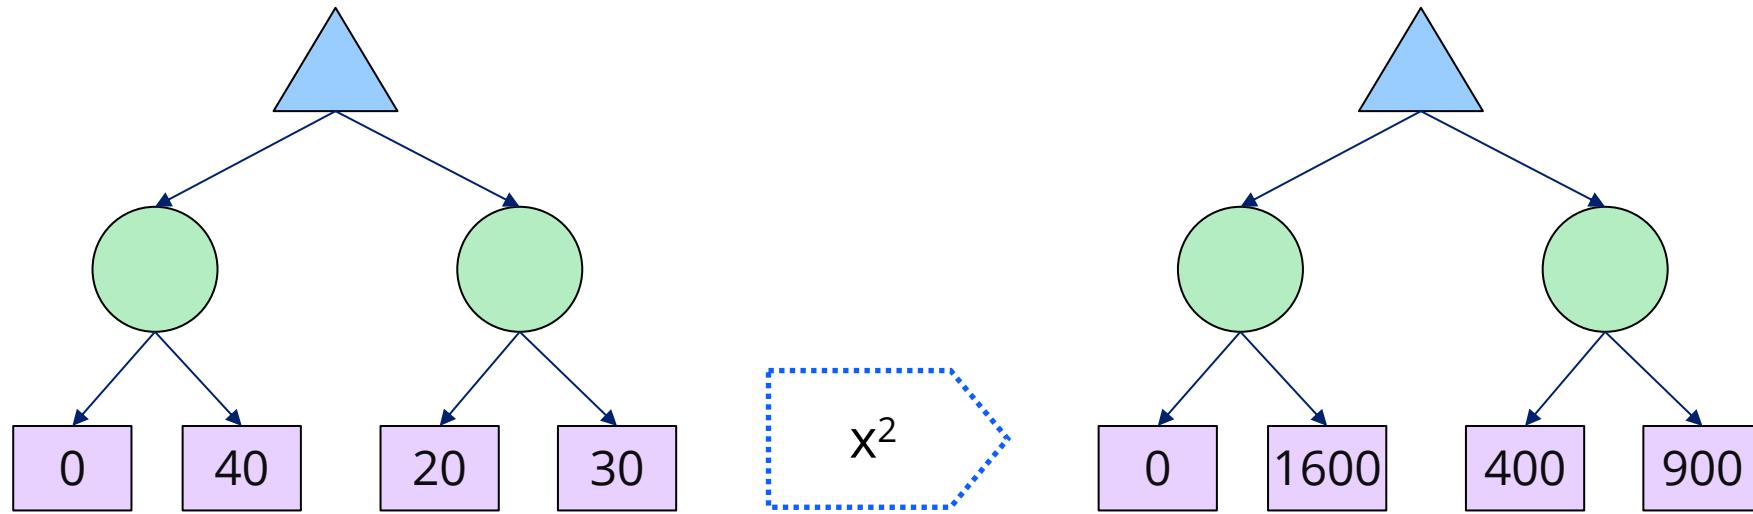
AIMA Chapter 16.1-16.3



Maximum Expected Utility

- Why should we **average** utilities? Why not minimax?
- **Principle of maximum expected utility:**
A rational agent should choose the action that **maximizes its expected utility, given its knowledge**
- Questions:
 - Where do utilities come from?
 - How do we know such utilities even exist?
 - How do we know that averaging even makes sense?
 - What if our behavior (preferences) can't be described by utilities?

Title



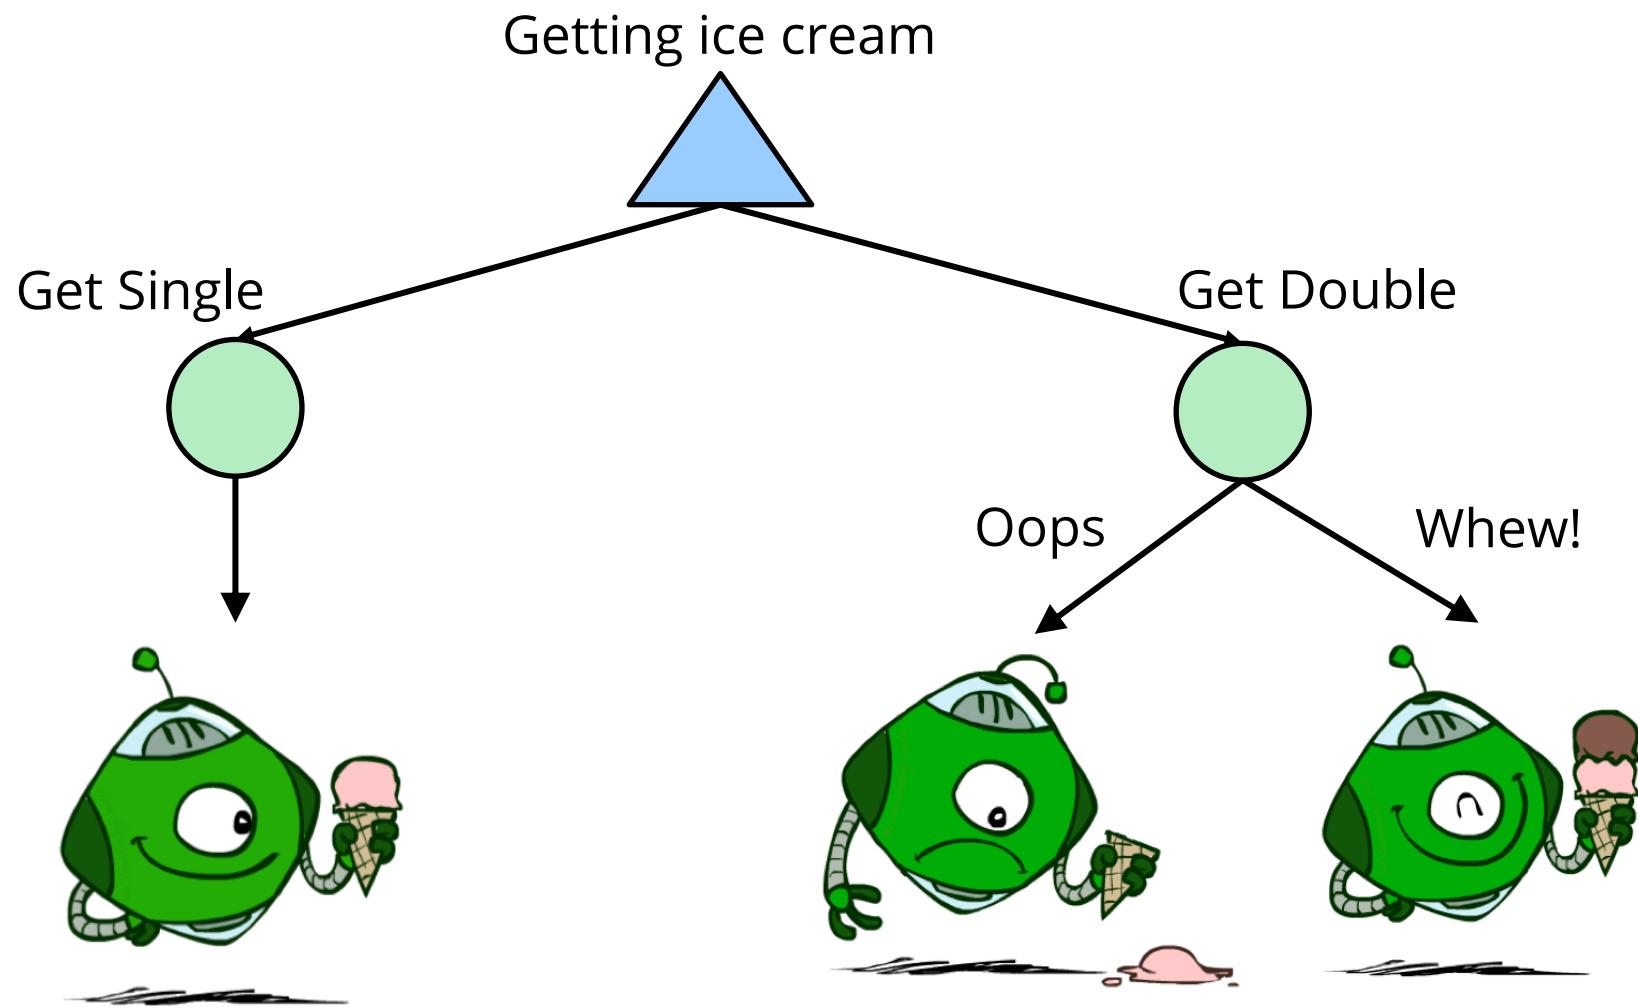
- **For worst-case minimax reasoning, terminal function scale doesn't matter**
We just want better states to have higher evaluations (get the ordering right)
We call this **insensitivity to monotonic transformations**
- **For average-case expectimax reasoning, we need *magnitudes* to be meaningful**

Title

- Utilities are functions from outcomes (states of the world) to real numbers that describe an agent's preferences
- Where do utilities come from?
 - In a game, may be simple (+1/-1)
 - Utilities summarize the agent's goals
 - Theorem: any “rational” preferences can be summarized as a utility function
- We hard-wire utilities and let behaviors emerge
 - Why don't we let agents pick utilities?
 - Why don't we prescribe behaviors?



Utilities: Uncertain Outcomes



Title

- An agent must have preferences among:

Prizes: A , B , etc.

Lotteries: situations with uncertain prizes

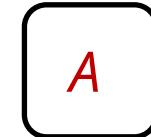
$$L = [p, A; (1 - p), B]$$

- Notation:

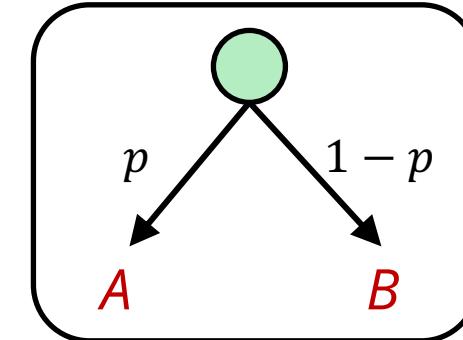
Preference: $A > B$

Indifference: $A \sim B$

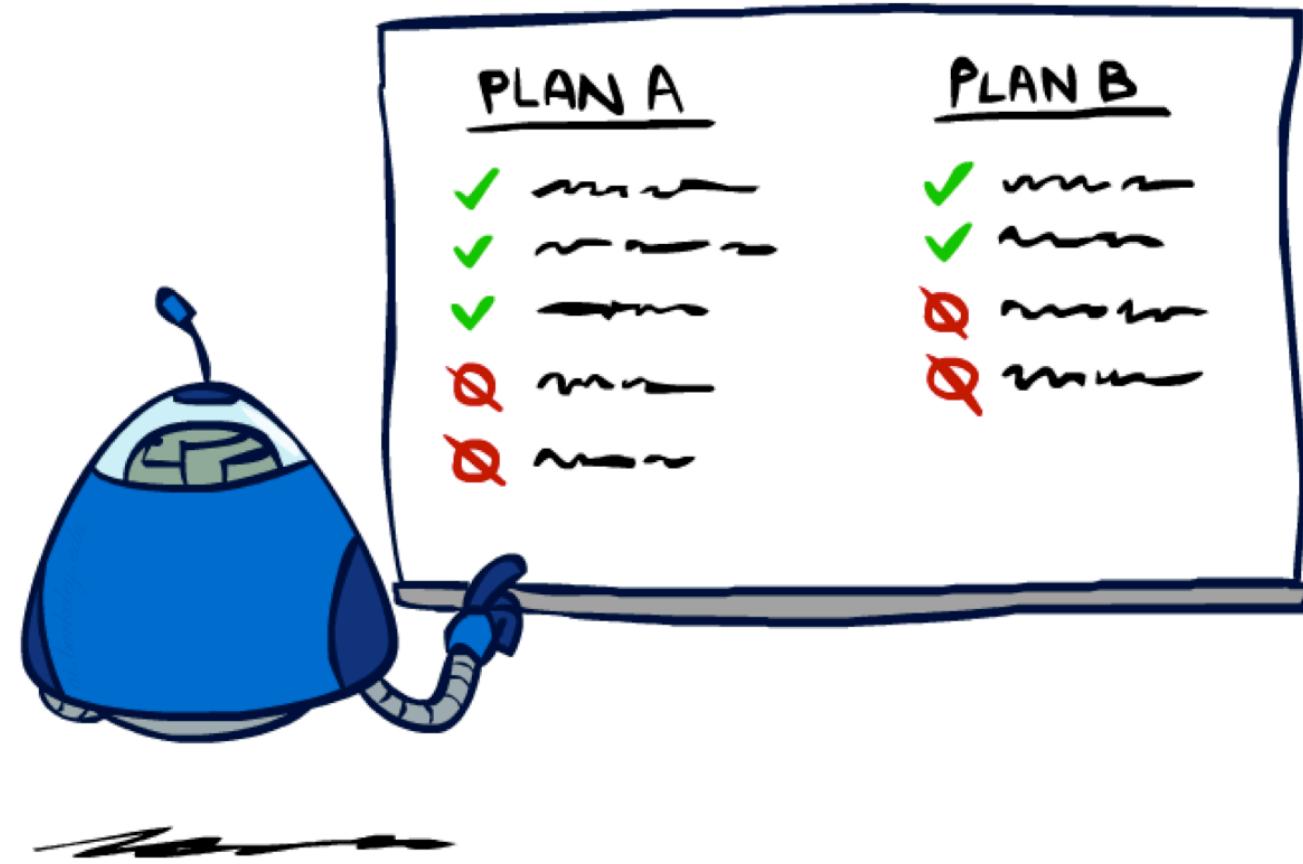
A Prize



A Lottery



Title



Title

- **We want some constraints on preferences before we call them rational, such as:**

Axiom of Transitivity: $(A > B) \wedge (B > C) \Rightarrow (A > C)$

- **For example: an agent with intransitive preferences can be induced to give away all of its money**

If $B > C$, then an agent with C would pay (say) 1 cent to get B

If $A > B$, then an agent with B would pay (say) 1 cent to get A

If $C > A$, then an agent with A would pay (say) 1 cent to get C

Title

The Axioms of Rationality

Orderability

$$(A \succ B) \vee (B \succ A) \vee (A \sim B)$$

Transitivity

$$(A \succ B) \wedge (B \succ C) \Rightarrow (A \succ C)$$

Continuity

$$A \succ B \succ C \Rightarrow \exists p [p, A; 1-p, C] \sim B$$

Substitutability

$$A \sim B \Rightarrow [p, A; 1-p, C] \sim [p, B; 1-p, C]$$

Monotonicity

$$A \succ B \Rightarrow$$

$$(p \geq q \Leftrightarrow [p, A; 1-p, B] \succeq [q, A; 1-q, B])$$



Theorem: Rational preferences imply behavior describable as maximization of expected utility

MEU Principle

- **Theorem [Ramsey, 1931; von Neumann & Morgenstern, 1944]**

Given any preferences satisfying these constraints, there exists a real-valued function U such that:

$$U(A) \geq U(B) \Leftrightarrow A \succcurlyeq B$$

$$U([p_1, S_1; \dots; p_n, S_n]) = \sum_i p_i U(S_i)$$

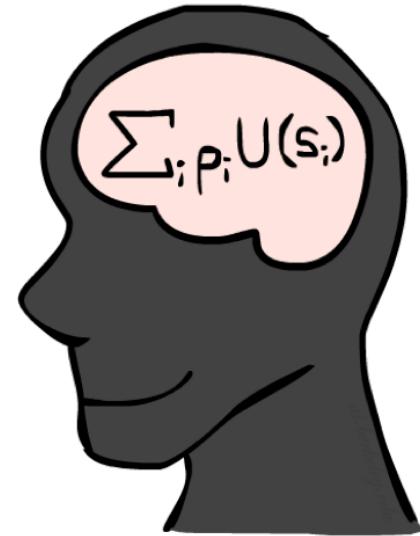
i.e. values assigned by U preserve preferences of both prizes and lotteries!

- **Maximum expected utility (MEU) principle:**

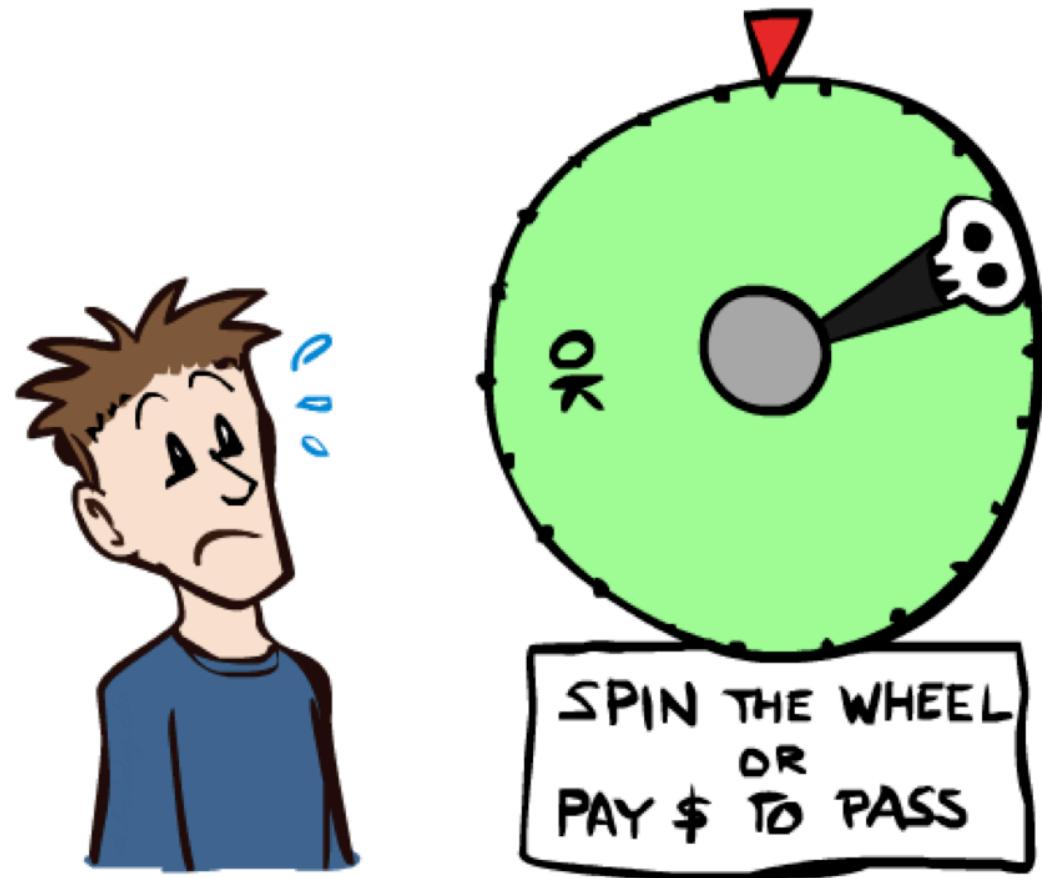
Choose the action that maximizes expected utility

Note: an agent can be entirely rational (consistent with MEU) without ever representing or manipulating utilities and probabilities

E.g., a lookup table for perfect tic-tac-toe, a reflex vacuum cleaner



Title



Utility Scales

- Normalized utilities: $u_+ = 1.0, u_- = 0.0$
- Micromorts: one-millionth chance of death, useful for paying to reduce product risks, etc.
- QALYs: quality-adjusted life years, useful for medical decisions involving substantial risk
- Note: behavior is invariant under positive linear transformation

$$U'(x) = k_1 U(x) + k_2 \text{ where } k_1 > 0$$

- With deterministic prizes only (no lottery choices), only ordinal utility can be determined, i.e., total order on prizes



Micromort examples

Death from	Micromorts per exposure
Scuba diving	5 per dive
Skydiving	7 per jump
Base-jumping	430 per jump
Climbing Mt. Everest	38,000 per ascent

1 micromort	
Train travel	6000 miles
Jet	1000 miles
Car	230 miles
Walking	17 miles
Bicycle	10 miles
Motorbike	6 miles



Title

- Utilities map states to real numbers. Which numbers?
- Standard approach to assessment (elicitation) of human utilities:

Compare a prize A to a standard lottery L_p between

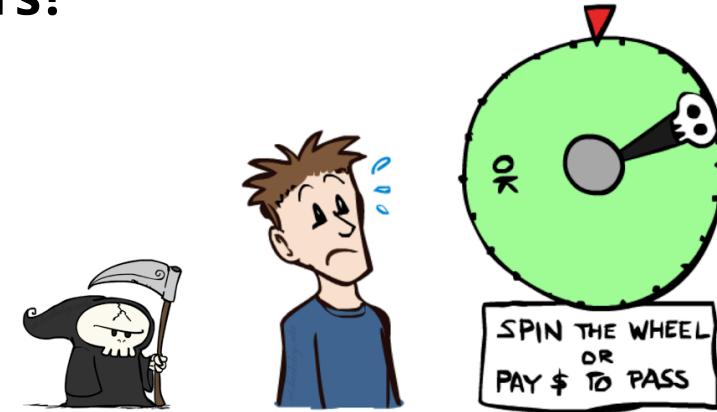
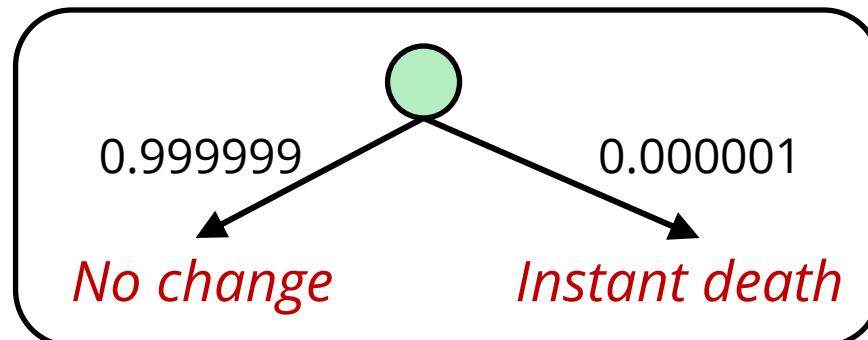
- “best possible prize” u_+ with probability p
- “worst possible catastrophe” u_- with probability $1 - p$

Adjust lottery probability p until indifference: $A - L_p$

Resulting p is a utility in $[0,1]$

Pay \$30

~



Money

- Money does not behave as a utility function, but we can talk about the utility of having money (or being in debt)
- Given a lottery $L = [p, \$X; (1 - p), \$Y]$

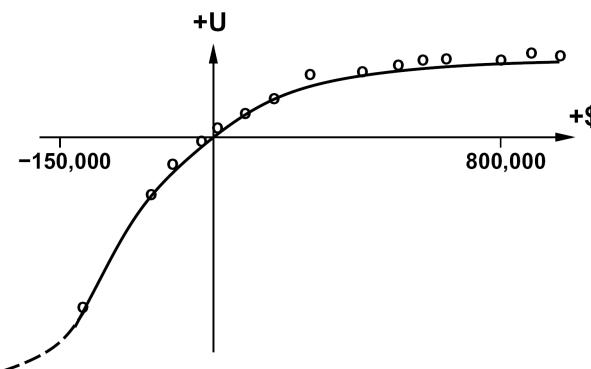
The expected monetary value $EMV(L) = p \cdot X + (1 - p) \cdot Y$

$$U(L) = p \cdot U(\$X) + (1 - p) \cdot U(\$Y)$$

Typically, $U(L) < U(EMV(L))$

In this sense, people are **risk-averse**

When deep in debt, people are **risk-prone**



Example: Insurance

- Consider the lottery [0.5, \$1000; 0.5, \$0]

What is its **expected monetary value**? (\$500)

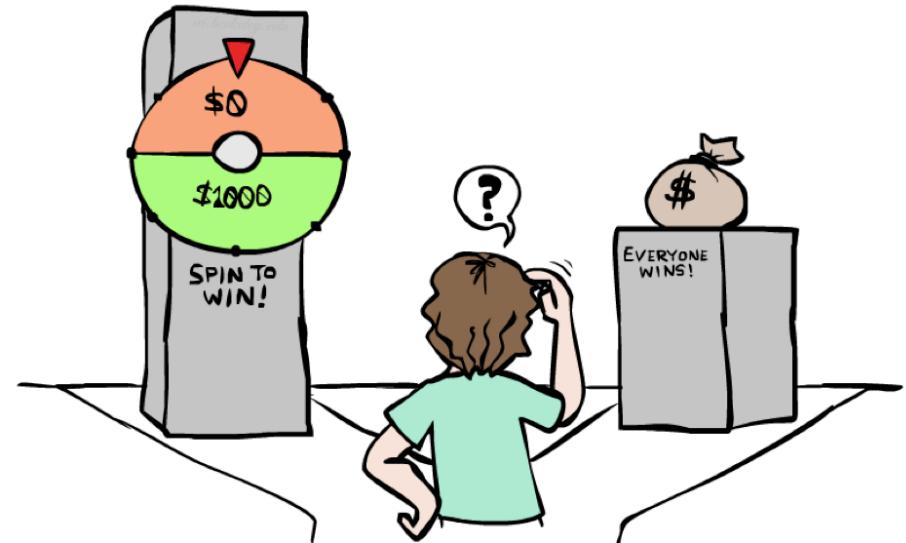
What is its **certainty equivalent**?

- Monetary value acceptable in lieu of lottery
- \$400 for most people

Difference of \$100 is the **insurance premium**

- There's an insurance industry because people will pay to reduce their risk
- If everyone were risk-neutral, no insurance needed!

It's win-win: you'd rather have the \$400 and the insurance company would rather have the lottery (their utility curve is linear and they have many lotteries)



Example: Human Rationality?

- **Famous example of Allais (1953)**

A: [0.8, \$4k; 0.2 \$0]

B: [1.0, \$3k; 0.0, \$0]

C: [0.2, \$4k; 0.8, \$0]

D: [0.2, \$3k; 0.75, \$0]

- **Most people prefer $B > A, C > D$**

- **But if $U(\$0) = 0$, then**

$$B > A \Rightarrow U(\$3k) > 0.8 \cdot U(\$4k)$$

$$C > D \Rightarrow 0.8 \cdot U(\$4k) > U(\$3k)$$

