

CIS 4210/5210:
ARTIFICIAL INTELLIGENCE

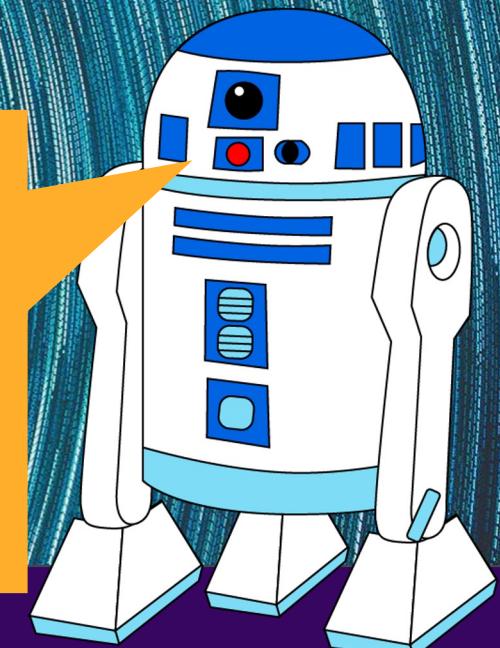
Search Problems

Professor Chris Callison-Burch

HW1 is due tonight

Optional HW2
Recitation on
Wednesday 8-9pm

Read the AIMA textbook
Chapter 3 “Problem Solving
by Search” (Sections 3.1-3.4).
Quiz on Canvas due by
Sunday.



Example Search Problems

AIMA 3.1-3.3

Reflex Agents

A simple reflex agent is one that selects an action based only on the **current percept**.

It **ignores** the rest of the **percept history**.



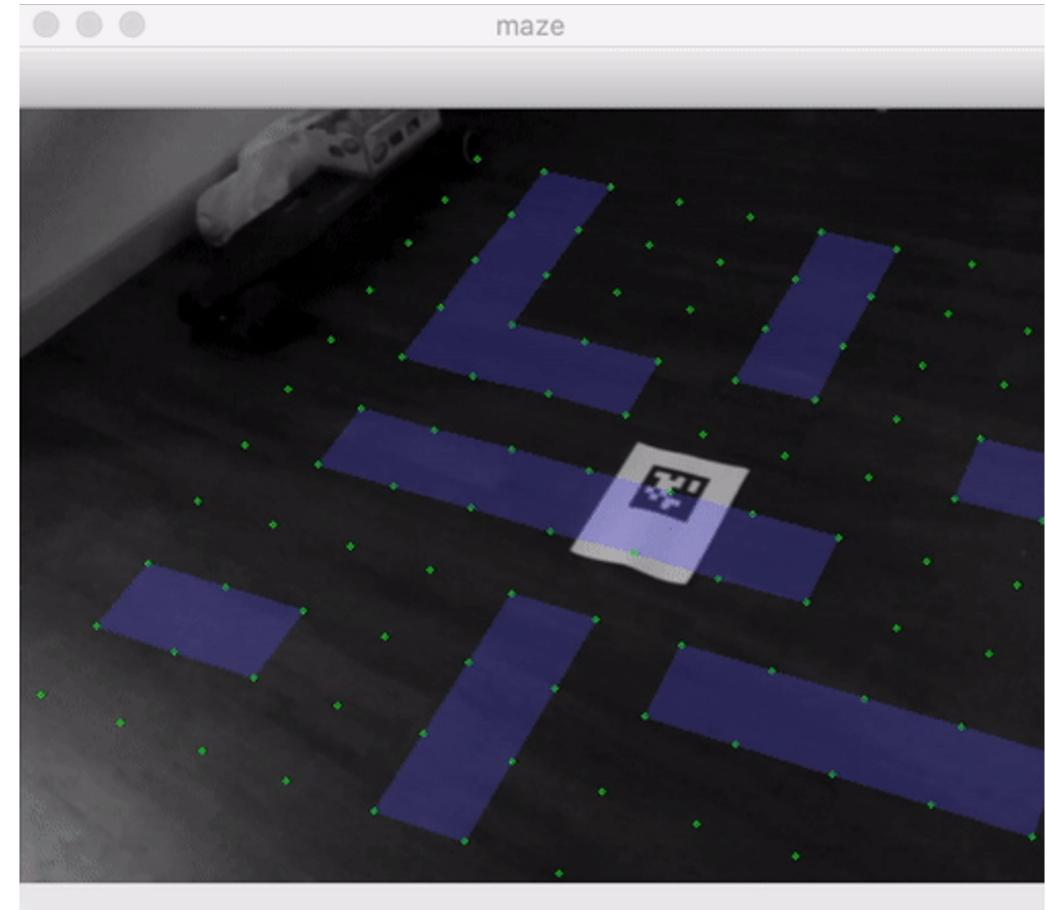
Problem-Solving Agent

A problem-solving agent must **plan ahead**.

The computational process that it undertakes is called **search**.

It will consider a **sequence of actions** that form a **path** to a **goal state**.

Such a sequence is called a **solution**.



Impact of Task Environments

The properties of the task environments change the types of solutions that we need.

If an environment is:

- **Fully observable**
- **Deterministic**
- **Known environment**

The solution to any problem in such an environment is a fixed sequence of actions.

In environments that are

- **Partially observable** or
- **Nondeterministic**

The solution must recommend different future actions depending on the what percepts it receives. This could be in the form of a *branching strategy*.

Example search problem: Holiday in Romania



You are here

Oradea

71

140

118

111

70

75

120

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

111

70

75

120

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

146

138

101

90

98

86

85

92

142

87

99

211

111

80

97

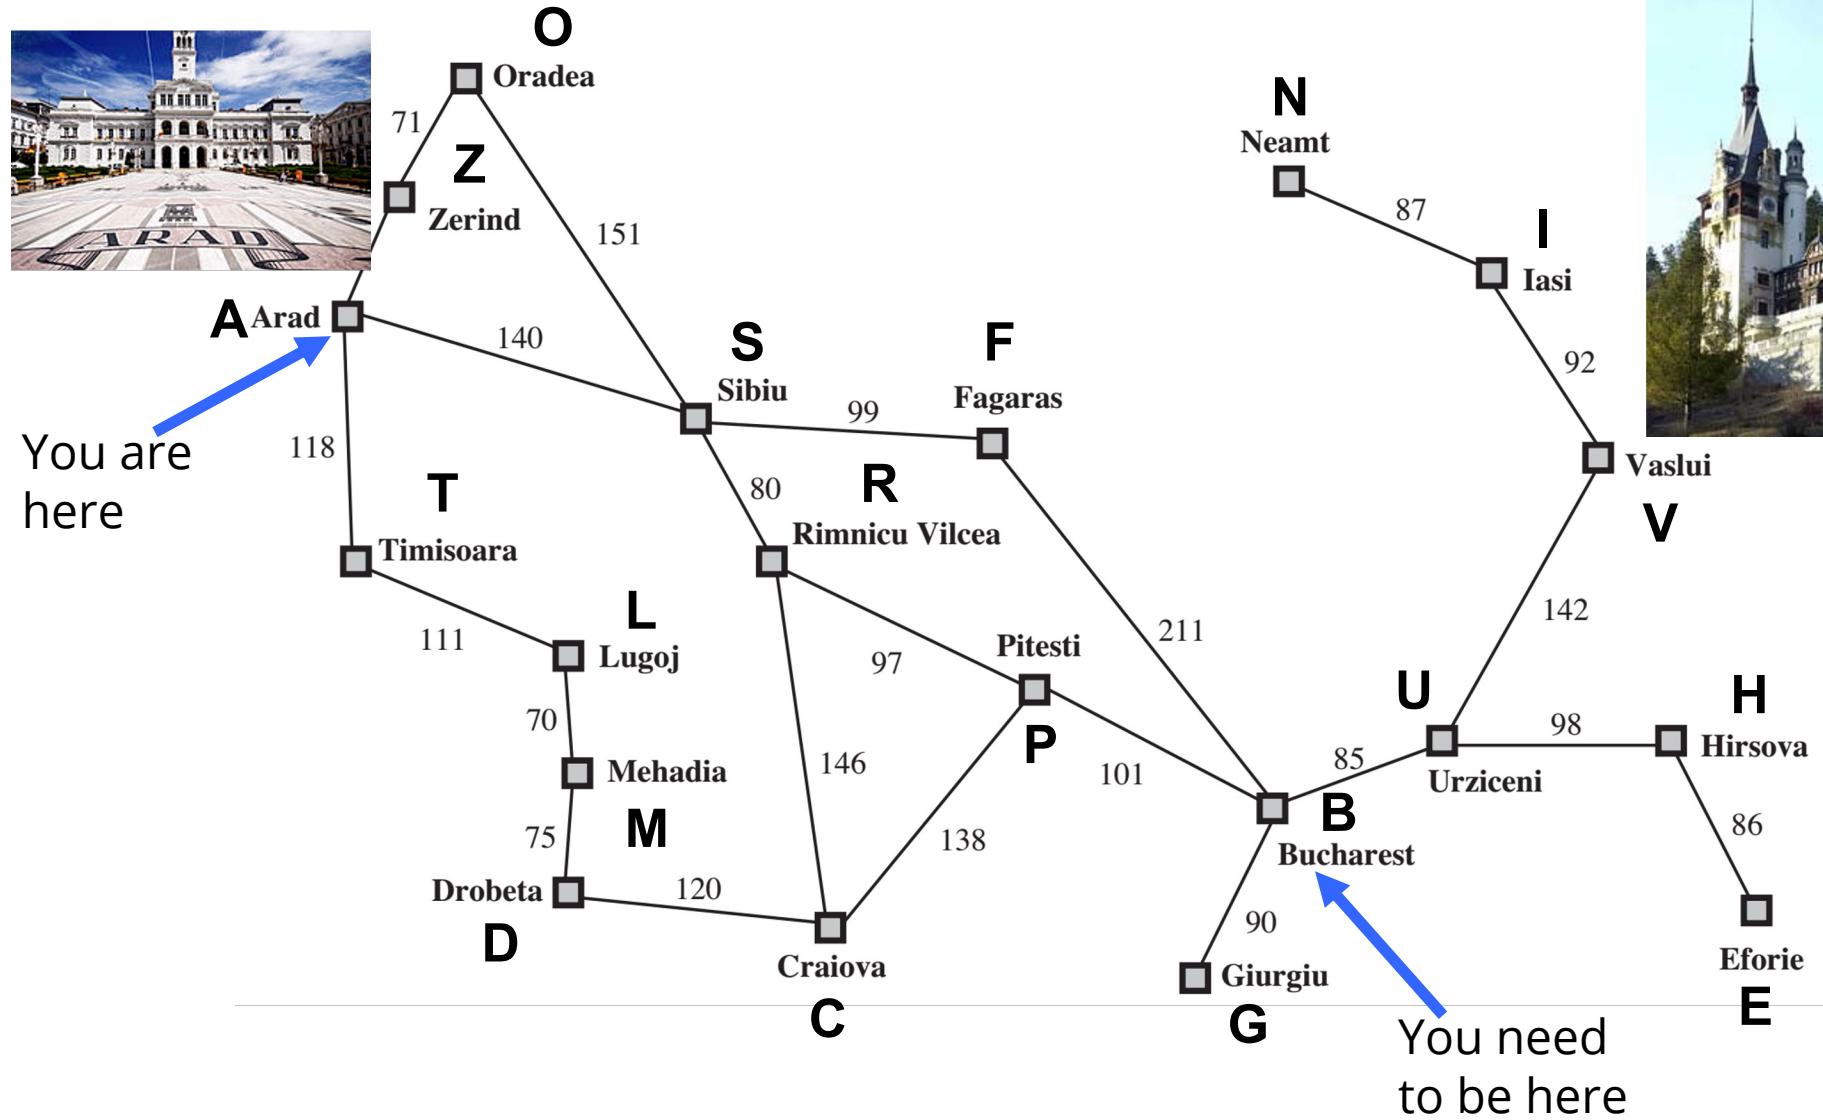
146

138

101

<

Example search problem: Holiday in Romania



Holiday in Romania

On holiday in Romania; currently in Arad

- Flight leaves tomorrow from Bucharest

Formulate ***goal***

- Be in Bucharest

Formulate ***search problem***

- States: various cities
- Actions: drive between cities
- Performance measure: minimize travel time / distance

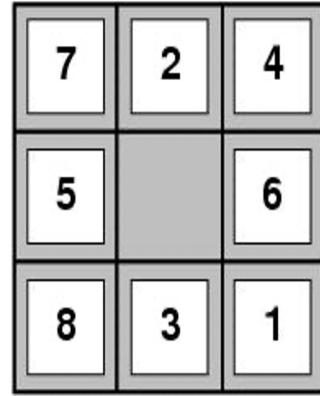
Find ***solution***

- Sequence of cities; e.g. Arad, Sibiu, Fagaras, Bucharest, ...

Example search problem: 8-puzzle

Formulate **goal**

- Pieces to end up in order as shown...



Start State



Goal State



Formulate **search problem**

- States:** configurations of the puzzle (9! configurations)
- Actions:** Move one of the movable pieces (≤ 4 possible)
- Performance measure:** minimize total moves

Find **solution**

- Sequence of pieces moved: 3,1,6,3,1,...

Why do we care about search? Doesn't ChatGPT solve everything?

Opinions and commentary from your Professor

What are Large Language Models?

Prompt

My favorite professor at the University of Pennsylvania is Chris Callison-Burch. He is a professor in the Computer and Information Science Department, and is the director of the Natural Language Processing Group at the Institute for Research in Cognitive Science. He is an incredibly accomplished professor and researcher, whose work has been widely published in top journals. He is also well-known for his research in machine translation, dialogue systems, and other areas of natural language processing. He is an inspiring professor who encourages his students to learn and provides a stimulating environment to learn.

knowledgeable = 47.95%
engaging = 19.51%
passionate = 10.00%
inspiring = 5.67%
talented = 2.67%
accomplished = 0.47%

Completion

f Natural Language Processing. He is fine translation, dialogue systems, an inspiring professor who ly, and provides a stimulating

Total: -5.35 logprob on 1 tokens
(86.28% probability covered in top 6 logits)

Submit



129

1

What are Large Language Models?

Chris Callison-Burch | University of Pennsylvania | Computer and Information Science
Department



Karen Rile | University of Pennsylvania | English Department

David Brainard | University of Pennsylvania | Psychology Department

Emily Falk | University of Pennsylvania | Annenberg School for Communication

Christopher Manning | Stanford University | Computer Science Department

Jeffrey Bigham | Carnegie Mellon University | Human-Computer Interaction Institute

Lera Boroditsky | U.C. San Diego | Cognitive Science Department

Few shot example

Prompts

Completions

"In-context Learning"

Submit



99



Penn Engineering

What are Large Language Models?

Write a biography for Chris Callison-Burch in a first person perspective. **Instruction following** 

Hi! My name is Chris Callison-Burch and I am a professor of Computer Science at the University of Pennsylvania. I have been in this field for over 15 years, and it has been an incredible journey.

I was born and raised in the United Kingdom, and I earned my undergraduate degree in Mathematics from the University of Cambridge. After that, I moved to the United States to pursue a doctorate in Computer Science from Stanford University.

My research focuses on natural language processing, which is the study of how computers can understand and interact with human language. I've developed several algorithms in this field, including the Syntax-based Machine Translation system (SMT), which is used by Google Translate and other translation services. I'm also the creator of the open-source software program Moses, which is used by researchers and academics around the world to make their own natural language processing tools.

Submit

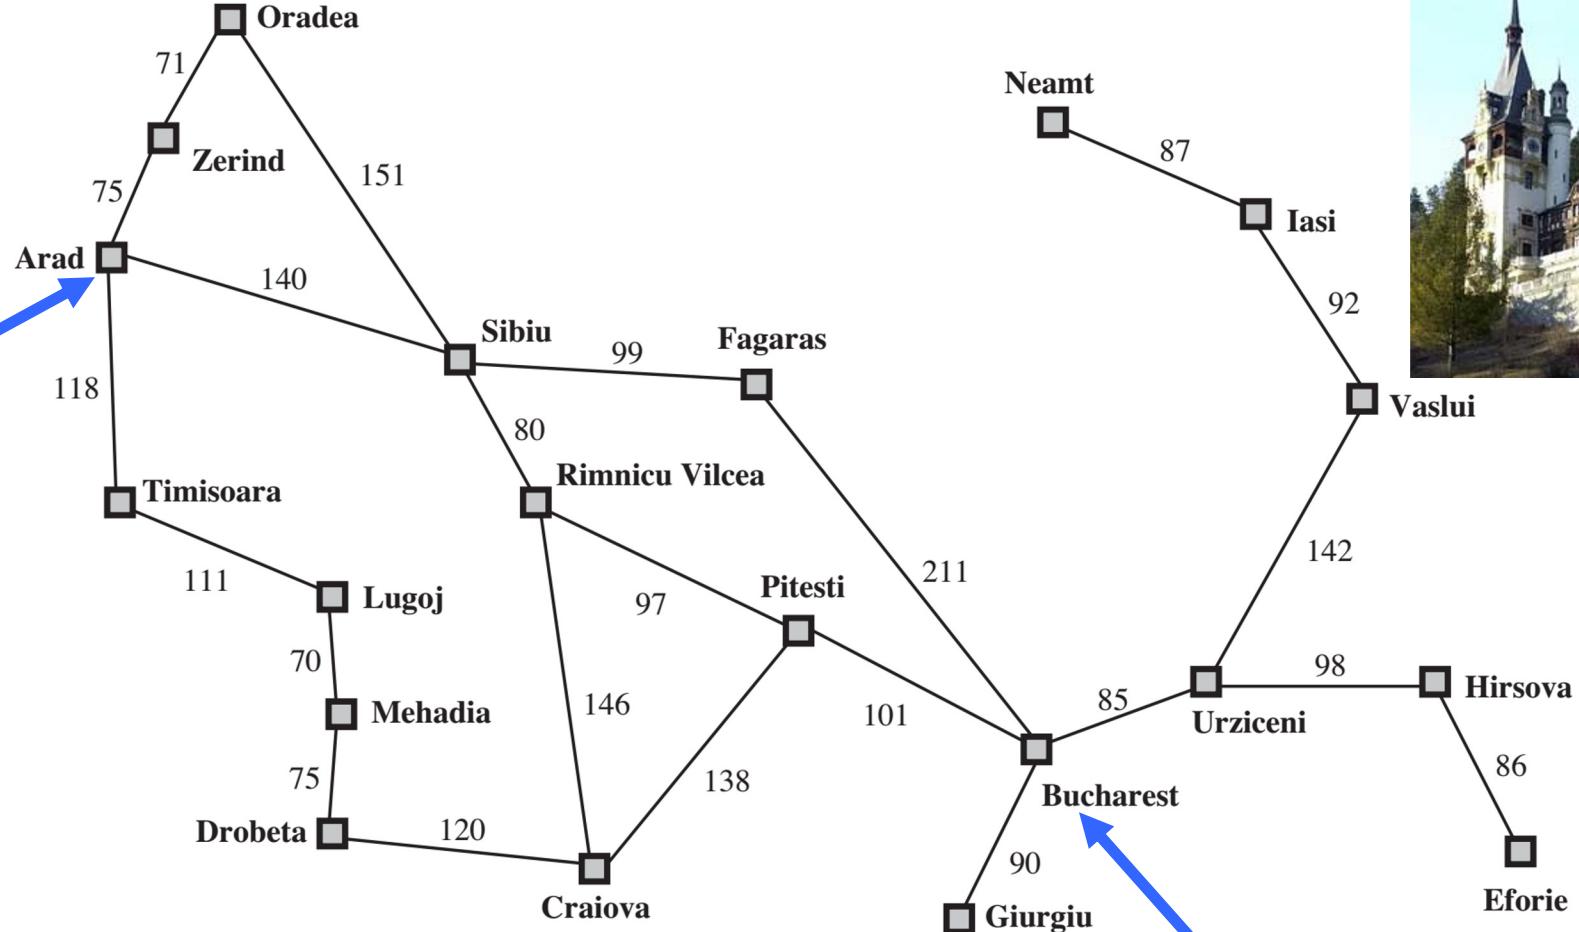


272

Are LLMs Planning Agents? Let's ask!



You are here



You need
to be here

What's the shortest path from Arad to Bucharest?

The shortest path from Arad to Bucharest is Arad -> Sibiu -> Rimnicu Vilcea -> Pitesti -> Bucharest.

What is the lowest path cost from Arad to Bucharest?

The lowest path cost from Arad to Bucharest is 366.

366 = 61.95%

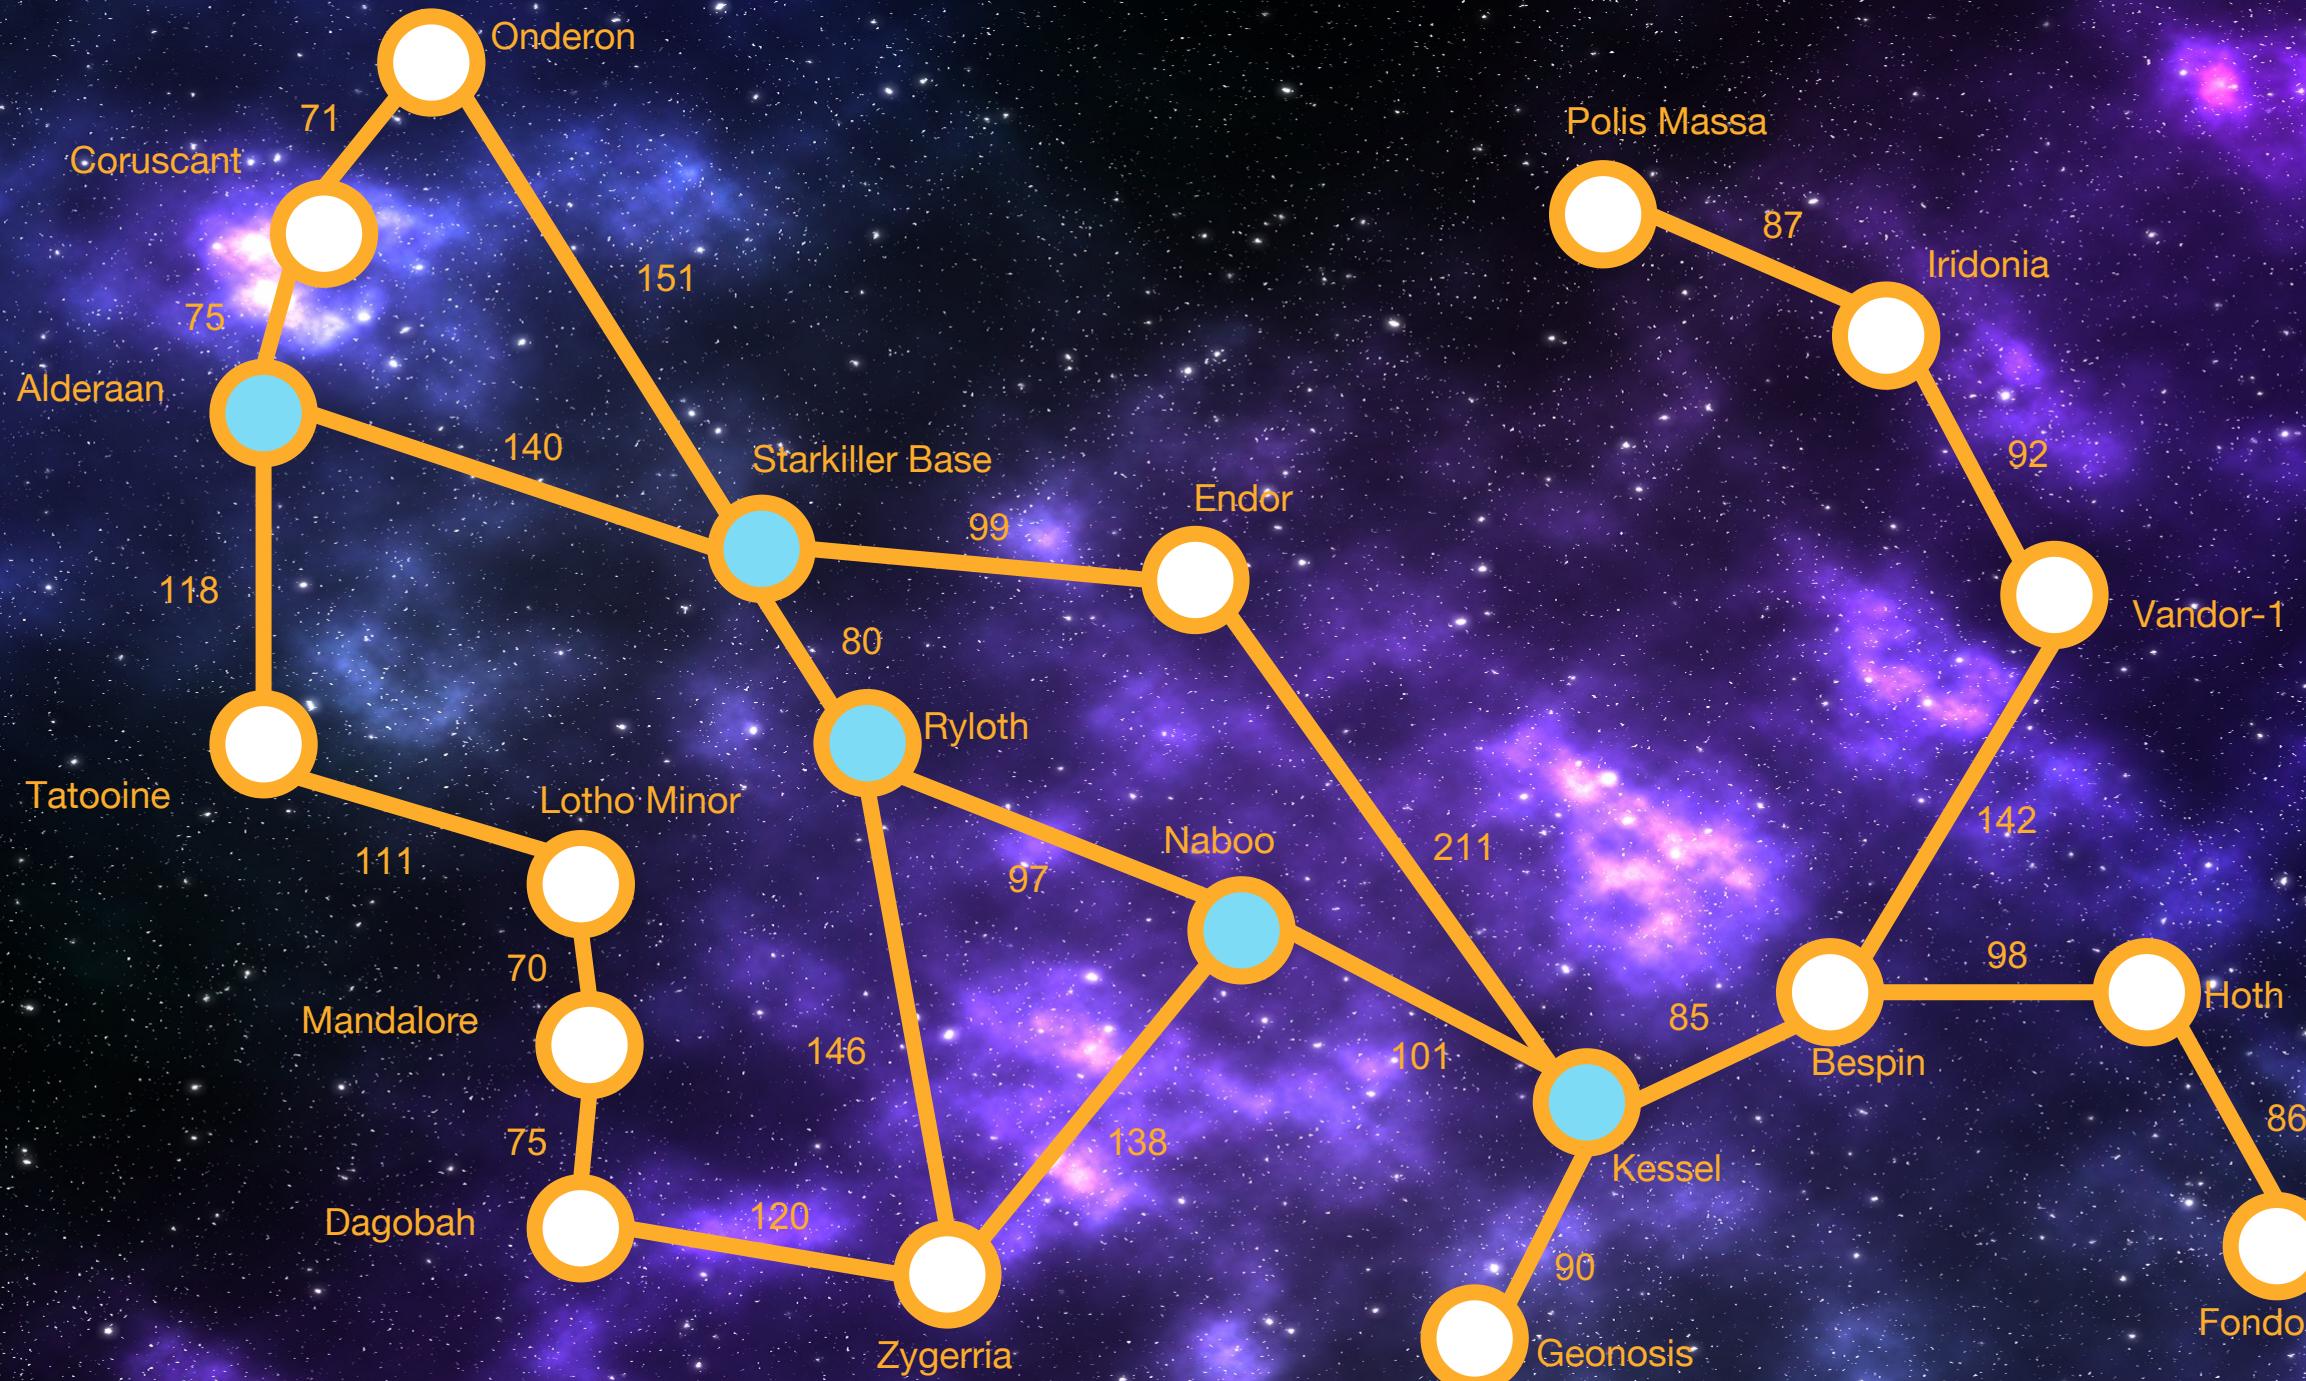
418 = 35.10%

118 = 0.72%

4 = 0.28%

318 = 0.25%

Total: -0.48 logprob on 1 tokens
(98.30% probability covered in top 5 logits)





Here's a list of distances between nodes in an undirected graph:

Alderaan		Coruscant		75
Alderaan		Starkiller Base		140
Alderaan		Tatooine		118
Coruscant	Onderon		71	
Onderon	Starkiller Base		151	
Tatooine		Lothal Minor		111
Lothal Minor		Mandalore		70
Mandalore		Dagobah		75
Dagobah		Zygerria		120
Zygerria		Ryloth		146
Zygerria		Naboo		138
Ryloth		Naboo		97
Ryloth		Starkiller Base		80
Starkiller Base		Endor		99
Kessel		Naboo		101
Kessel		Endor		211

Zygerria | Naboo | 138



Ryloth | Naboo | 97

Ryloth | Starkiller Base | 80

Starkiller Base | Endor | 99

Kessel | Naboo | 101

Kessel | Endor | 211

Kessel | Geonosis | 90

Kessel | Bespin | 85

Bespin | Hoth | 98

Hoth | Fondor | 86

Vandor-1 | Bespin | 142

Iridonia | Vandor-1 | | 92

Polis Massa | Iridonia | 87

What is the shortest path from Alderaan to Kessel?

Alderaan -> Tatooine -> Lotho Minor -> Mandalore -> Dagobah -> Zygerria -> Ryloth -> Starkiller

Base -> Endor -> Kessel (236 units)

CIS 4210/5210:
ARTIFICIAL INTELLIGENCE

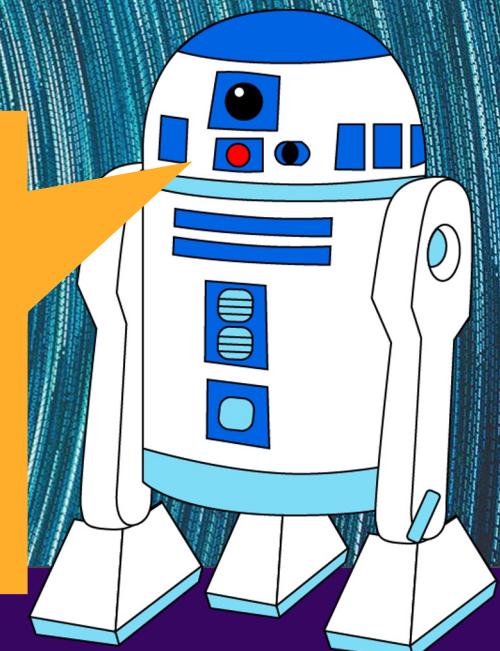
Search Problems

Professor Chris Callison-Burch

HW1 is due tonight

Optional HW2
Recitation on
Wednesday 8-9pm

Read the AIMA textbook
Chapter 3 “Problem Solving
by Search” (Sections 3.1-3.4).
Quiz on Canvas due by
Sunday.



Defining Search Problems

Formal Definition

1. **States:** a set S

2. An *initial state* $s_i \in S$

3. **Actions:** a set A

$\forall s \text{ } Actions(s) = \text{the set of actions that can be executed in } s, \text{ that are applicable in } s.$

4. **Transition Model:** $\forall s \forall a \in Actions(s) \text{ } Result(s, a) \rightarrow s_r$

s_r is called a *successor* of s

$\{s_i\} \cup Successors(s_i)^*$ = *state space*

5. **Path cost (Performance Measure):** Must be additive, e.g. sum of distances, number of actions executed, ...

$c(x, a, y)$ is the step cost, assumed ≥ 0

- (where action a goes from state x to state y)

6. **Goal test:** $Goal(s)$

Can be implicit, e.g. $checkmate(s)$

s is a *goal state* if $Goal(s)$ is true

Vacuum World

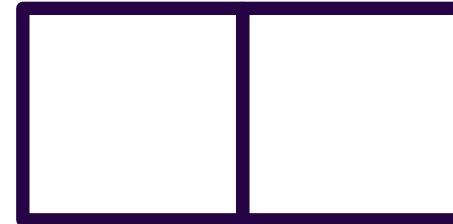
States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.



Vacuum World

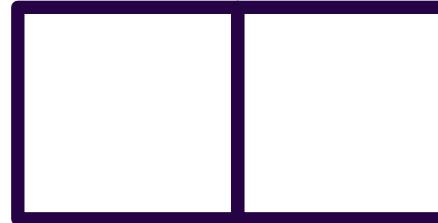
States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.



Vacuum World

States: A state of the world says which objects are in which cells.

In a simple two cell version,

- the agent can be in one cell at a time
- each cell can have dirt or not

2 positions for agent * 2^2 possibilities for dirt = 8 states.

With n cells, there are $n*2^n$ states.

Goal states: States where everything is clean.



One state is designated as the **initial state**

Vacuum World



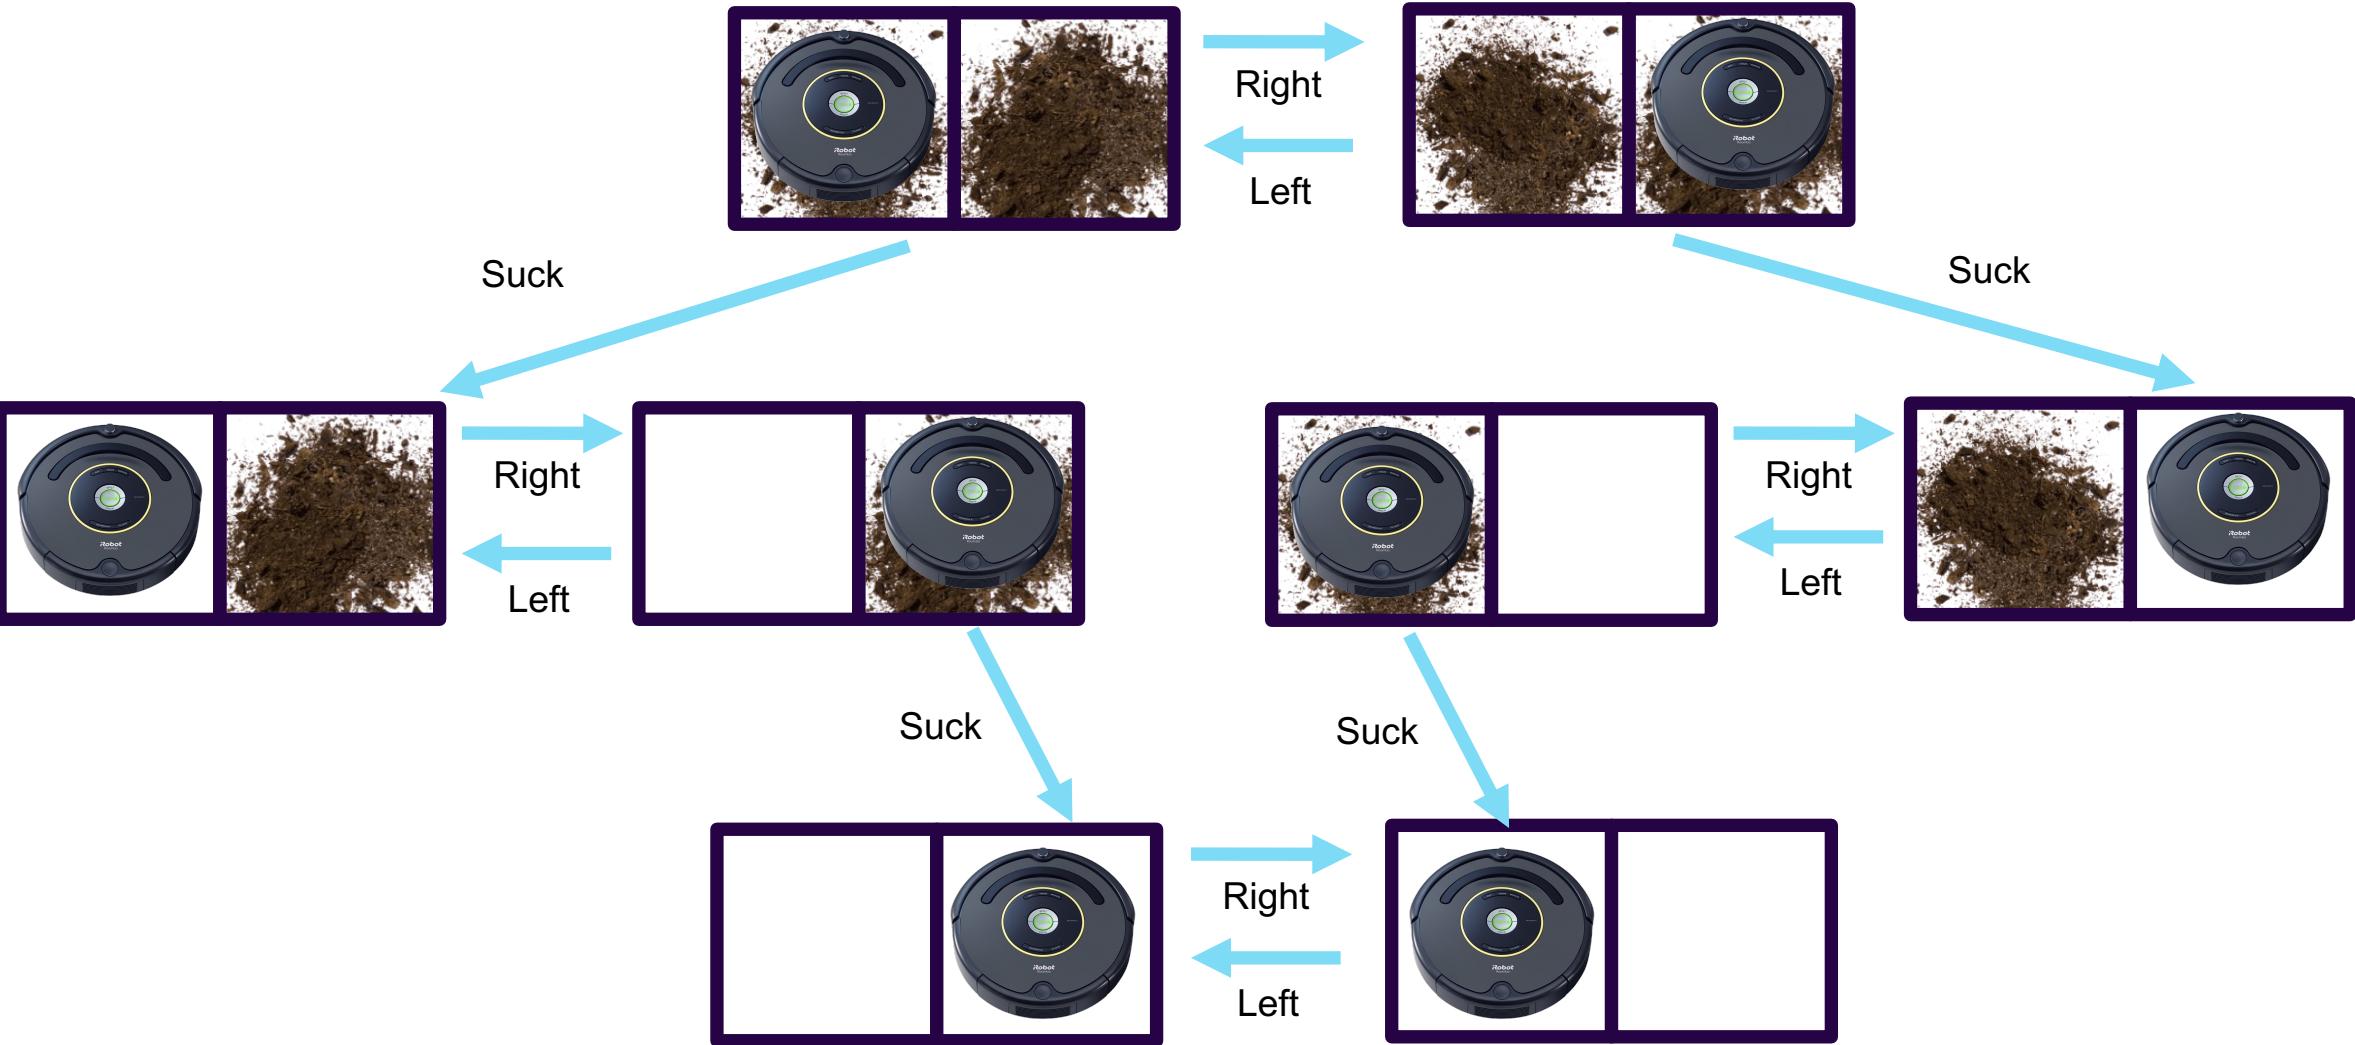
Actions:

- *Suck*
- *Move Left*
- *Move Right*
- *(Move Up)*
- *(Move Down)*

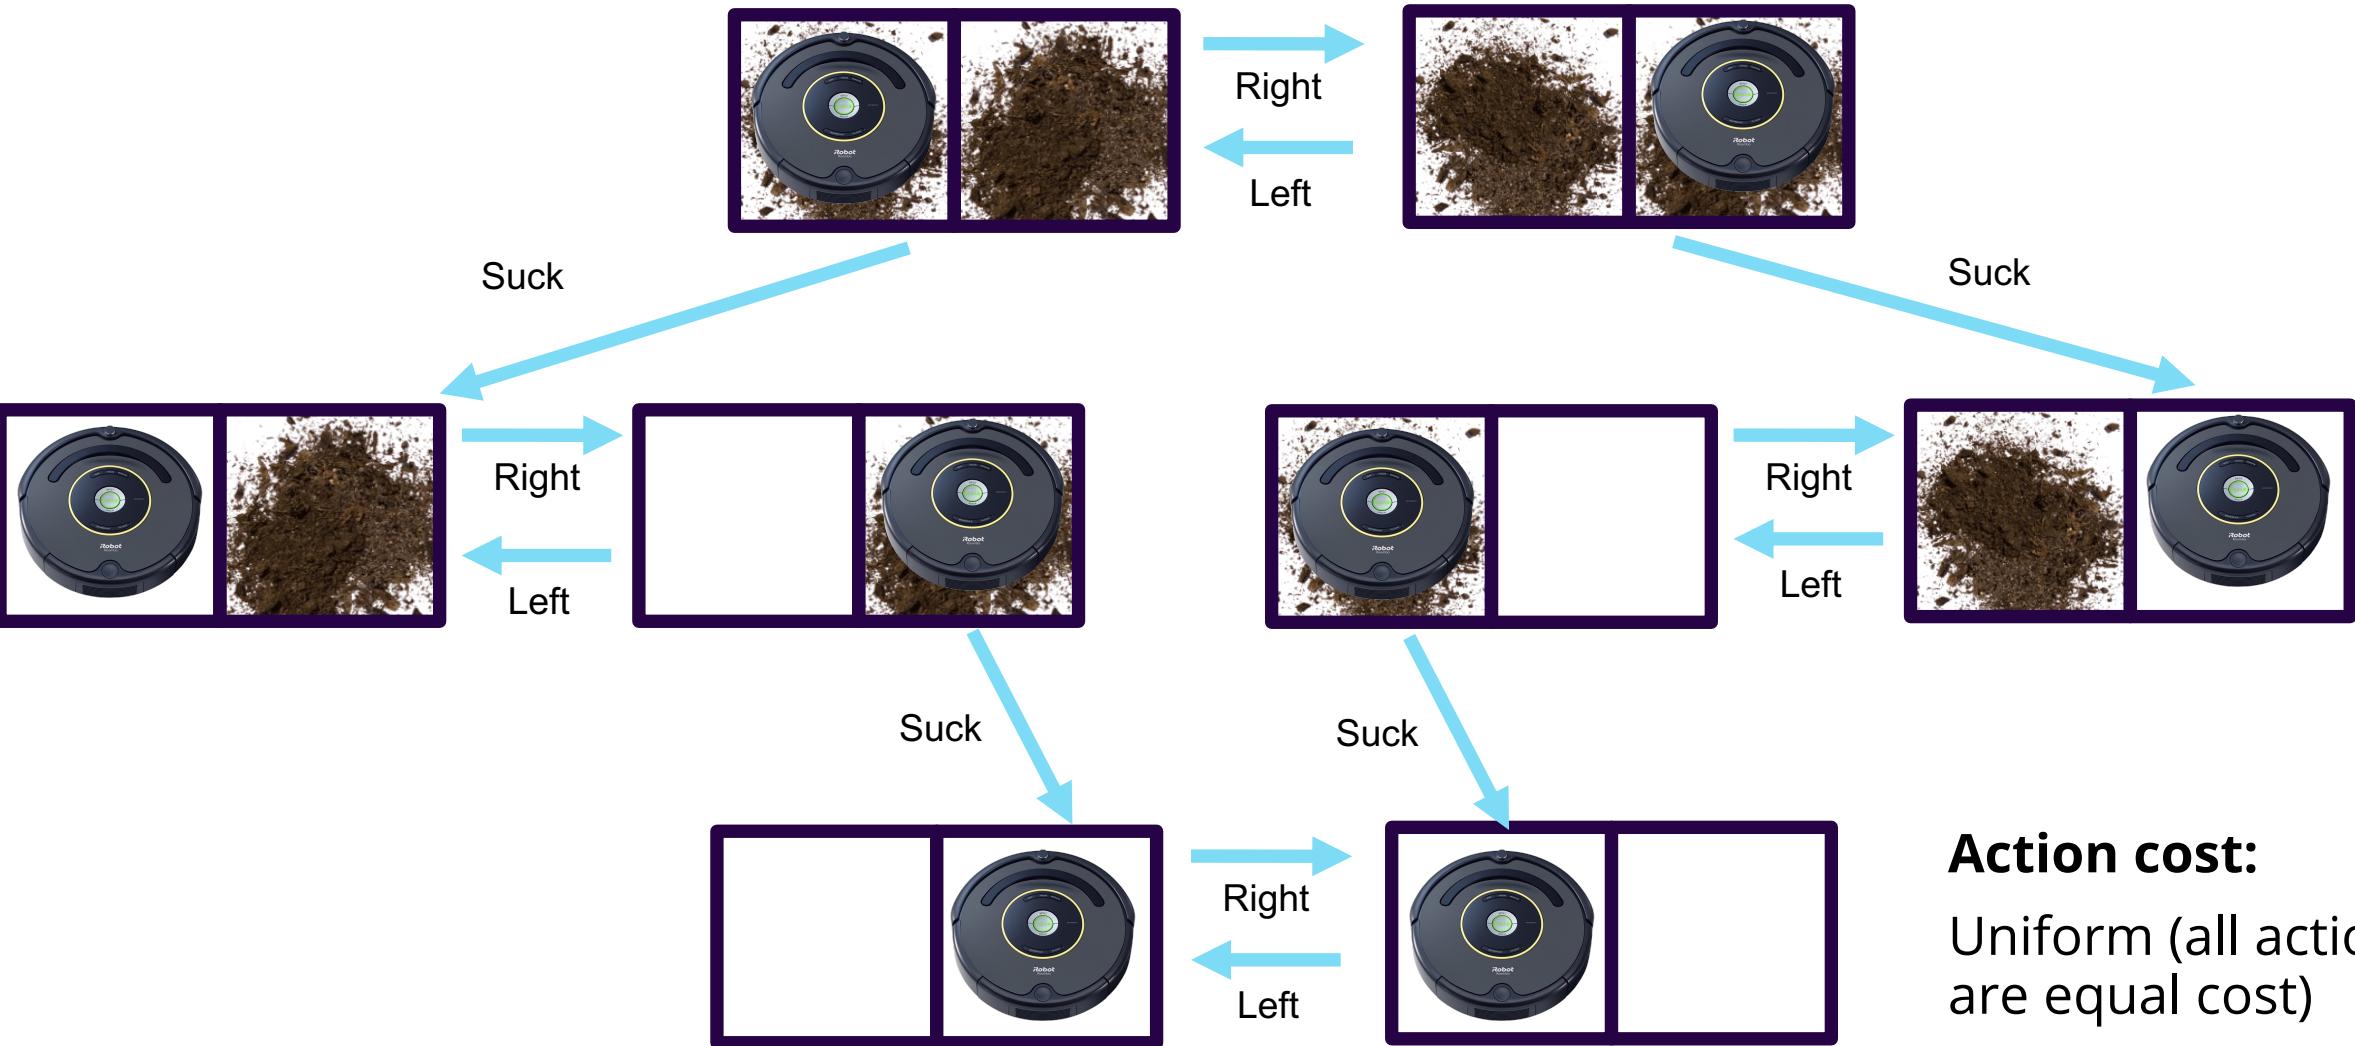
Transition:

- Suck – removes dirt
- Move – moves in that direction, unless agent hits a wall, in which case it stays put.

Vacuum World

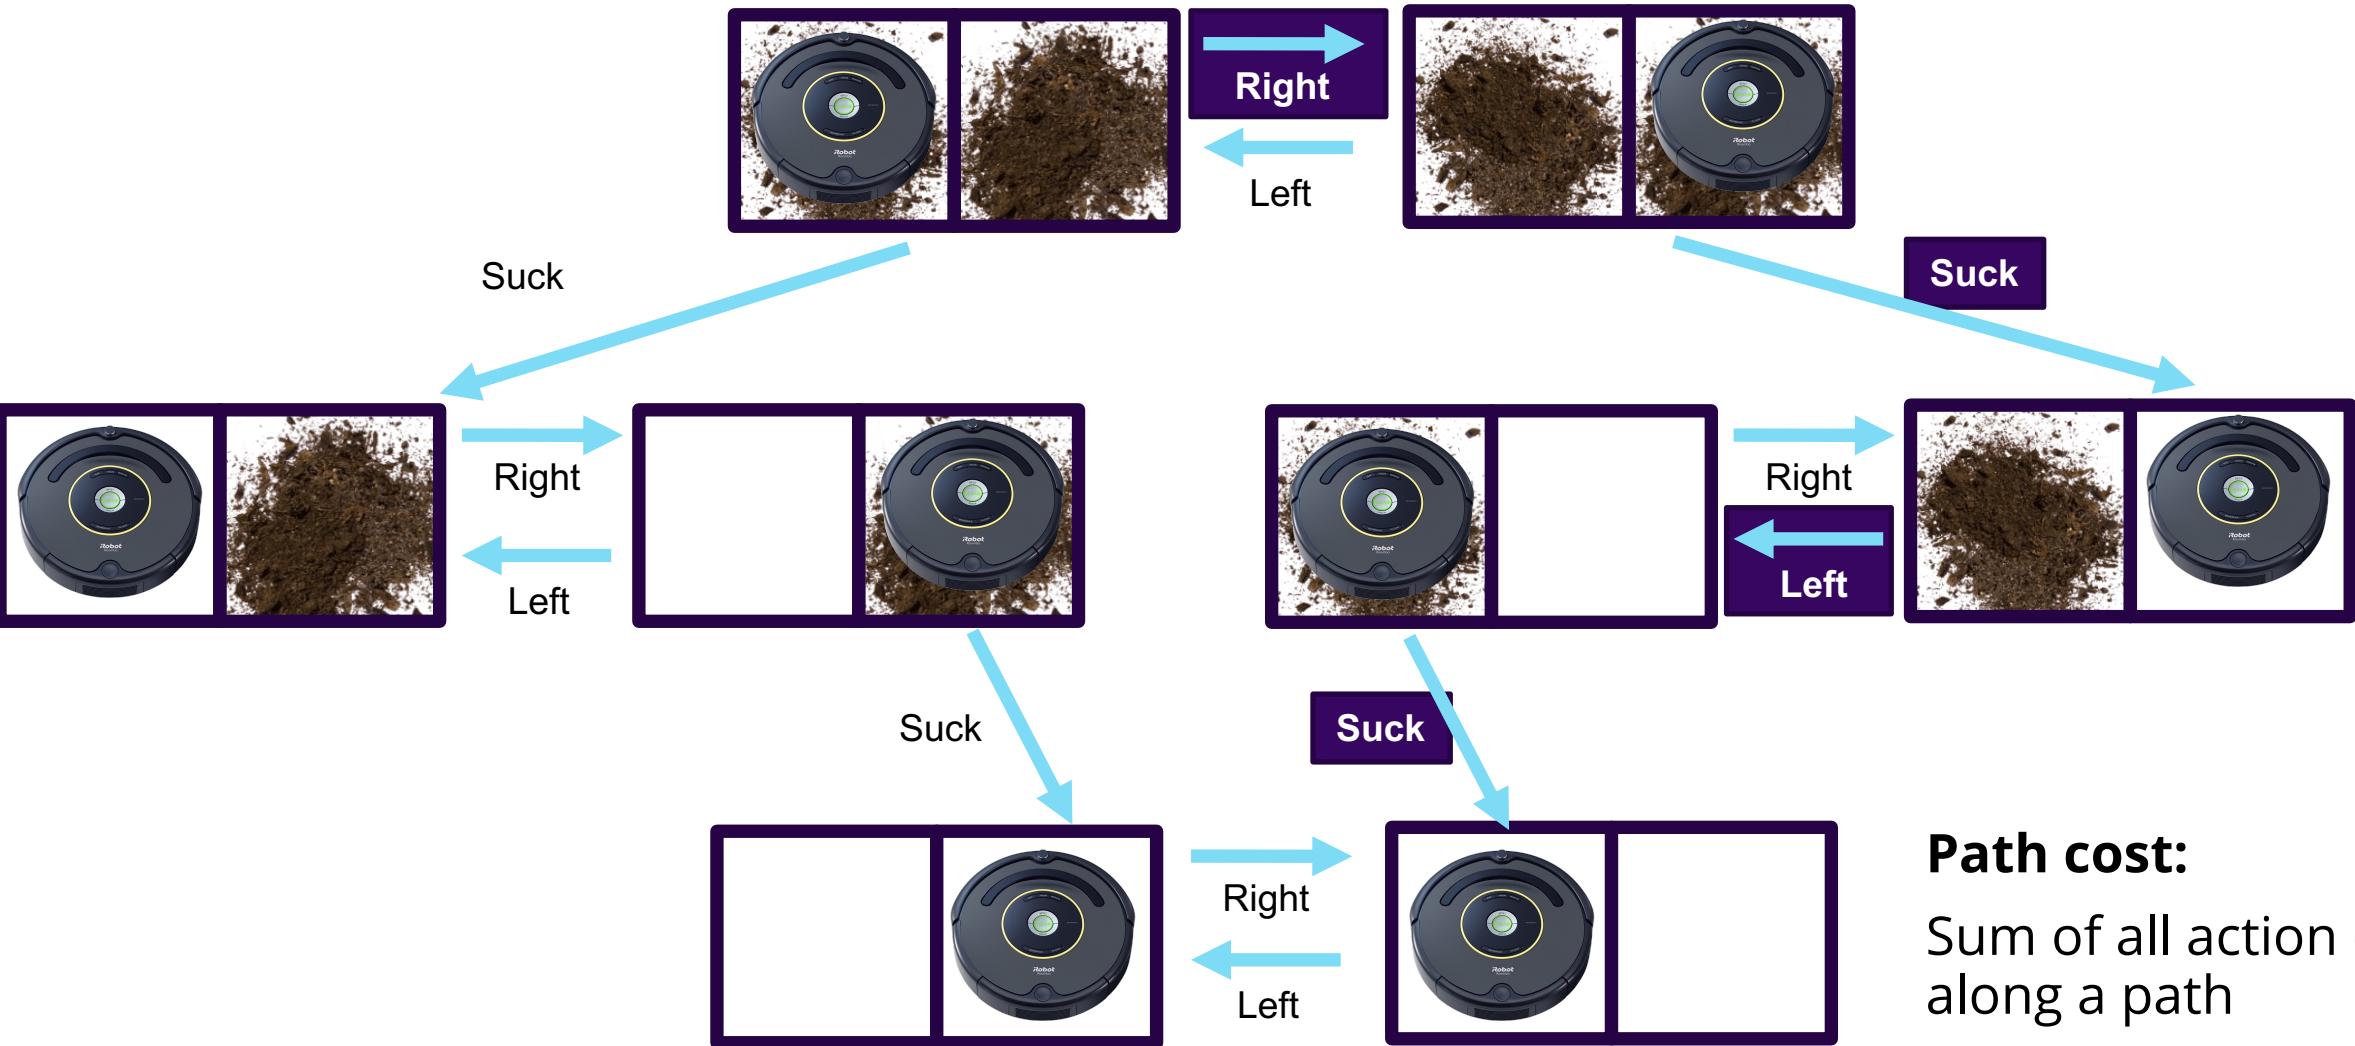


Vacuum World



Action cost:
Uniform (all actions
are equal cost)

Vacuum World



Vacuum World

Initial state



Right
Left



Suck



Right

Left



Suck



Right

Left



Suck

Suck

Right
Left



Solution:

A path from the initial state to a goal state

Goal states

Art: Formulating a Search Problem

Decide:

Which properties matter & how to represent

- *Initial State, Goal State, Possible Intermediate States*

Which actions are possible & how to represent

- *Operator Set: Actions and Transition Model*

Which action is next

- *Path Cost Function*

Formulation greatly affects combinatorics of search space and therefore speed of search

Hard subtask: Selecting a state space

Real world is absurdly complex

State space must be **abstracted** for problem solving

(abstract) **State** = set (equivalence class) of real-world states

(abstract) **Action** = equivalence class of combinations of real-world actions

- e.g. *Arad → Zerind* represents a complex set of possible routes, detours, rest stops, etc
- The abstraction is valid if the path between two states is reflected in the real world

Each abstract action should be “easier” than the real problem

Useful Concepts

State space: the set of all states reachable from the initial state by *any* sequence of actions

- *When several operators can apply to each state, this gets large very quickly*
- *Might be a proper subset of the set of configurations*

Path: a sequence of actions leading from one state s_j to another state s_k

Solution: a path from the initial state s_i to a state s_f that satisfies the goal test

Search tree: a way of representing the paths that a search algorithm has explored. The root is the initial state, leaves of the tree are successor states.

Frontier: those states that are available for *expanding* (for applying legal actions to)

Solutions and *Optimal* Solutions

A **solution** is a sequence of **actions** from the **initial state** to a **goal state**.

Optimal Solution: A solution is **optimal** if no solution has a lower **path cost**.

CIS 4210/5210:
ARTIFICIAL INTELLIGENCE

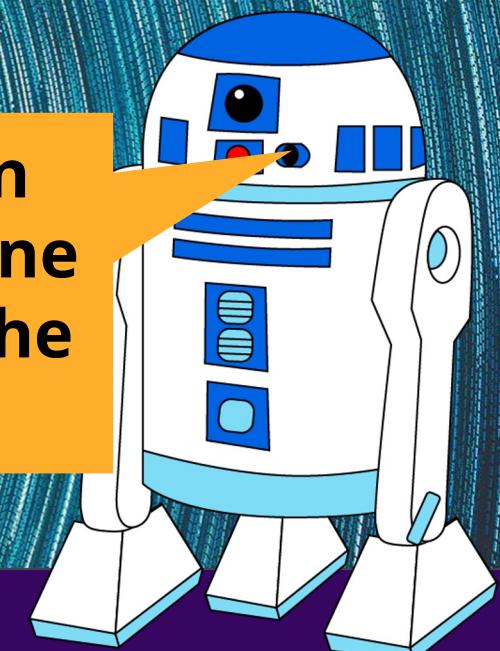
Search Problems

Professor Chris Callison-Burch

HW2 is due
Wednesday, Quiz is
Due Sunday

HW2 Recitation was
Wednesday 8-9pm.
Recording is available

My OHs will be in-person on
Mondays from 2-4pm in Towne
M70 (mezzanine level near the
fishbowl computing lab)



Basic Search Algorithms

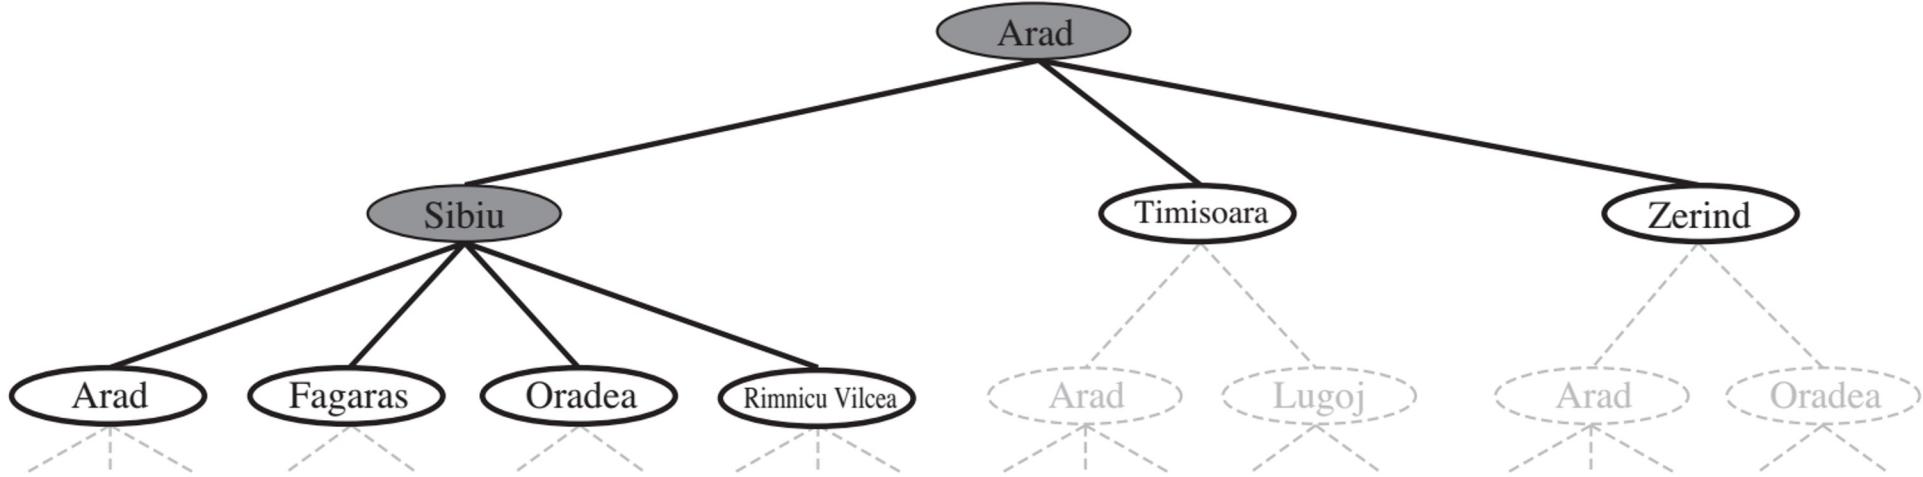
Basic search algorithms: *Tree Search*

Generalized algorithm to solve search problems

Enumerate in some order all possible paths from the initial state

- Here: search through ***explicit tree generation***
 - ROOT= initial state.
 - Nodes in search tree generated through ***transition model***
 - Tree search treats different paths to the same node as distinct

Generalized tree search



```
function TREE-SEARCH(problem, strategy) return a solution or failure
    Initialize frontier to the initial state of the problem
    do
        if the frontier is empty then return failure
        choose leaf node for expansion according to strategy & remove from frontier
        if node contains goal state then return solution
        else expand the node and add resulting nodes to the frontier
```

The strategy determines search process!

States Versus Nodes

A **state** is a (representation of a) **physical configuration**

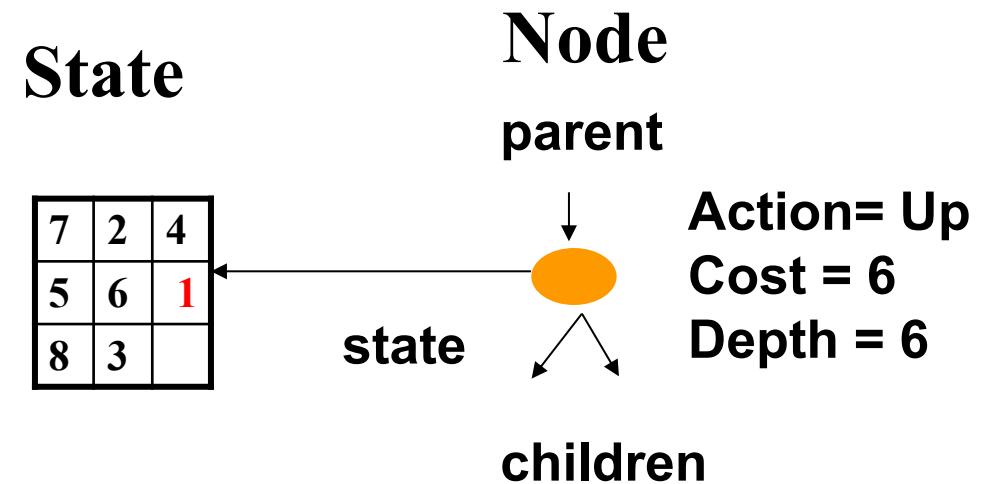
A **node** is a data structure constituting **part of a search tree**

- Also includes *parent, children, depth, path cost $g(x)$*
- Here *node = <state, parent-node, children, action, path-cost, depth>*

States do not have parents, children, depth or path cost!

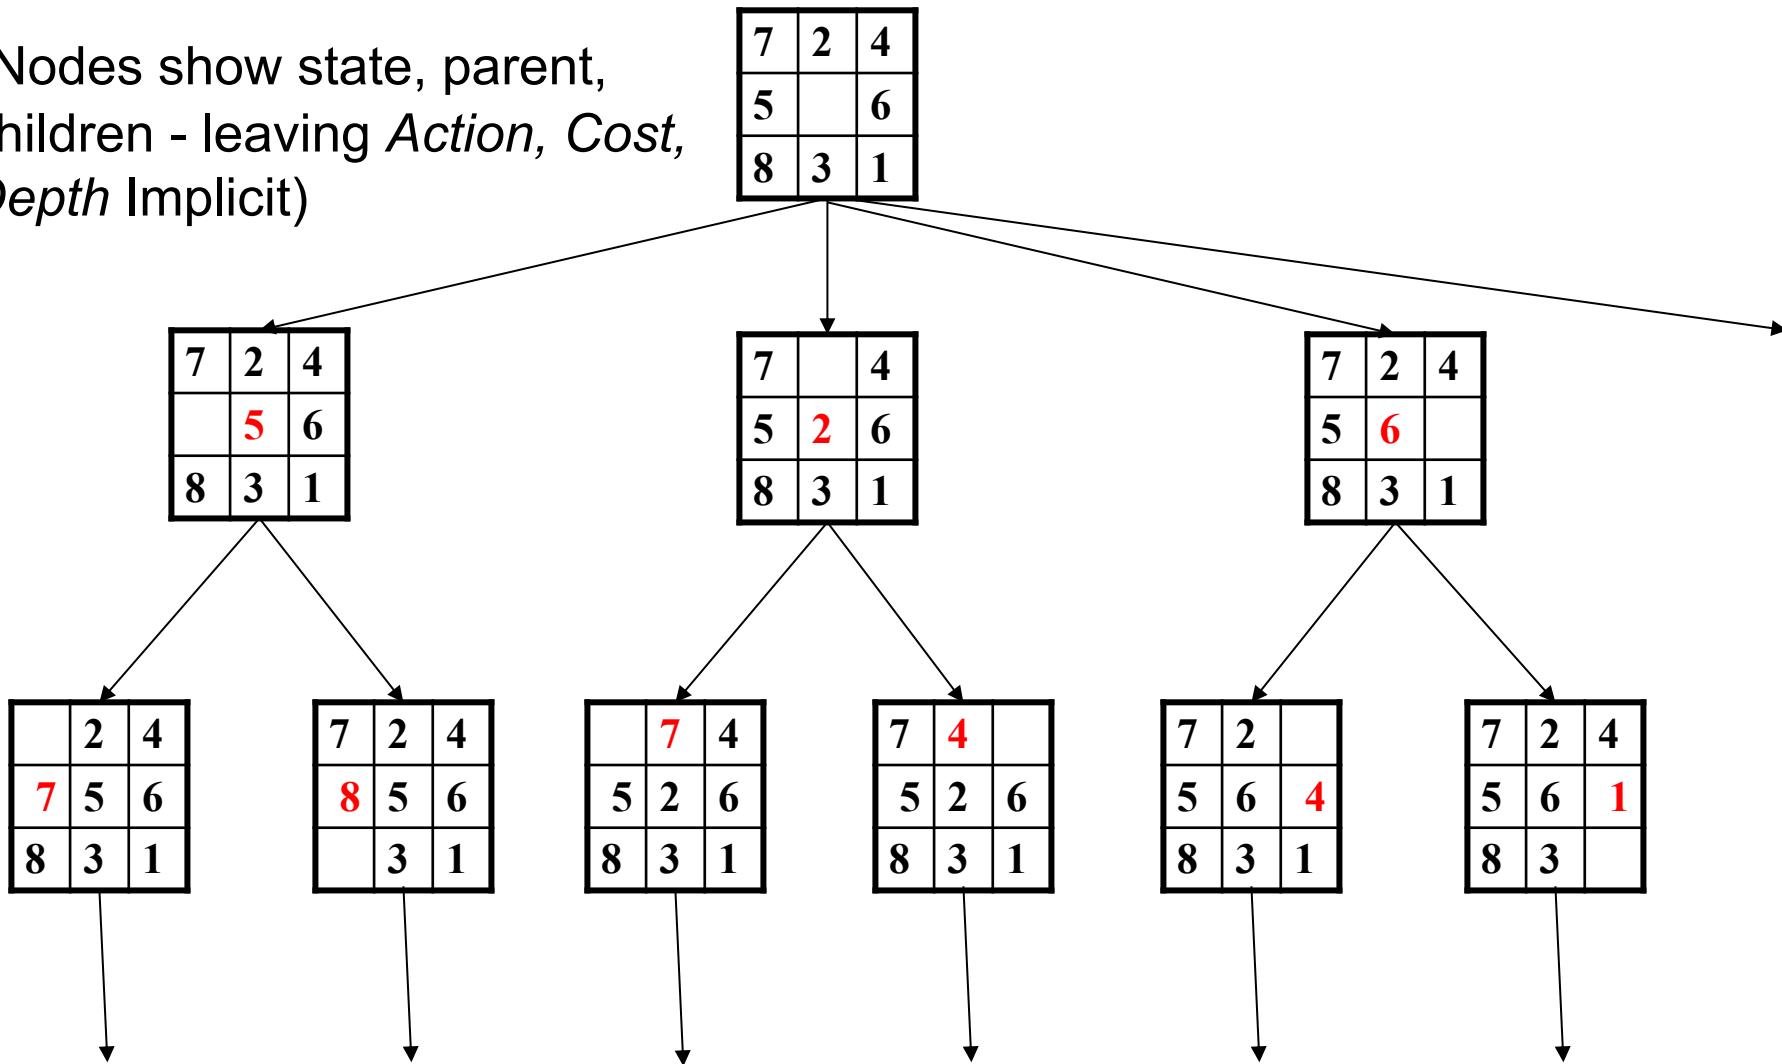
The EXPAND function

- uses the Actions and Transition Model to create the corresponding states
 - creates new nodes,
 - fills in the various fields



8-Puzzle Search Tree

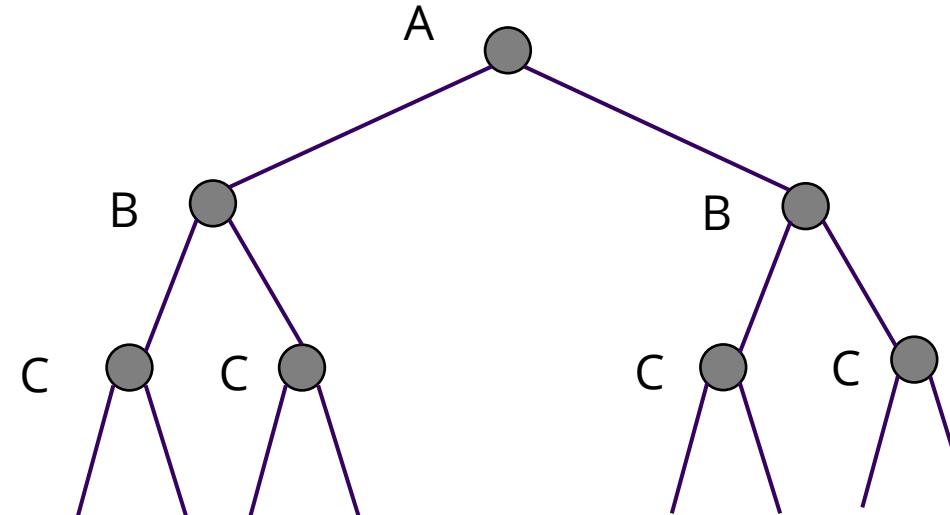
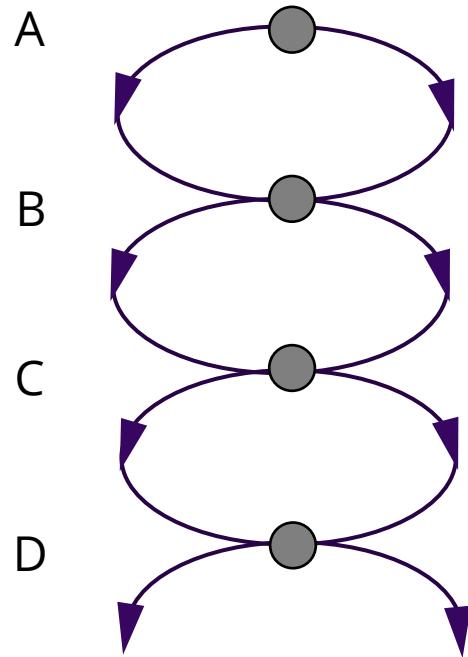
(Nodes show state, parent,
children - leaving *Action, Cost,
Depth* Implicit)



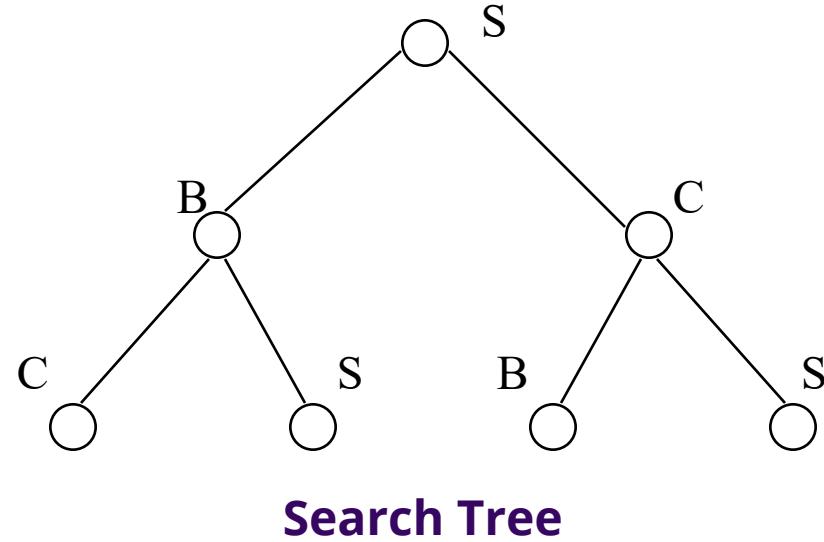
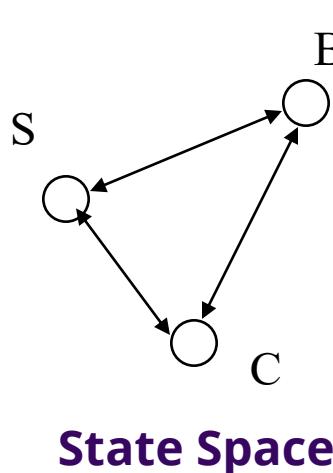
(Suppressing
useless
“backwards”
moves)

Problem: Repeated states

Failure to detect ***repeated states*** can turn a linear problem into an ***exponential*** one!



Solution: Graph Search!



Graph search

- Simple Mod from tree search: *Check to see if a node has been visited before adding to search queue*
 - must keep track of all possible states (can use a lot of memory)
 - e.g., 8-puzzle problem, we have $9!/2 \approx 182K$ states

Graph Search vs Tree Search

function TREE-SEARCH(*problem*) **returns** a solution, or failure

 initialize the frontier using the initial state of *problem*

loop do

if the frontier is empty **then return** failure

 choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

 expand the chosen node, adding the resulting nodes to the frontier

function GRAPH-SEARCH(*problem*) returns a solution, or failure

 initialize the frontier using the initial state of *problem*

initialize the explored set to be empty

loop do

if the frontier is empty **then return** failure

 choose a leaf node and remove it from the frontier

if the node contains a goal state **then return** the corresponding solution

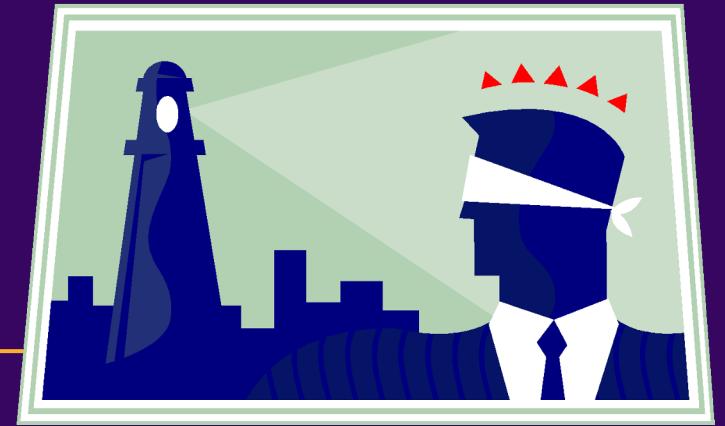
add node to the explored set

 expand the chosen node, adding the resulting nodes to the frontier

only if not in the frontier of explored set

Uninformed Search Strategies

AIMA 3.3-3.4



***Uninformed* search strategies:**

AKA “Blind search”

Uses only information available in problem definition

Informally:

Uninformed search: All non-goal nodes in frontier look equally good

Informed search: Some non-goal nodes can be ranked above others.

Search Strategies

Review: **Strategy** = order of tree expansion

- Implemented by different queue structures (LIFO, FIFO, priority)

Dimensions for evaluation

- **Completeness**- always find the solution?
- **Optimality** - finds a least cost solution (lowest path cost) first?
- **Time complexity** - # of nodes generated (*worst case*)
- **Space complexity** - **# of nodes simultaneously in memory** (*worst case*)

Time/space complexity variables

- b , *maximum branching factor* of search tree
- d , *depth* of the shallowest goal node
- m , maximum length of any path in the state space (potentially ∞)

Introduction to *space* complexity

You know about:

- “Big O” notation
- *Time complexity*

Space complexity is analogous to time complexity

Units of space are arbitrary

- Doesn’t matter because Big O notation ignores constant multiplicative factors
- Plausible Space units:
 - One Memory word
 - Size of any fixed size data structure
 - For example, size of fixed size node in search tree

Breadth-First Search and Depth-First Search

Breadth-first search

Idea:

- Expand *shallowest* unexpanded node

Implementation:

- *frontier* is FIFO (First-In-First-Out) Queue:
 - Put successors at the *end* of *frontier* successor list.

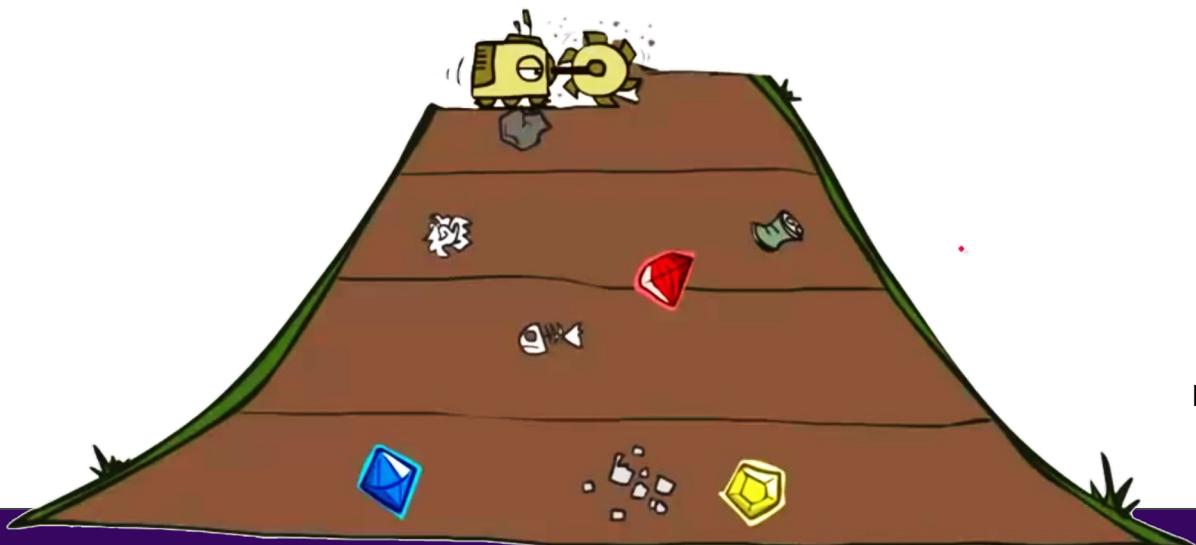


Image credit: Dan Klein and
Pieter Abbeel
<http://ai.berkeley.edu>

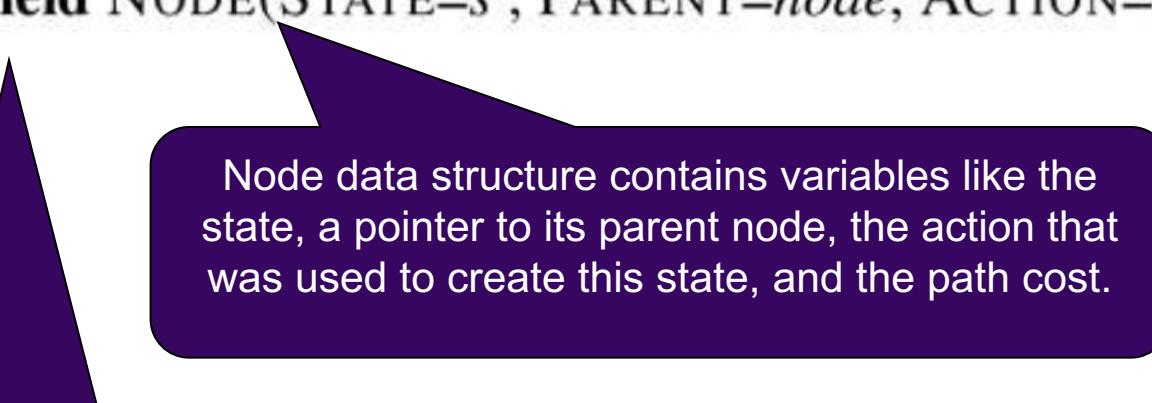
Breadth-first search

```
function BREADTH-FIRST-SEARCH(problem) returns a solution node or failure
  node  $\leftarrow$  NODE(problem.INITIAL)
  if problem.IS-GOAL(node.STATE) then return node
  frontier  $\leftarrow$  a FIFO queue, with node as an element
  reached  $\leftarrow$  {problem.INITIAL}
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if problem.IS-GOAL(s) then return child
      if s is not in reached then
        add s to reached
        add child to frontier
  return failure
```

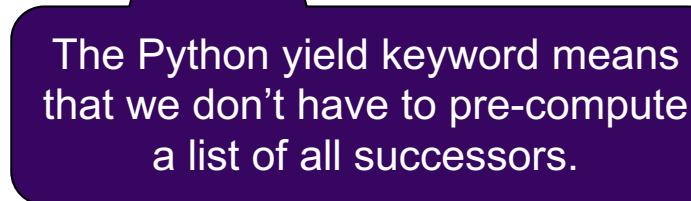
Position within
queue of new items
determines search
strategy

Breadth-first search

```
function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```



Node data structure contains variables like the state, a pointer to its parent node, the action that was used to create this state, and the path cost.



The Python `yield` keyword means that we don't have to pre-compute a list of all successors.

Breadth-first search

```
function BREADTH-FIRST-SEARCH(problem) returns a solution node or failure
  node  $\leftarrow$  NODE(problem.INITIAL)
  if problem.IS-GOAL(node.STATE) then return node
  frontier  $\leftarrow$  a FIFO queue, with node as an element
  reached  $\leftarrow$  {problem.INITIAL}
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if problem.IS-GOAL(s) then return child
      if s is not in reached then
        add s to reached
        add child to frontier
  return failure
```

Subtle: Node inserted into queue only after testing to see if it is a goal state

Properties of breadth-first search

Complete?

Yes (if b is finite)

Optimal?

Yes, if cost = 1 per step
(not optimal in general)

Time Complexity?

$1+b+b^2+b^3+\dots+b^d = O(b^d)$

Space Complexity?

$O(b^d)$ (keeps every node in memory)

b : maximum branching factor of search tree

d : depth of the least cost solution

m : maximum depth of the state space (∞)

Exponential Space (and time) Is Not Good...

- Exponential complexity uninformed search problems *cannot* be solved for any but the smallest instances.
- (*Memory* requirements are a bigger problem than *execution* time.)

DEPTH	NODES	TIME	MEMORY
2	110	0.11 milliseconds	10 ⁷ kilobytes
4	11110	11 milliseconds	10.6 megabytes
6	10 ⁶	1.1 seconds	1 gigabytes
8	10 ⁸	2 minutes	103 gigabytes
10	10 ¹⁰	3 hours	10 terabytes
12	10 ¹²	13 days	1 petabytes
14	10 ¹⁴	3.5 years	99 petabytes

Assumes b=10, 1M nodes/sec, 1000 bytes/node

Depth-first search

Idea:

- Expand *deepest* unexpanded node

Implementation:

- *frontier* is LIFO (Last-In-First-Out) Queue:
 - Put successors at the *front* of *frontier* successor list.



Image credit: Dan Klein and
Pieter Abbeel
<http://ai.berkeley.edu>

Properties of depth-first search

Complete? No. It fails in infinite-depth spaces, spaces with loops

If we modify it to avoid repeated states along path, then it is complete in finite spaces

Optimal? No

Time? $O(b^m)$. This is terrible if m is much larger than d

- but if solutions are dense, may be much faster than breadth-first

Space? $O(b^*m)$, i.e., linear space!

b : maximum branching factor of search tree

d : depth of the least cost solution

m : maximum depth of the state space (∞)

Depth-first vs Breadth-first

Use depth-first if

- *Space is restricted*
- There are many possible solutions with long paths and wrong paths are usually terminated quickly
- Search can be fine-tuned quickly

Use breadth-first if

- *Possible infinite paths*
- Some solutions have short paths
- Can quickly discard unlikely paths

Search Conundrum

Breadth-first

- Complete,
- Optimal
- but uses $O(b^d)$ space*

Depth-first

- Not complete *unless m is bounded*
- Not optimal
- Uses $O(b^m)$ time; terrible if $m \gg d$
- but only uses $O(b*m)$ space*

How can we get the best of both?

Depth-limited search: A building block

Depth-First search *but with depth limit l .*

- i.e. nodes at depth l *have no successors.*
- No infinite-path problem!

If $l = d$ (by luck!), then optimal

- But:
 - If $l < d$ then incomplete 😞
 - If $l > d$ then not optimal 😞

Time complexity: $O(b^l)$

Space complexity: $O(bl)$ 😊

Iterative deepening search

A general strategy to find best depth limit l .

- Key idea: use *Depth-limited search* as subroutine, with increasing l .

```
For  $l = 0$  to  $\infty$  do
    depth-limited-search to level l
    if it succeeds
        then return solution
```

- *Complete & optimal*: Goal is always found at depth d , the depth of the shallowest goal-node.

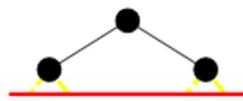
Could this possibly be efficient?

Nodes constructed at each deepening

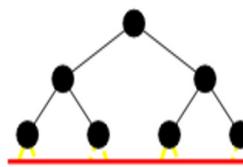
Depth 0: 0 (Given the node, doesn't *construct* it.)



Depth 1: b^1 nodes



Depth 2: b nodes + b^2 nodes



Depth 3: b nodes + b^2 nodes + b^3 nodes

...

Total nodes constructed:

Depth 0: 0 (Given the node, doesn't *construct* it.)

Depth 1: $b^1 = b$ nodes

Depth 2: b nodes + b^2 nodes

- Depth 3: b nodes + b^2 nodes + b^3 nodes
- ...

Suppose the first solution is the last node at depth 3:

Total nodes constructed:

3* b nodes + **2*** b^2 nodes + **1*** b^3 nodes

ID search, Evaluation: Time Complexity

- More generally, the time complexity is
 - $(d)b + (d-1)b^2 + \dots + (1)b^d = O(b^d)$

As efficient in terms of $O(\dots)$ as Breadth First Search:

- $b + b^2 + \dots + b^d = O(b^d)$

ID search, Evaluation

Complete: YES (no infinite paths) 😊

Time complexity: $O(b^d)$

Space complexity: $O(bd)$ 😊

Optimal: YES if step cost is 1. 😊

Summary of algorithms

Criterion	Breadth-First	Depth-First	Depth-limited	Iterative deepening
Complete?	YES	NO	NO	YES
Time	b^d	b^m	b^l	b^d
Space	b^d	bm	bl	bd
Optimal?	YES	NO	NO	YES

Next Up: Informed Search

AIMA 3.5-3.6

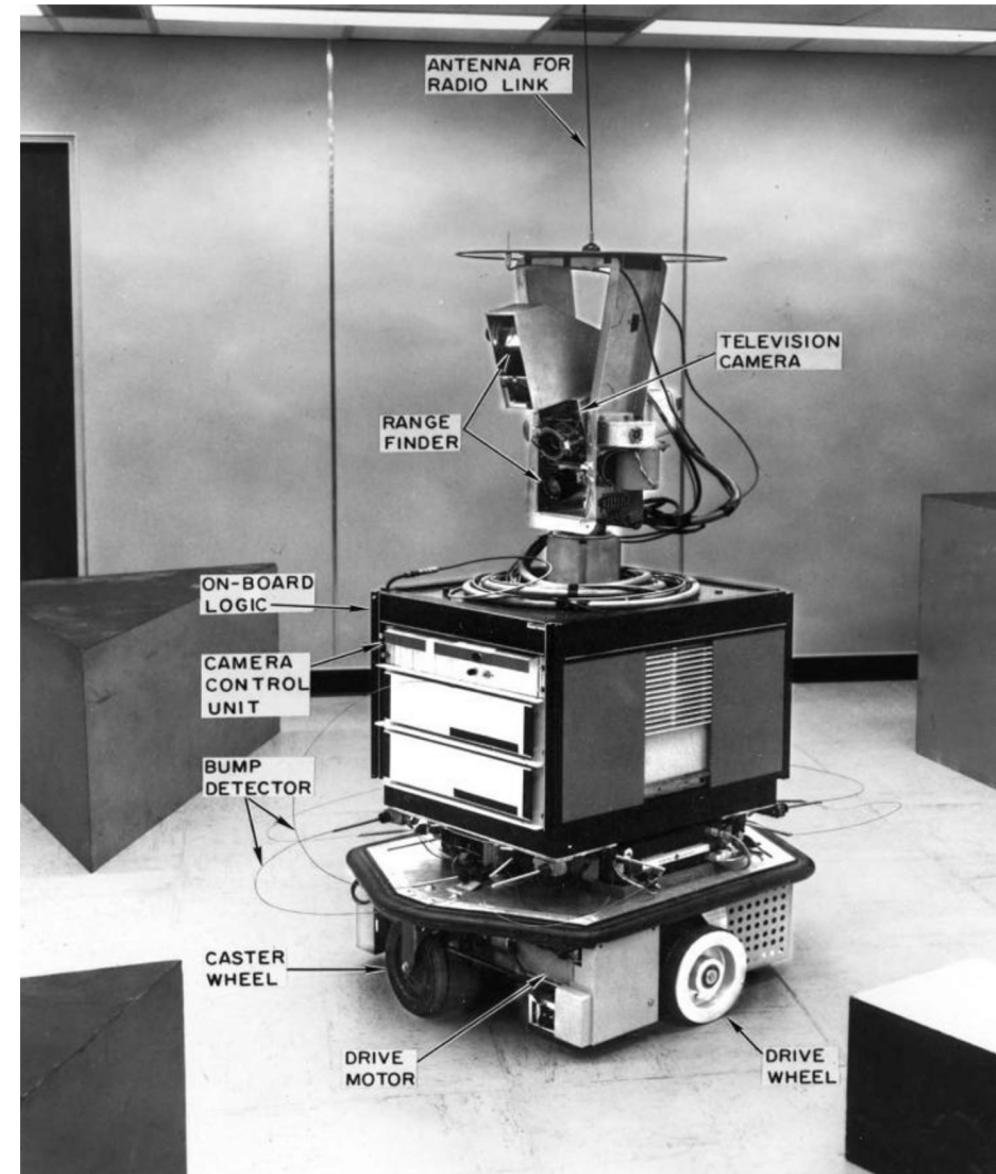
Informed Search

An **informed search** strategy uses **domain-specific information** about the location of the goals in order to find a solution **more efficiently** than uninformed search.

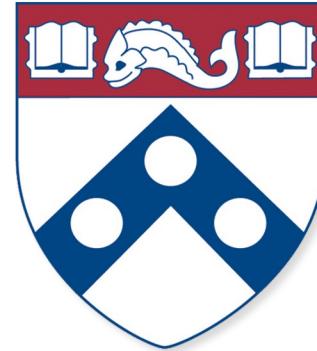
Hints will come as part of a **heuristic function** denoted $h(n)$.

One of the most famous informed search algorithms is **A*** which was developed for **robot navigation**.

Shakey the robot was developed at the Stanford Research Institute from 1966 to 1972.



<https://www.youtube.com/watch?v=7bsEN8mwUB8>



Penn Engineering

Copyright 2021 University of Pennsylvania
No reproduction or distribution without permission.